

## 4. GREEDY ALGORITHMS I

---

- ▶ *coin changing*
- ▶ *interval scheduling*
- ▶ *interval partitioning*
- ▶ *scheduling to minimize lateness*

Lecture slides by Kevin Wayne

Copyright © 2005 Pearson–Addison Wesley

<http://www.cs.princeton.edu/~wayne/kleinberg-tardos>

## 4. GREEDY ALGORITHMS I

---



- ▶ *coin changing*
- ▶ *interval scheduling*
- ▶ *interval partitioning*
- ▶ *scheduling to minimize lateness*

Greedy algorithms

# Coin changing

---

**Goal.** Given U. S. currency denominations { 1, 5, 10, 25, 100 }, devise a method to pay amount to customer using fewest coins.

**Ex.** 34¢.



**Cashier's algorithm.** At each iteration, add coin of the largest value that does not take us past the amount to be paid.

**Ex.** \$2.89.



# Cashier's algorithm

---

At each iteration, add coin of the largest value that does not take us past the amount to be paid.

**CASHIERS-ALGORITHM** ( $x, c_1, c_2, \dots, c_n$ )

**SORT**  $n$  coin denominations so that  $0 < c_1 < c_2 < \dots < c_n$ .

$S \leftarrow \emptyset$ . ← multiset of coins selected

**WHILE** ( $x > 0$ )

$k \leftarrow$  largest coin denomination  $c_k$  such that  $c_k \leq x$ .

**IF** (no such  $k$ )

**RETURN** “*no solution.*”

**ELSE**

$x \leftarrow x - c_k$ .

$S \leftarrow S \cup \{ k \}$ .

**RETURN**  $S$ .



## Is the cashier's algorithm optimal?

- A. Yes, greedy algorithms are always optimal.
- B. Yes, for any set of coin denominations  $c_1 < c_2 < \dots < c_n$  provided  $c_1 = 1$ .
- C. Yes, because of special properties of U.S. coin denominations.
- D. No.



# Cashier's algorithm (for arbitrary coin denominations)

---

Q. Is cashier's algorithm optimal for any set of denominations?

A. No. Consider U.S. postage: 1, 10, 21, 34, 70, 100, 350, 1225, 1500.

- Cashier's algorithm:  $140\text{¢} = 100 + 34 + 1 + 1 + 1 + 1 + 1$ .
- Optimal:  $140\text{¢} = 70 + 70$ .



A. No. It may not even lead to a feasible solution if  $c_1 > 1$ : 7, 8, 9.

- Cashier's algorithm:  $15\text{¢} = 9 + ?$ .
- Optimal:  $15\text{¢} = 7 + 8$ .

# Properties of any optimal solution (for U.S. coin denominations)

---

**Property.** Number of pennies  $\leq 4$ .

Pf. Replace 5 pennies with 1 nickel.

**Property.** Number of nickels  $\leq 1$ .

**Property.** Number of quarters  $\leq 3$ .

**Property.** Number of nickels + number of dimes  $\leq 2$ .

Pf.

- Recall:  $\leq 1$  nickel.
- Replace 3 dimes and 0 nickels with 1 quarter and 1 nickel;
- Replace 2 dimes and 1 nickel with 1 quarter.



**dollars**  
(100¢)

**quarters**  
(25¢)

**dimes**  
(10¢)

**nickels**  
(5¢)

**pennies**  
(1¢)

# Optimality of cashier's algorithm (for U.S. coin denominations)

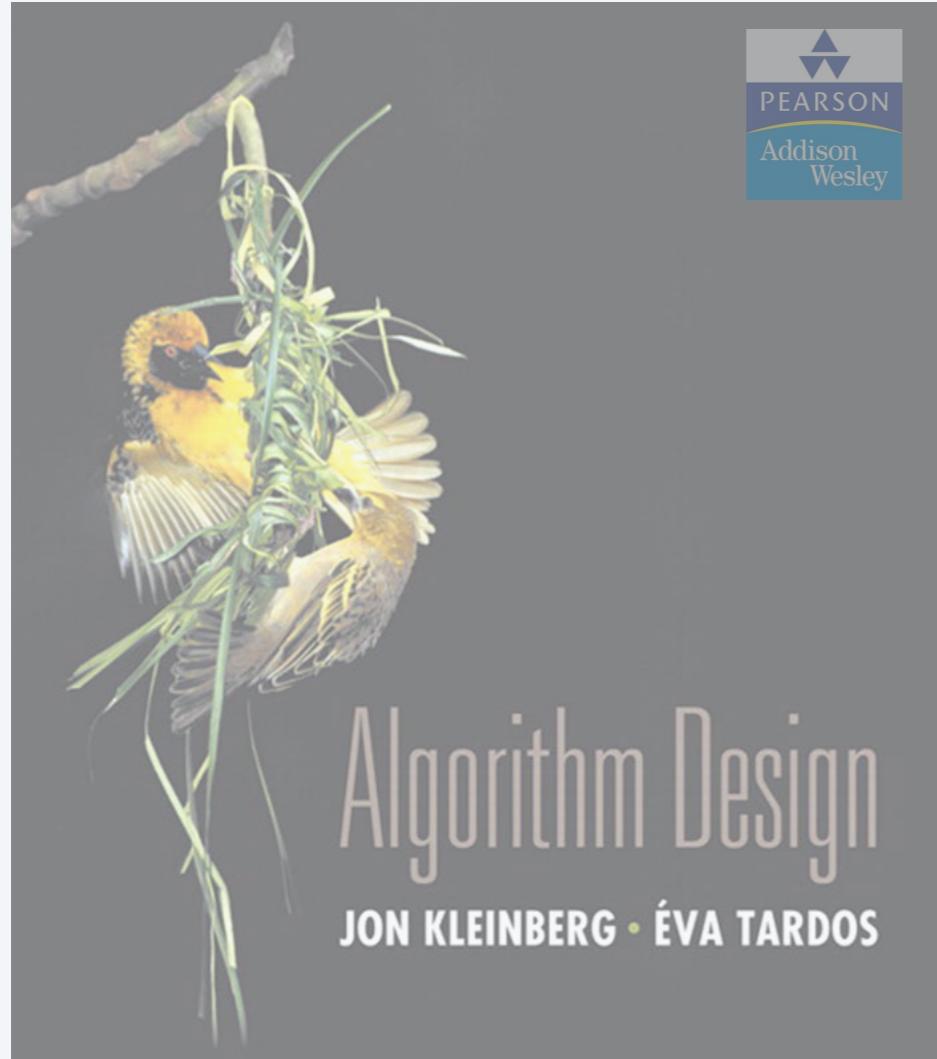
---

**Theorem.** Cashier's algorithm is optimal for U.S. coins { 1, 5, 10, 25, 100 }.

Pf. [ by induction on amount to be paid  $x$  ]

- Consider optimal way to change  $c_k \leq x < c_{k+1}$  : greedy takes coin  $k$ .
- We claim that any optimal solution must take coin  $k$ .
  - if not, it needs enough coins of type  $c_1, \dots, c_{k-1}$  to add up to  $x$
  - table below indicates no optimal solution can do this
- Problem reduces to coin-changing  $x - c_k$  cents, which, by induction, is optimally solved by cashier's algorithm. ■

$k$	$c_k$	all optimal solutions must satisfy	max value of coin denominations $c_1, c_2, \dots, c_{k-1}$ in any optimal solution
1	1	$P \leq 4$	—
2	5	$N \leq 1$	4
3	10	$N + D \leq 2$	$4 + 5 = 9$
4	25	$Q \leq 3$	$20 + 4 = 24$
5	100	<i>no limit</i>	$75 + 24 = 99$



## SECTION 4.1

# 4. GREEDY ALGORITHMS I

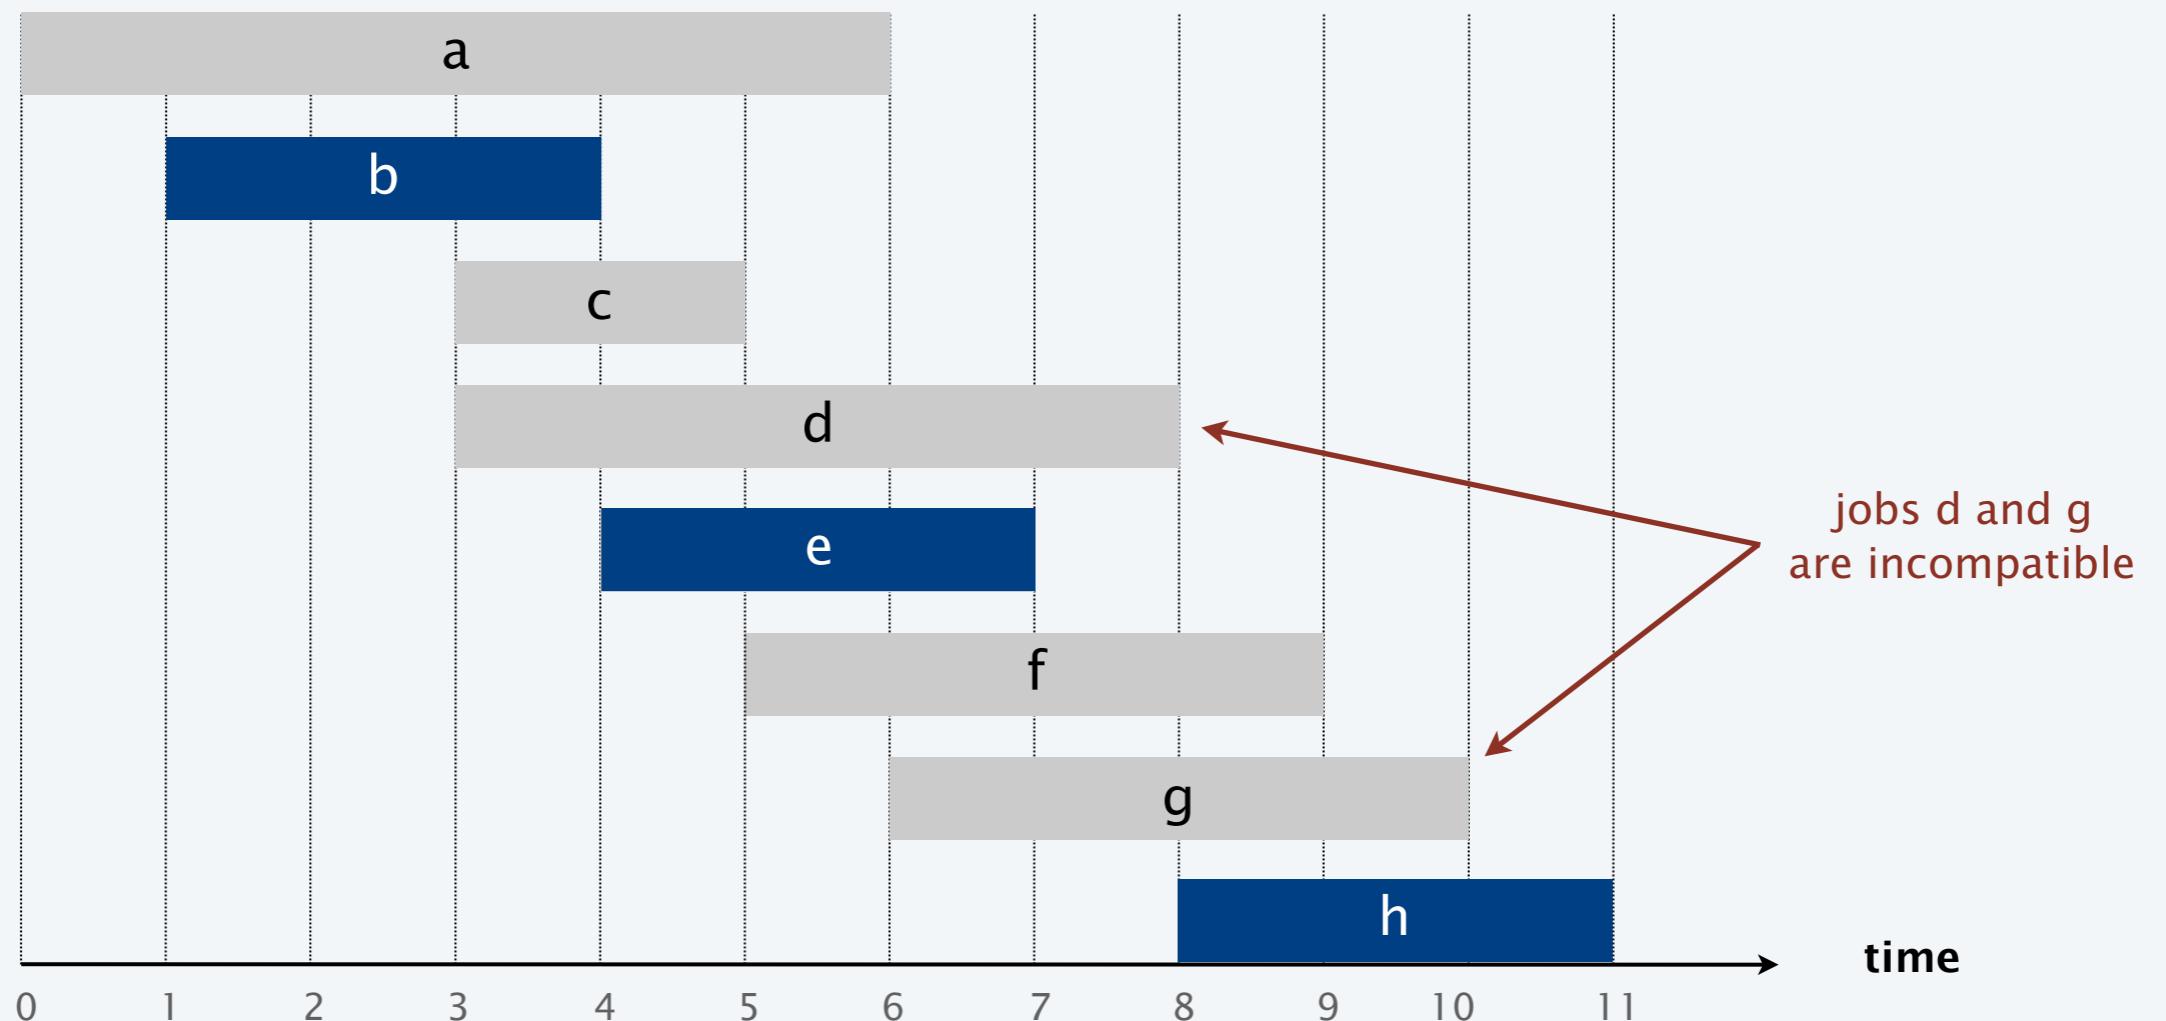
---

- ▶ *coin changing*
- ▶ *interval scheduling*
- ▶ *interval partitioning*
- ▶ *scheduling to minimize lateness*
- ▶ *optimal caching*

# Interval scheduling

---

- Job  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Two jobs are **compatible** if they don't overlap.
- Goal: find maximum subset of mutually compatible jobs.





**Consider jobs in some order, taking each job provided it's compatible with the ones already taken. Which rule is optimal?**

- A. [Earliest start time] Consider jobs in ascending order of  $s_j$ .
- B. [Earliest finish time] Consider jobs in ascending order of  $f_j$ .
- C. [Shortest interval] Consider jobs in ascending order of  $f_j - s_j$ .
- D. None of the above.

# Interval scheduling: earliest-finish-time-first algorithm



EARLIEST-FINISH-TIME-FIRST ( $n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$ )

SORT jobs by finish times and renumber so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .

$S \leftarrow \emptyset$ . ← set of jobs selected

FOR  $j = 1$  TO  $n$

IF (job  $j$  is compatible with  $S$ )

$S \leftarrow S \cup \{ j \}$ .

RETURN  $S$ .

Proposition. Can implement earliest-finish-time first in  $O(n \log n)$  time.

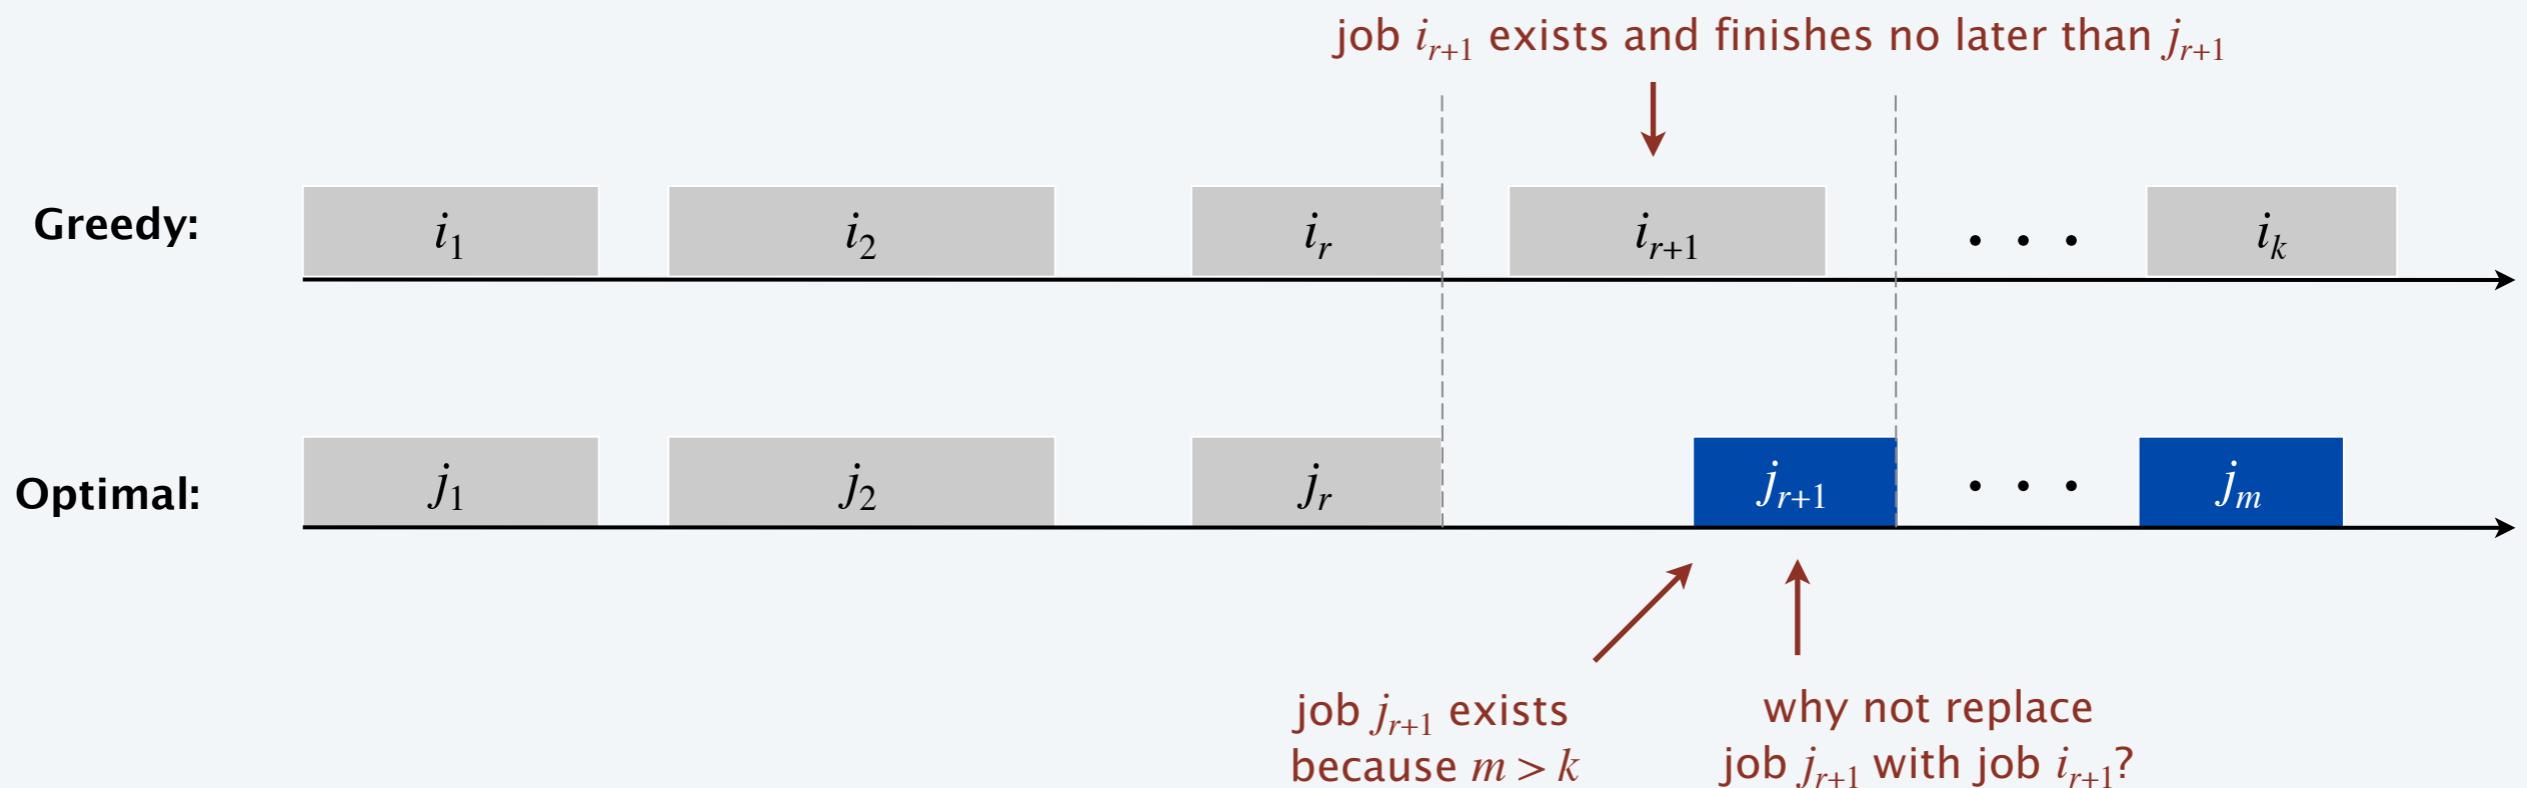
- Keep track of job  $j^*$  that was added last to  $S$ .
- Job  $j$  is compatible with  $S$  iff  $s_j \geq f_{j^*}$ .
- Sorting by finish times takes  $O(n \log n)$  time.

# Interval scheduling: analysis of earliest-finish-time-first algorithm

Theorem. The earliest-finish-time-first algorithm is optimal.

Pf. [by contradiction]

- Assume greedy is not optimal, and let's see what happens.
- Let  $i_1, i_2, \dots, i_k$  denote set of jobs selected by greedy.
- Let  $j_1, j_2, \dots, j_m$  denote set of jobs in an optimal solution with  $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$  for the largest possible value of  $r$ .

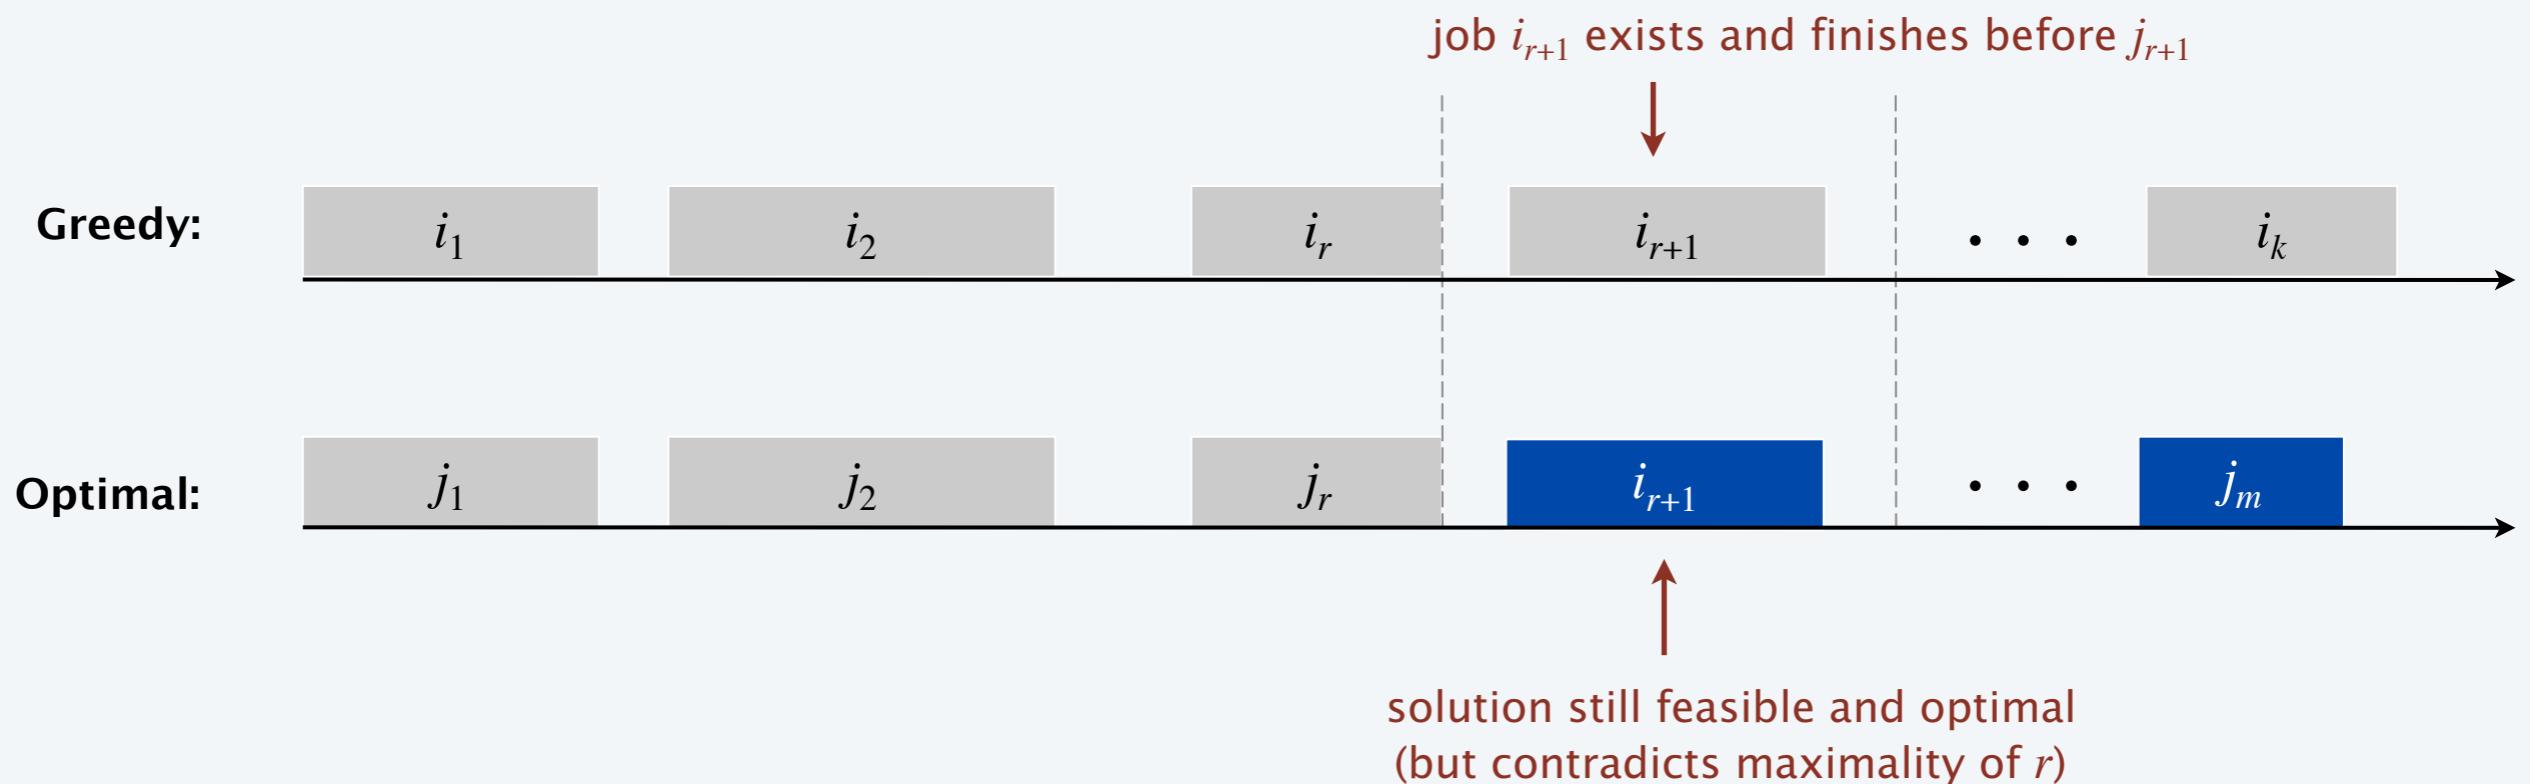


# Interval scheduling: analysis of earliest-finish-time-first algorithm

**Theorem.** The earliest-finish-time-first algorithm is optimal.

**Pf.** [by contradiction]

- Assume greedy is not optimal, and let's see what happens.
- Let  $i_1, i_2, \dots, i_k$  denote set of jobs selected by greedy.
- Let  $j_1, j_2, \dots, j_m$  denote set of jobs in an optimal solution with  $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$  for the largest possible value of  $r$ .



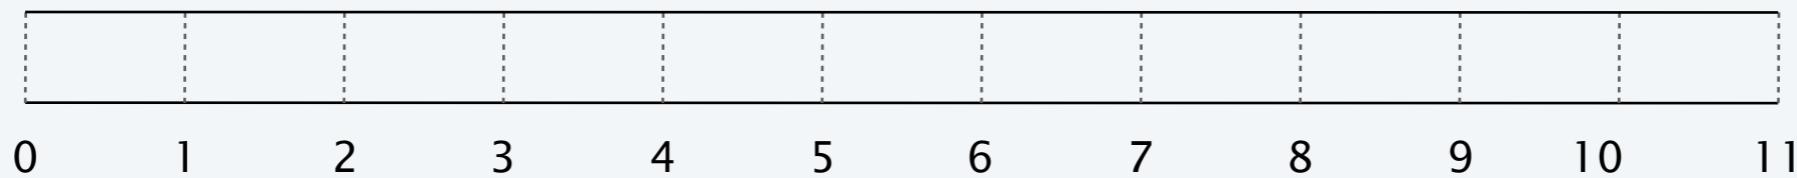
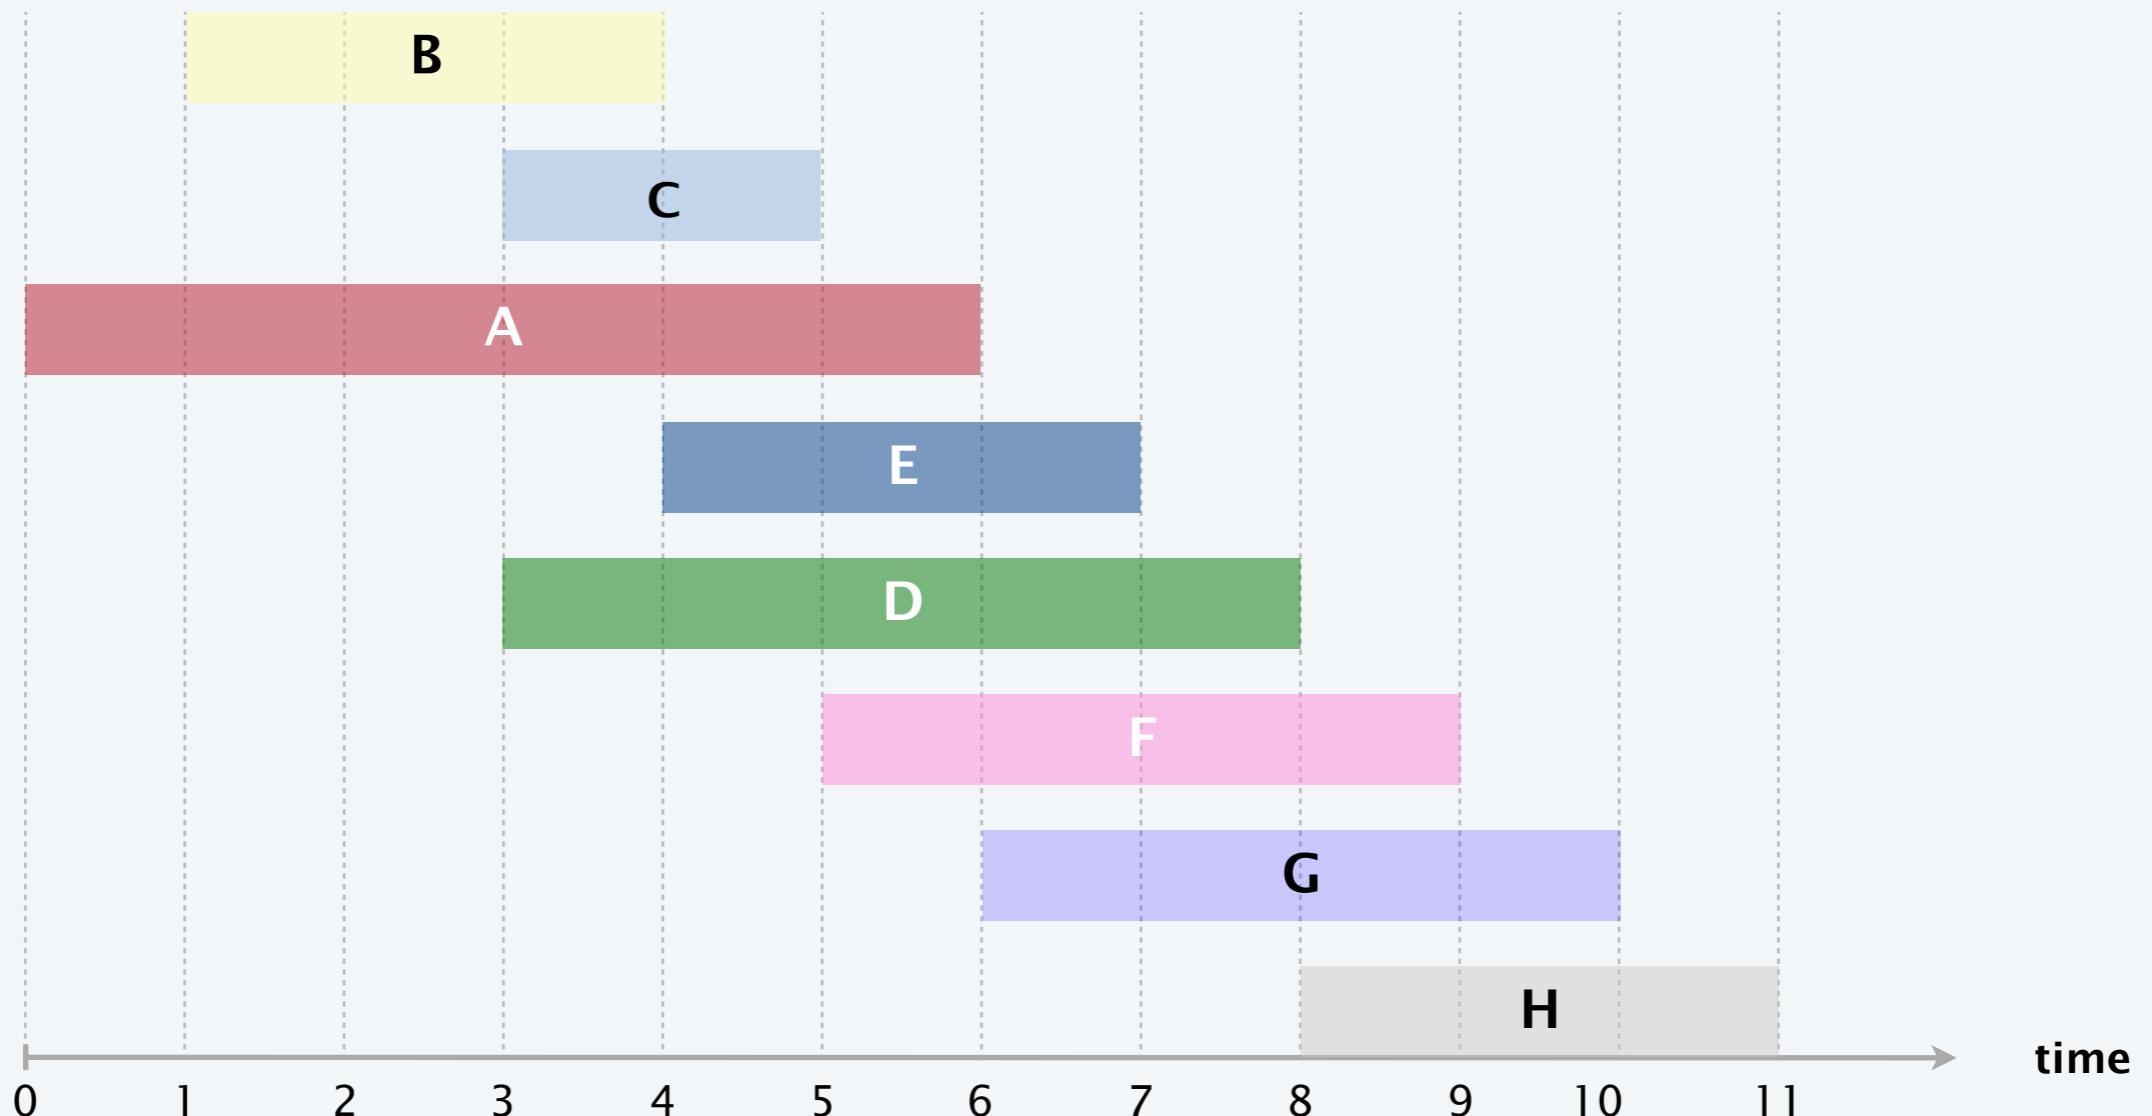


Suppose that each job also has a positive weight and the goal is to find a maximum weight subset of mutually compatible intervals.  
Is the earliest-finish-time-first algorithm still optimal?

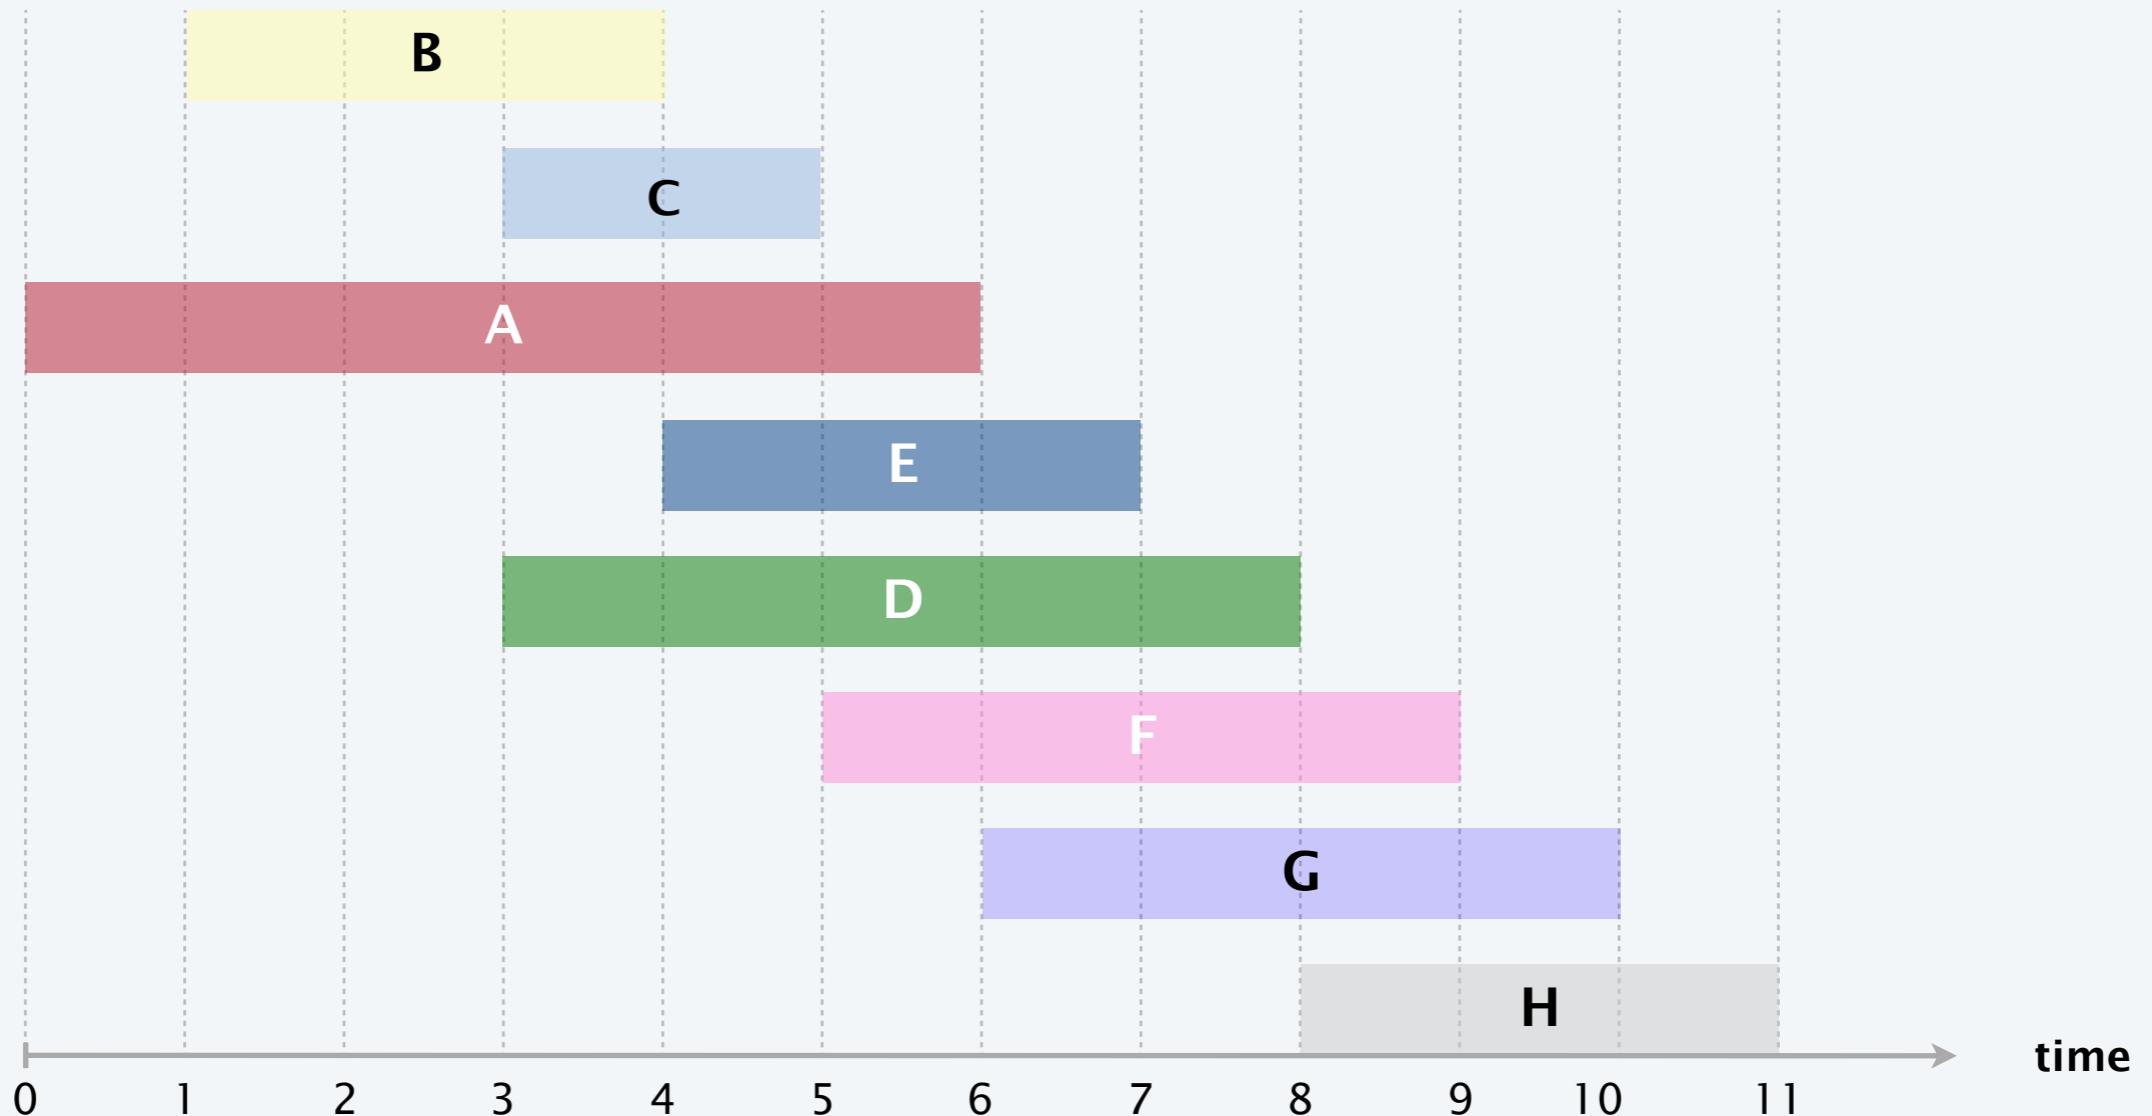
- A. Yes, because greedy algorithms are always optimal.
- B. Yes, because the same proof of correctness is valid.
- C. No, because the same proof of correctness is no longer valid.
- D. No, because you could assign a huge weight to a job that overlaps the job with the earliest finish time.

# Earliest-finish-time-first algorithm demo

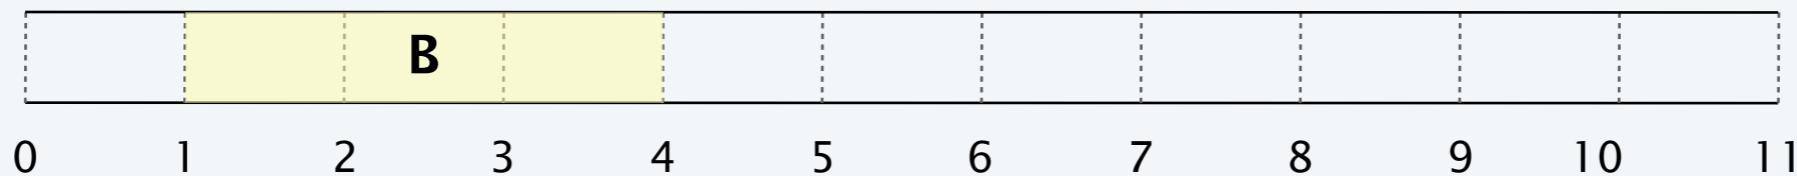
---



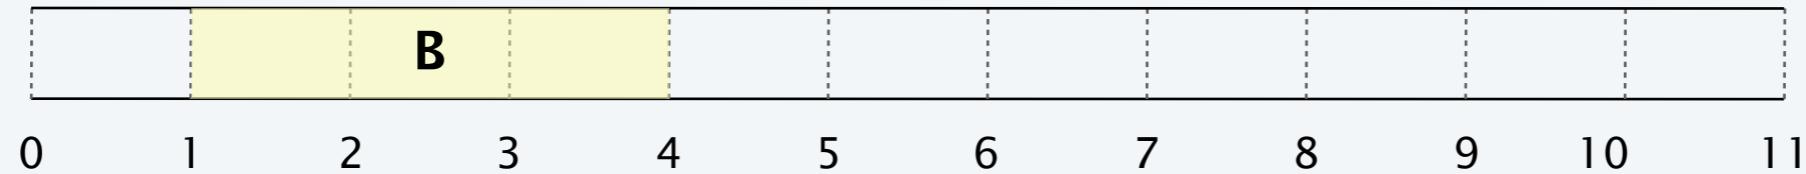
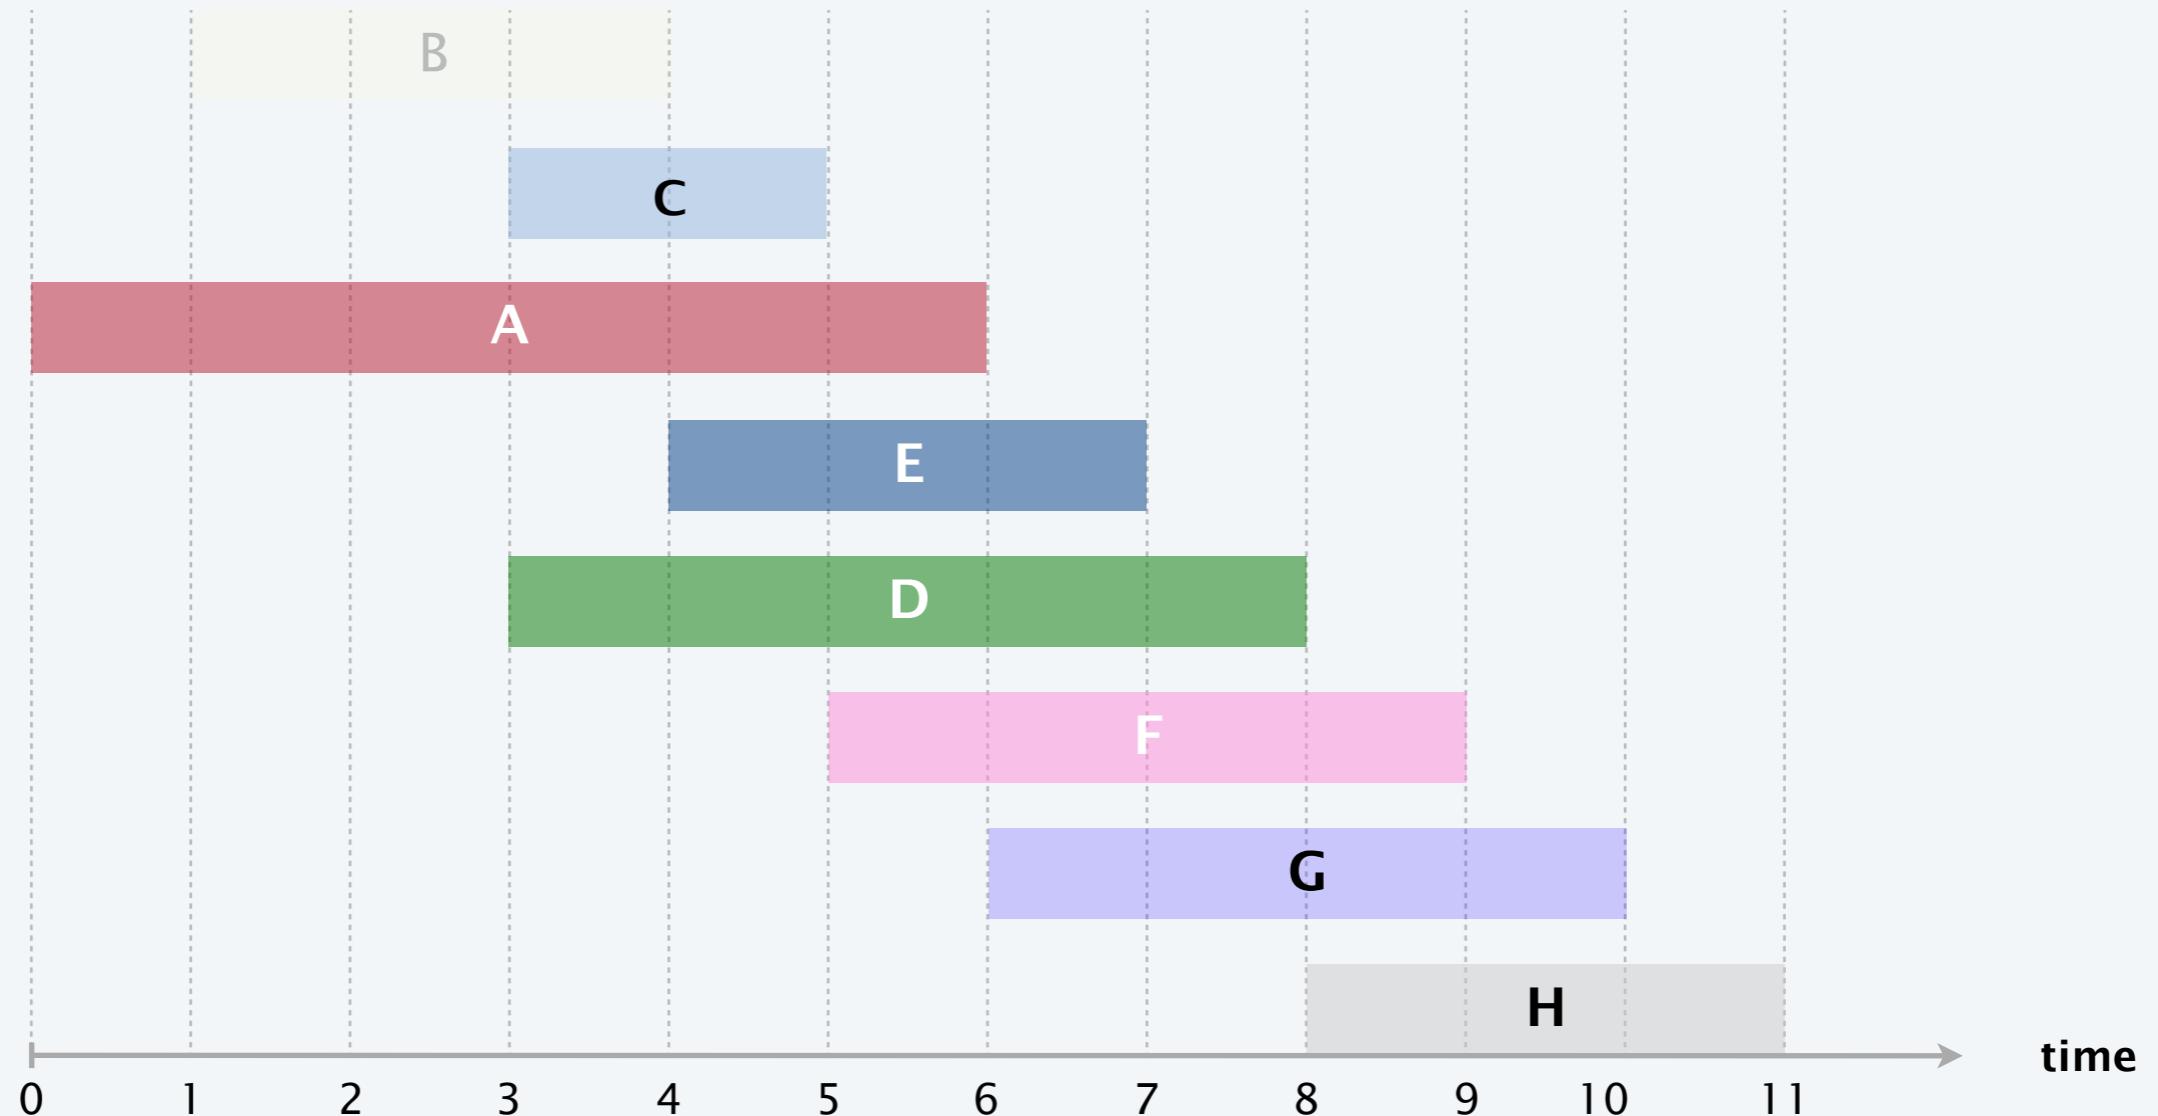
# Earliest-finish-time-first algorithm demo



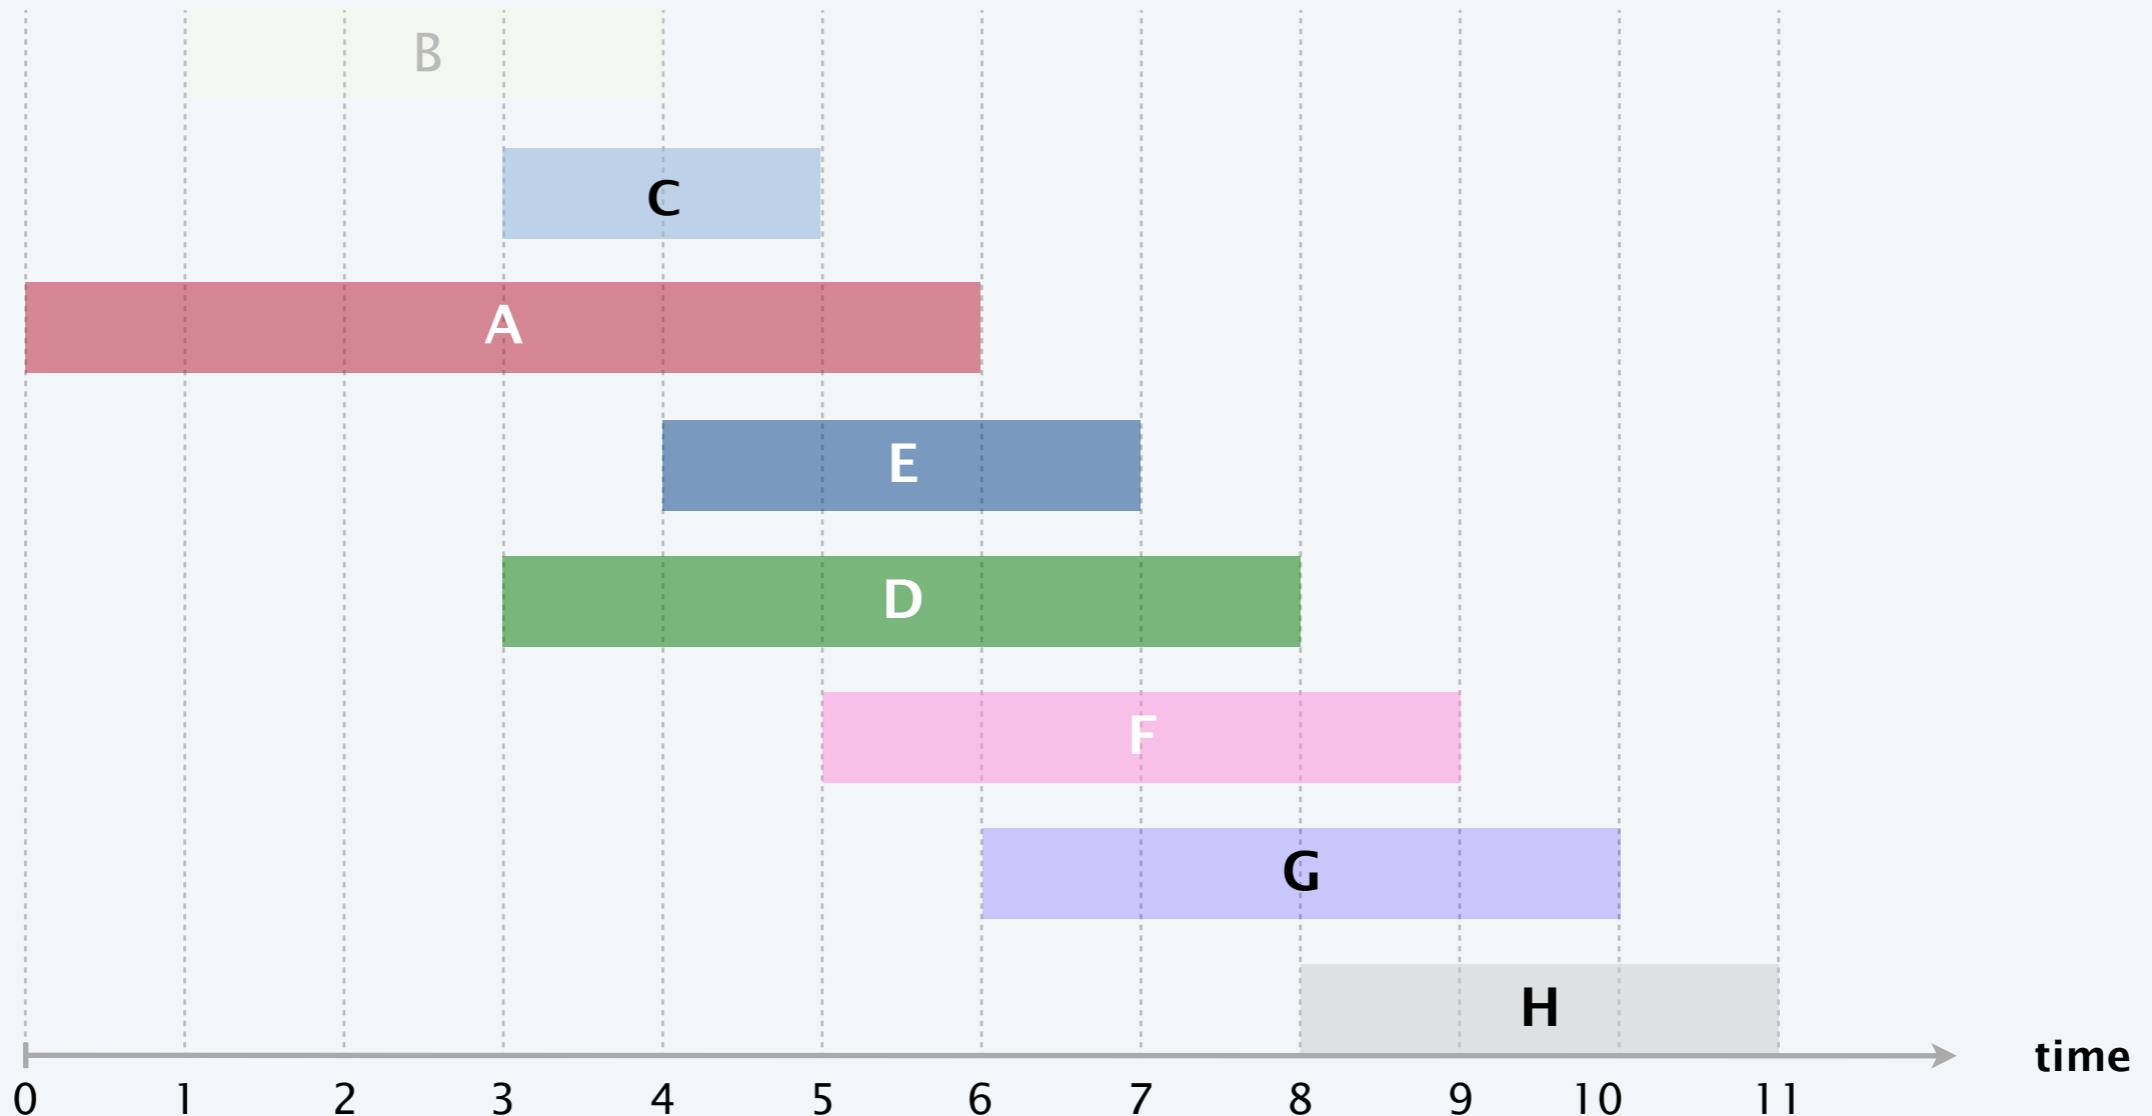
job B is compatible (add to schedule)



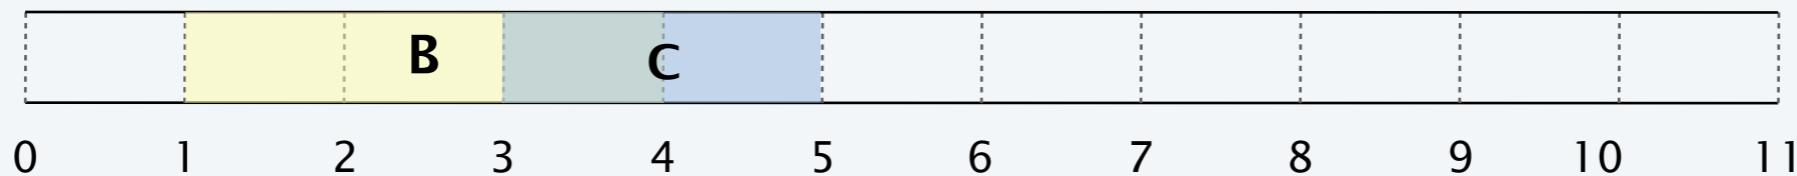
# Earliest-finish-time-first algorithm demo



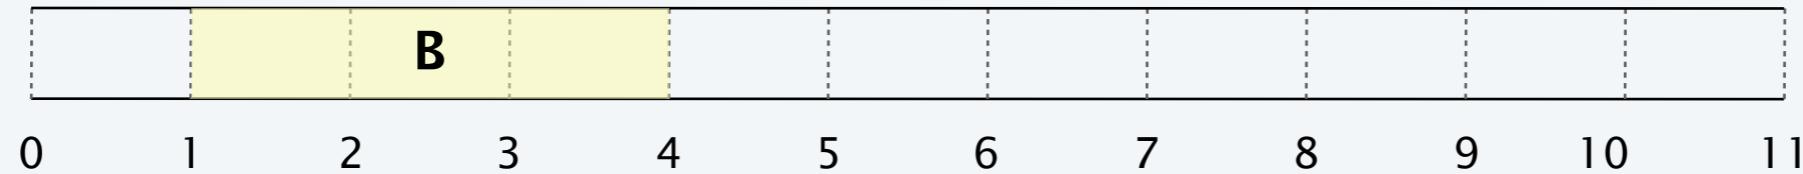
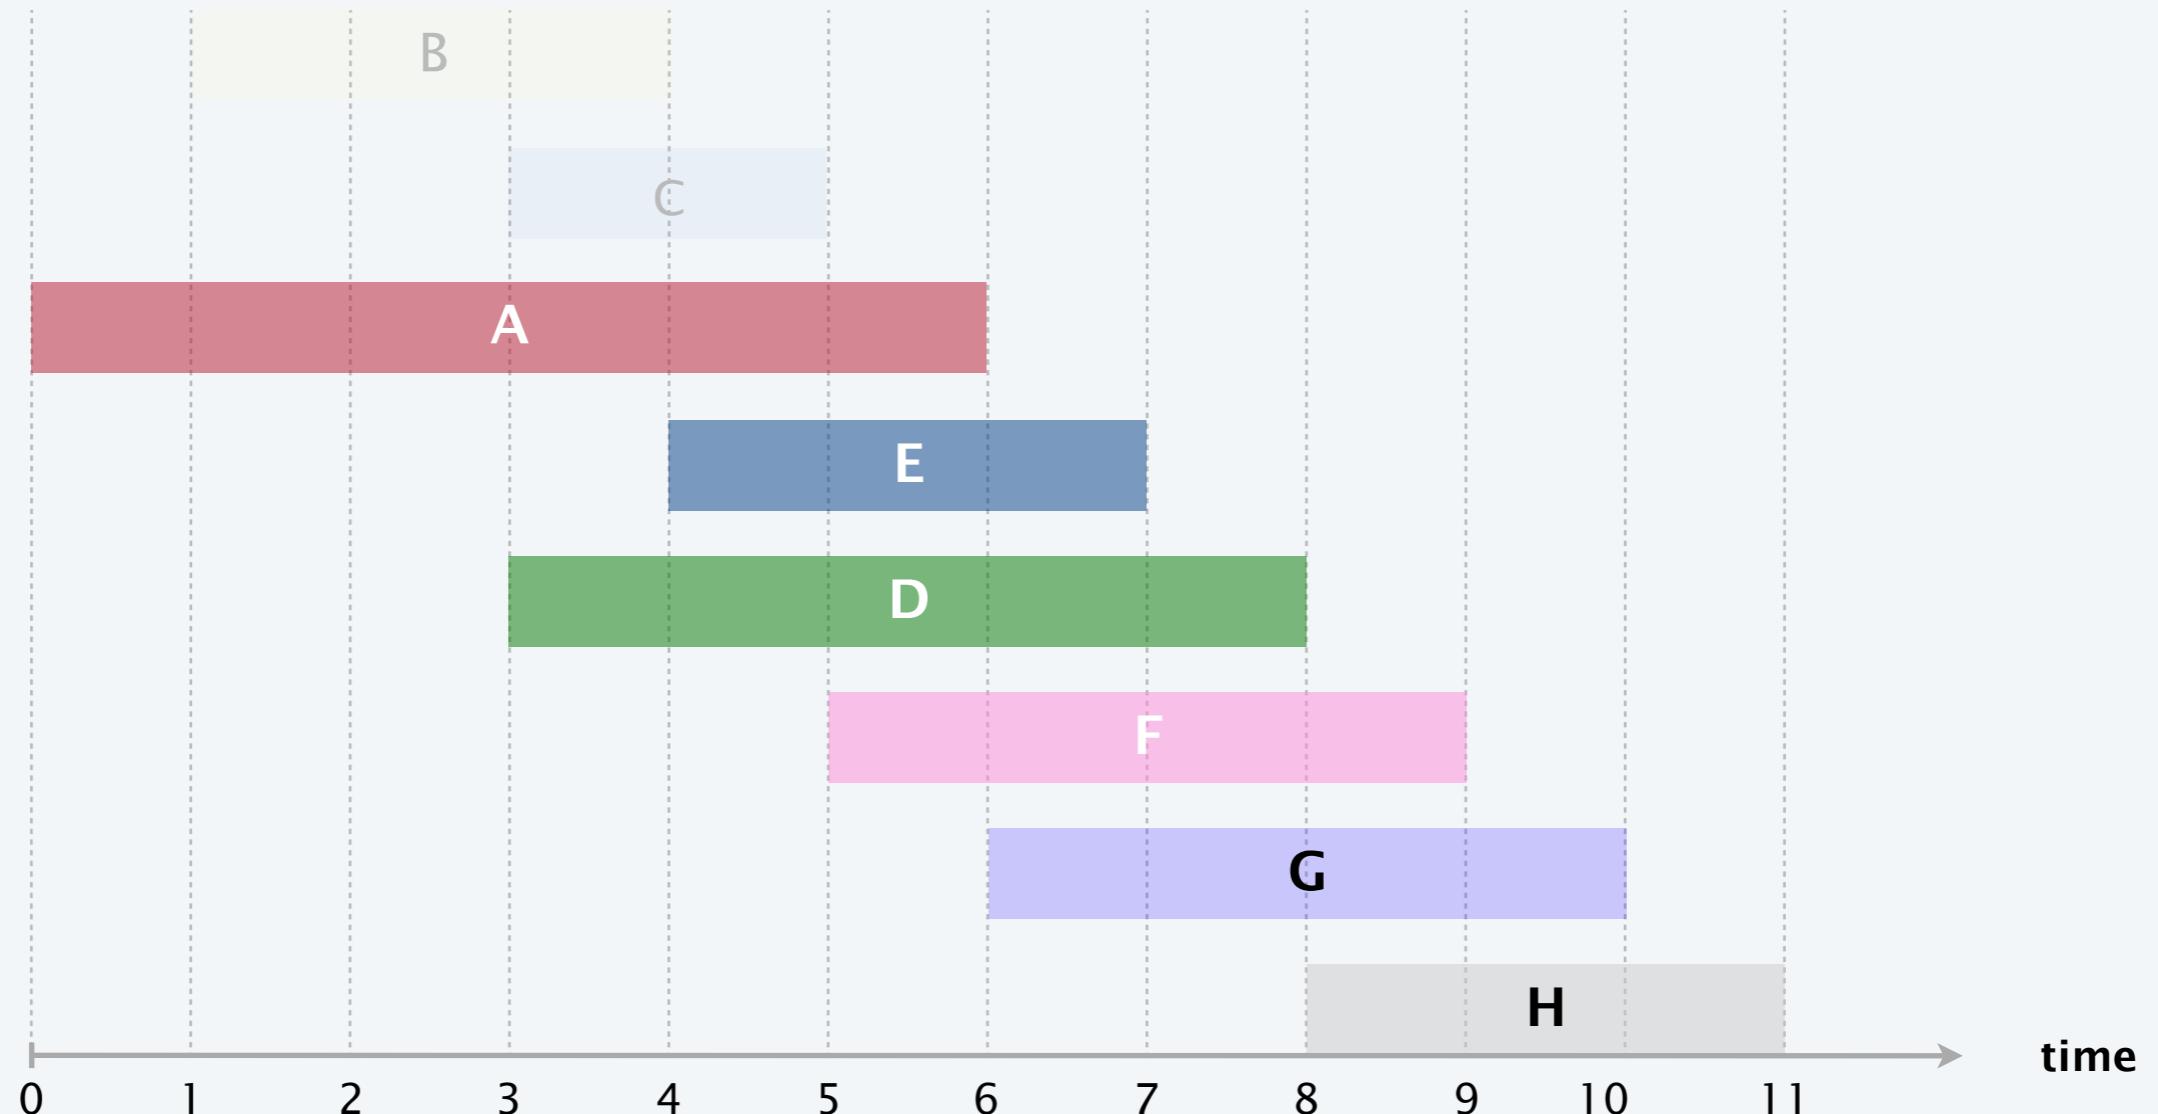
# Earliest-finish-time-first algorithm demo



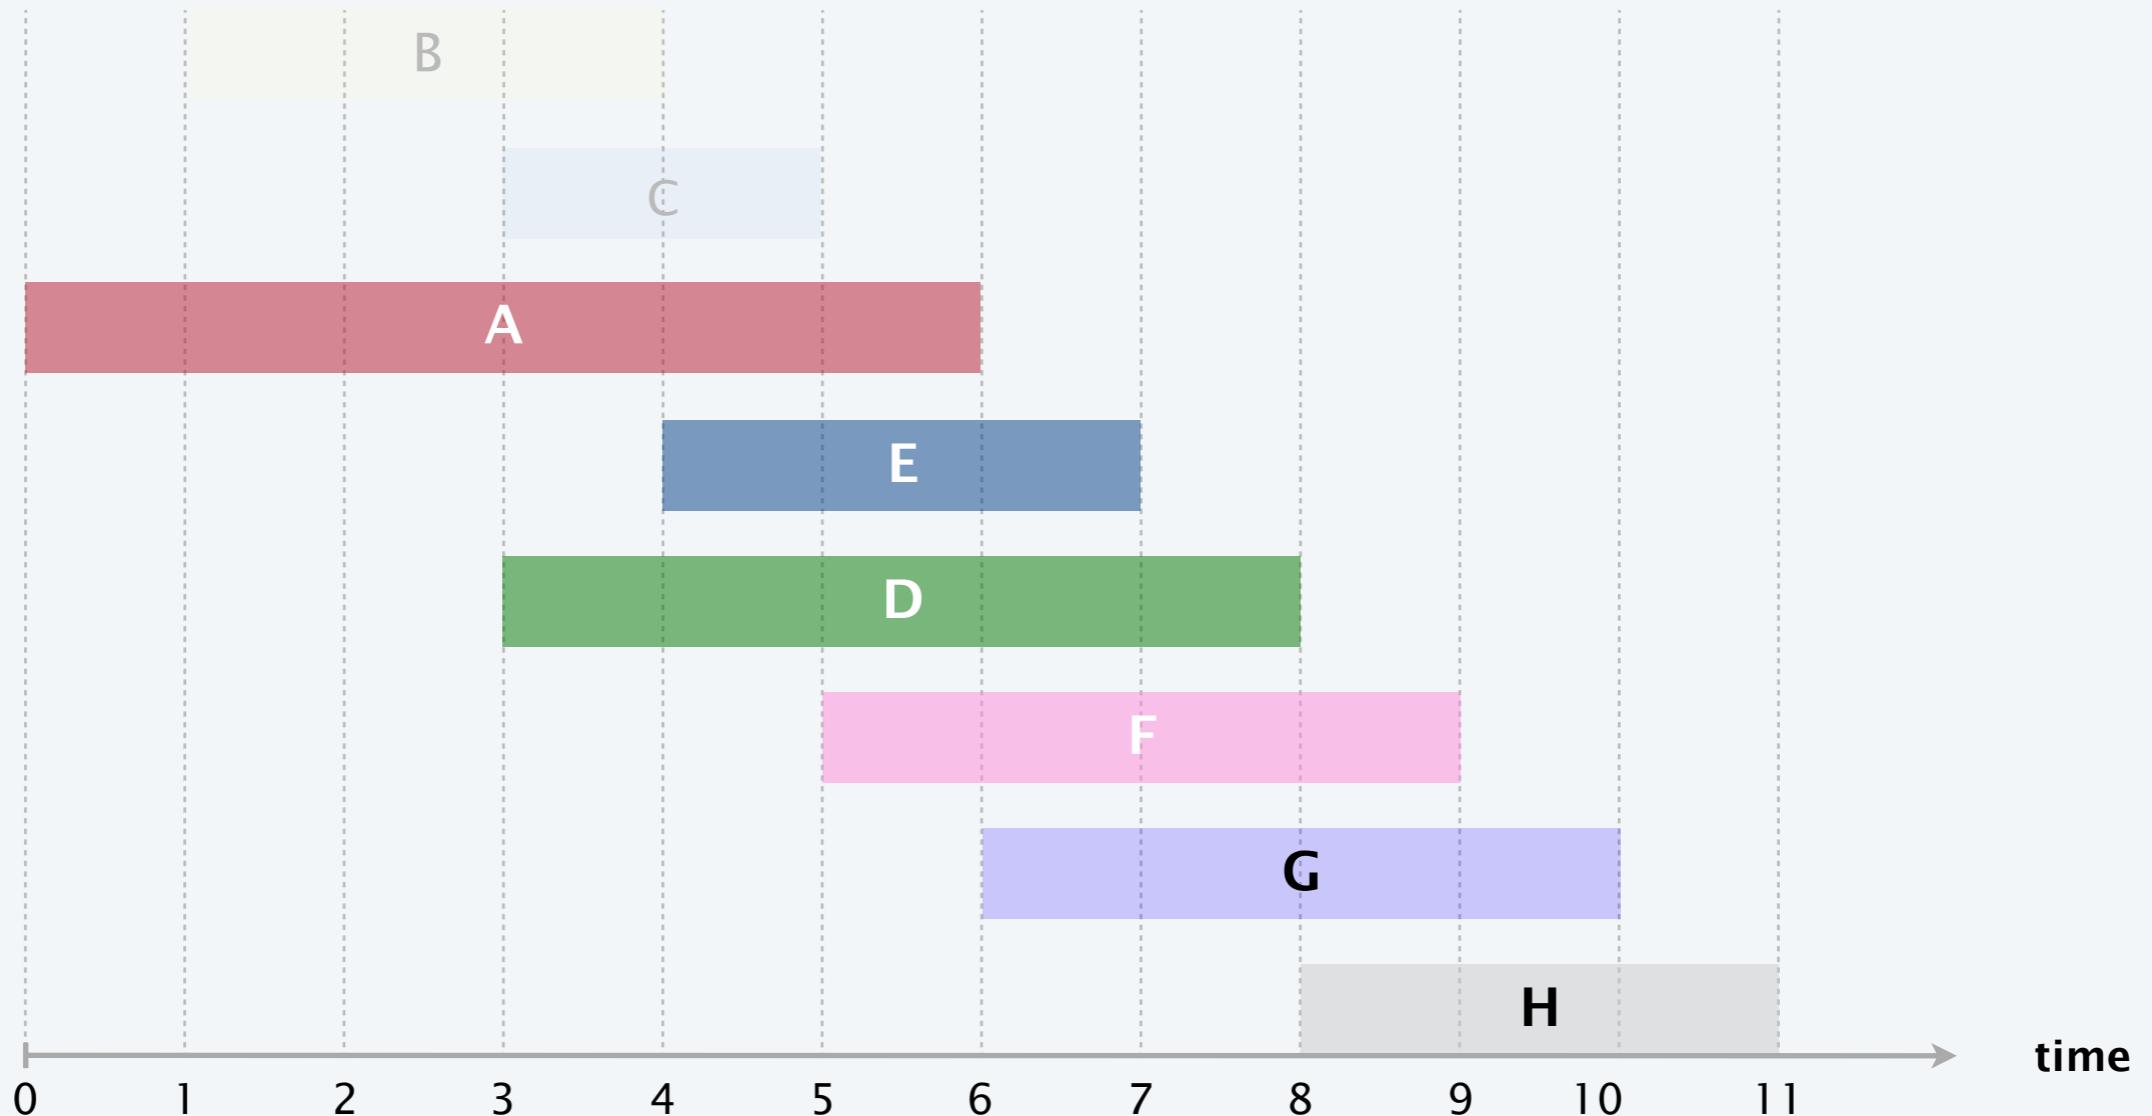
job C is incompatible (do not add to schedule)



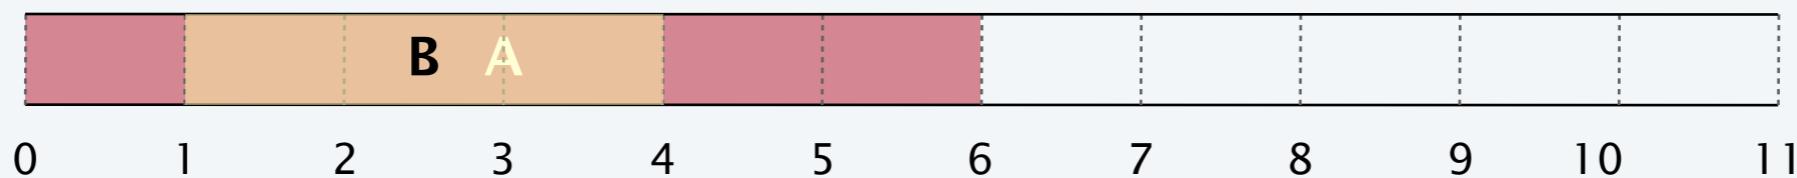
# Earliest-finish-time-first algorithm demo



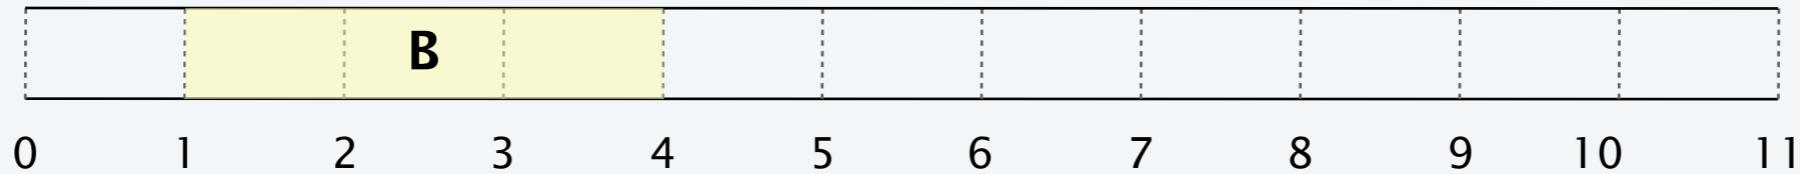
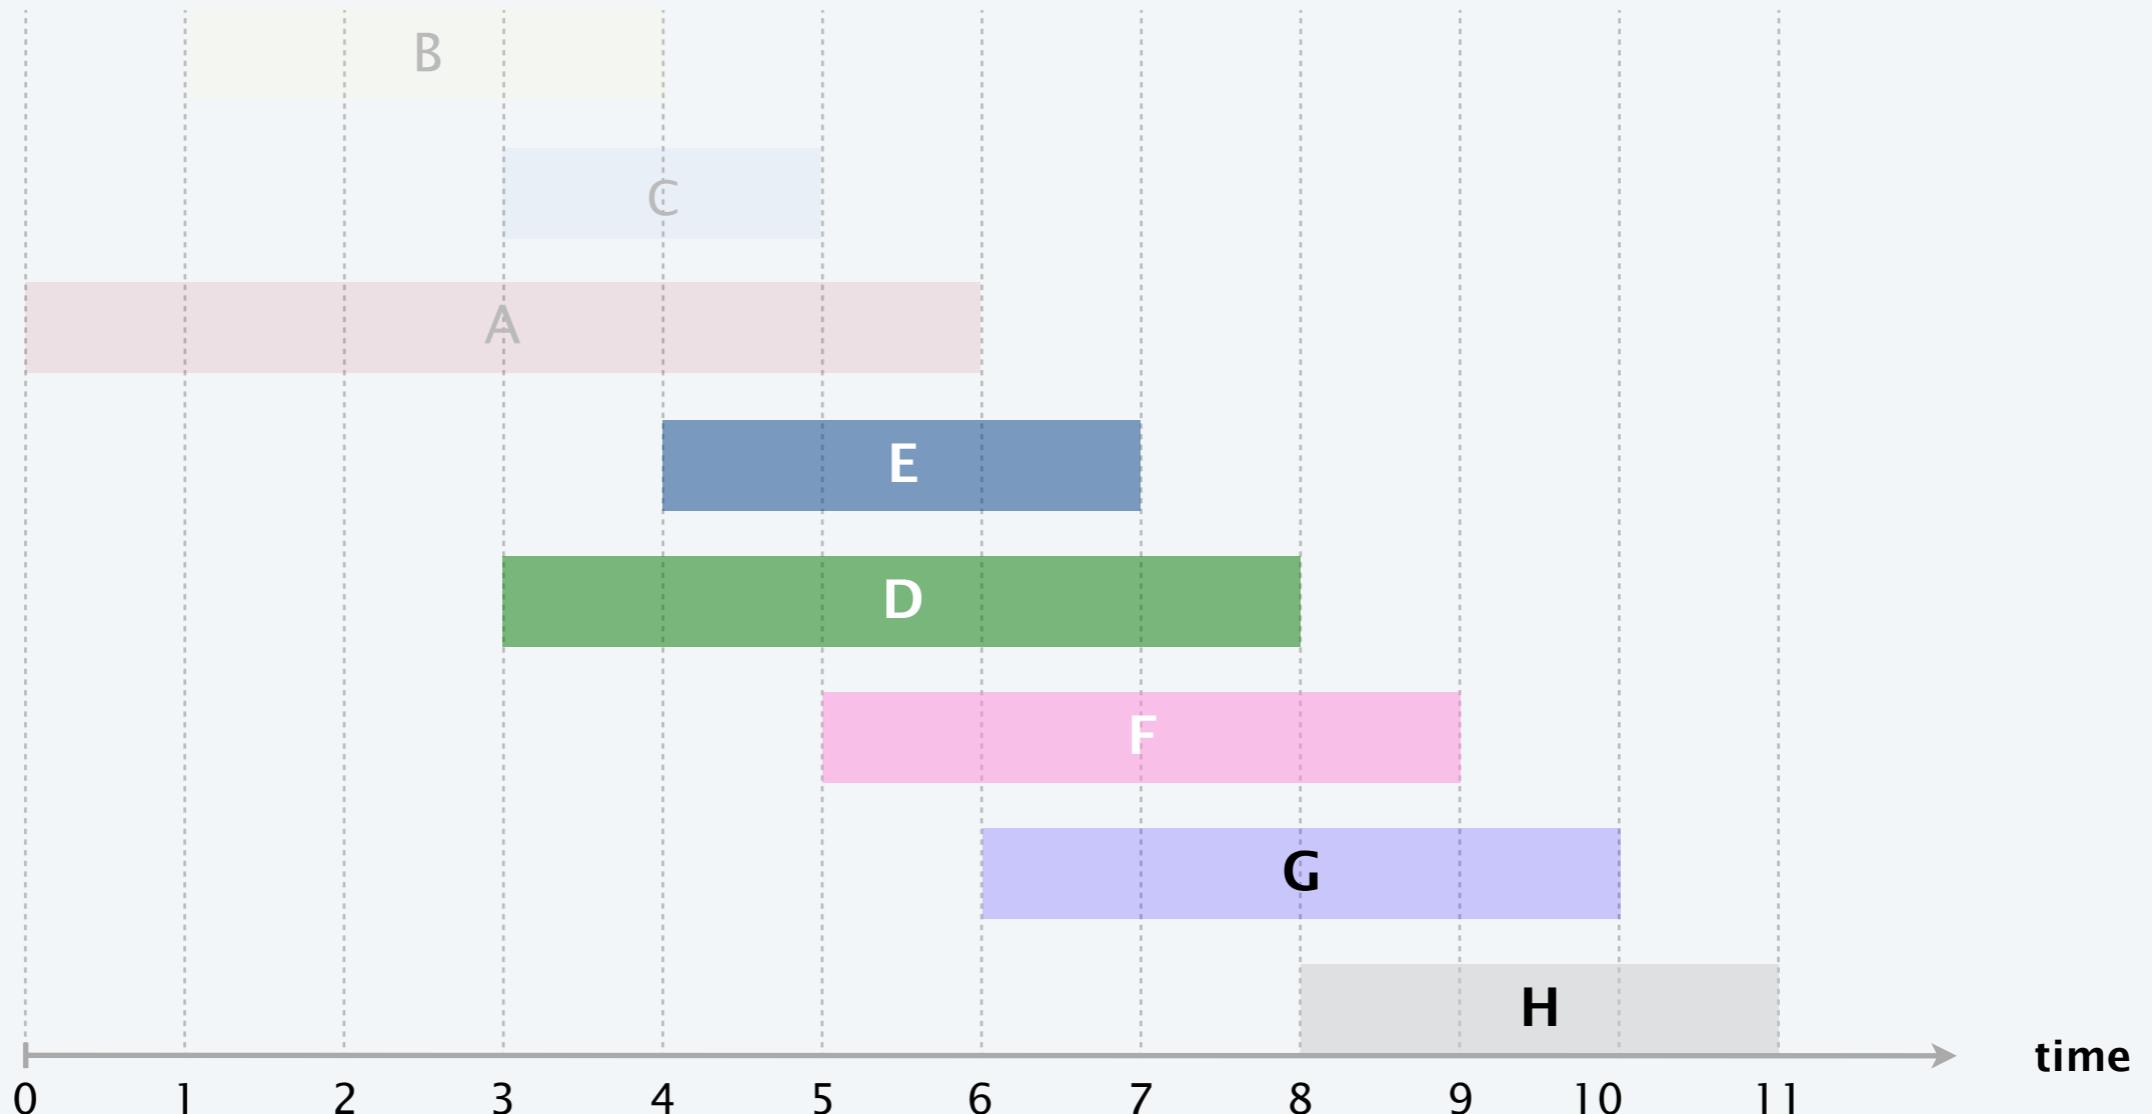
# Earliest-finish-time-first algorithm demo



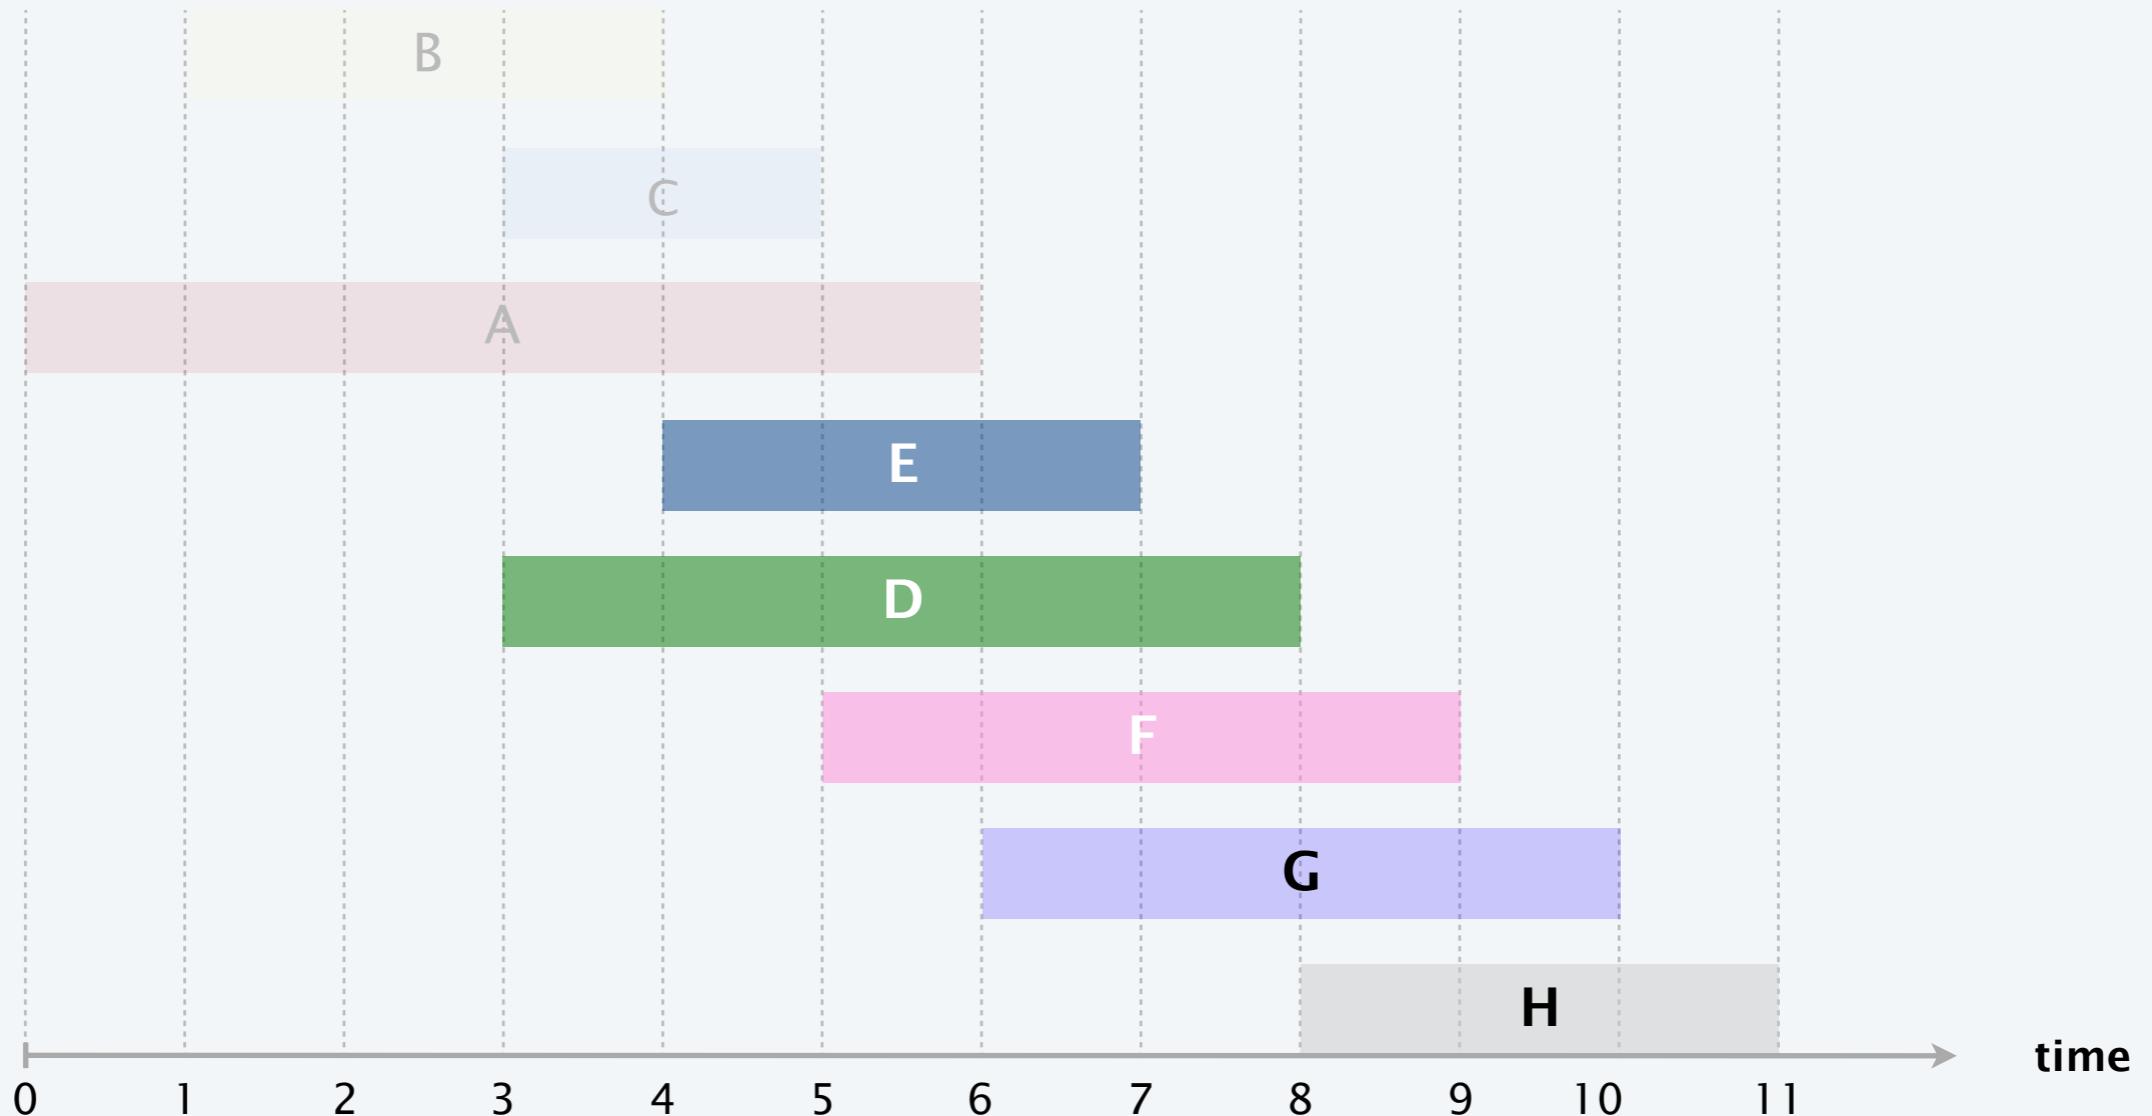
job A is incompatible (do not add to schedule)



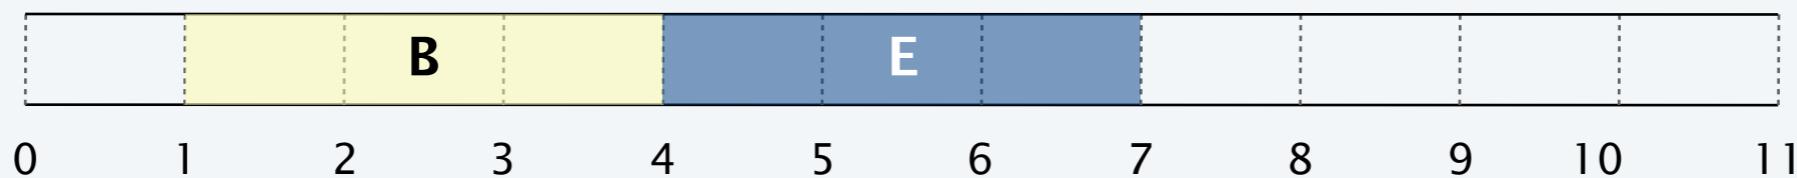
# Earliest-finish-time-first algorithm demo



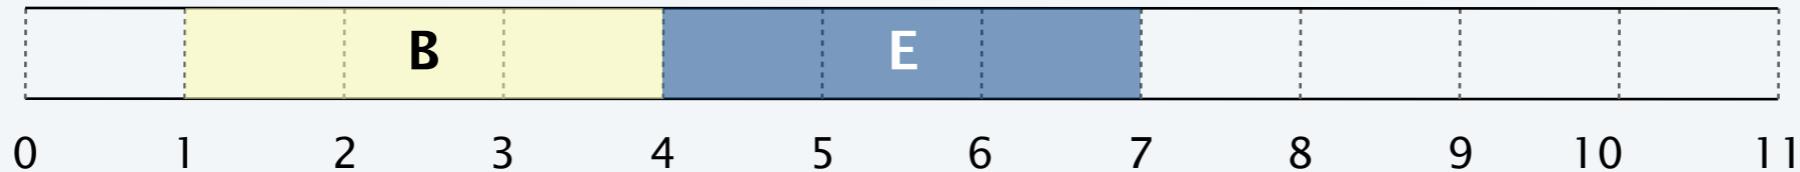
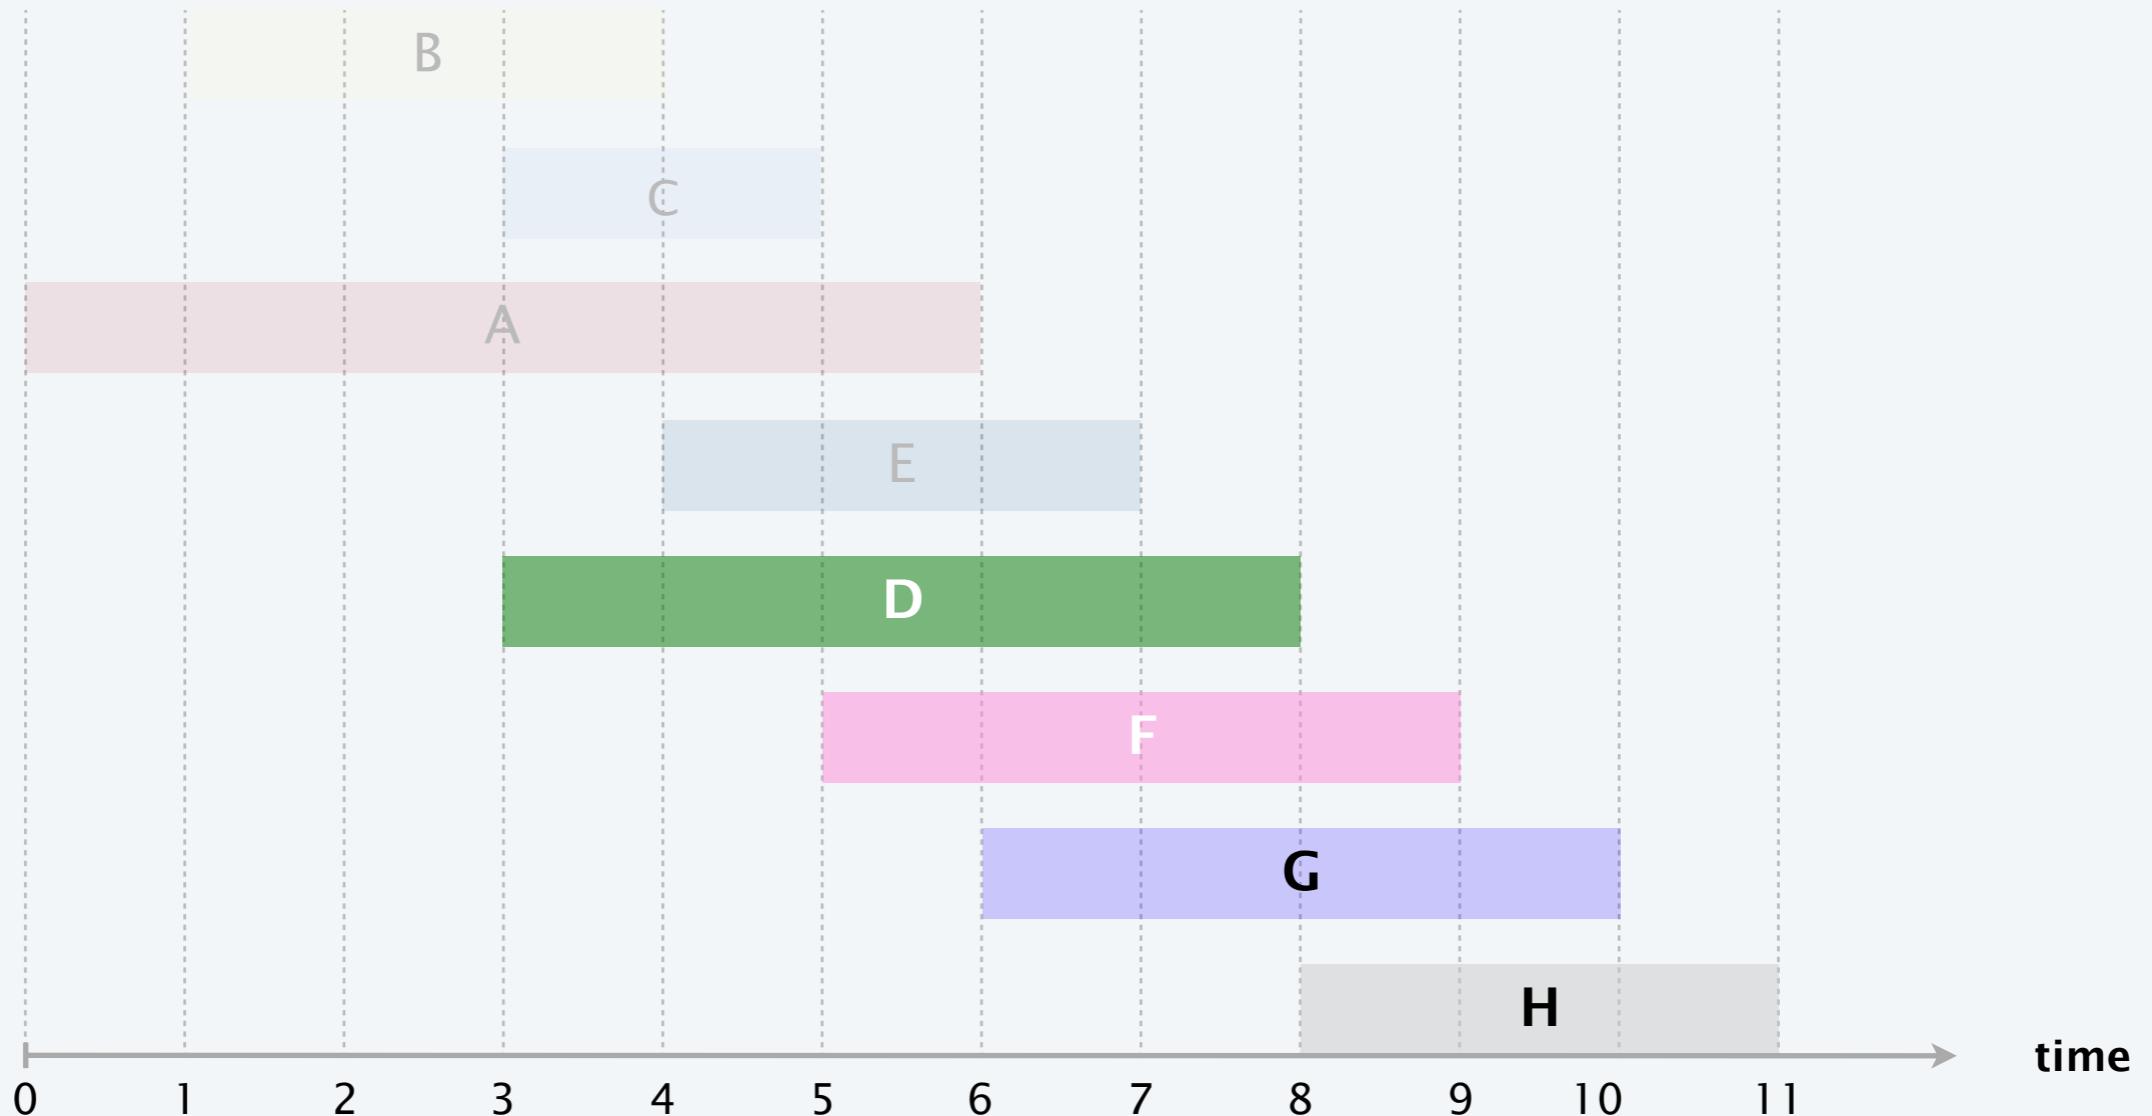
# Earliest-finish-time-first algorithm demo



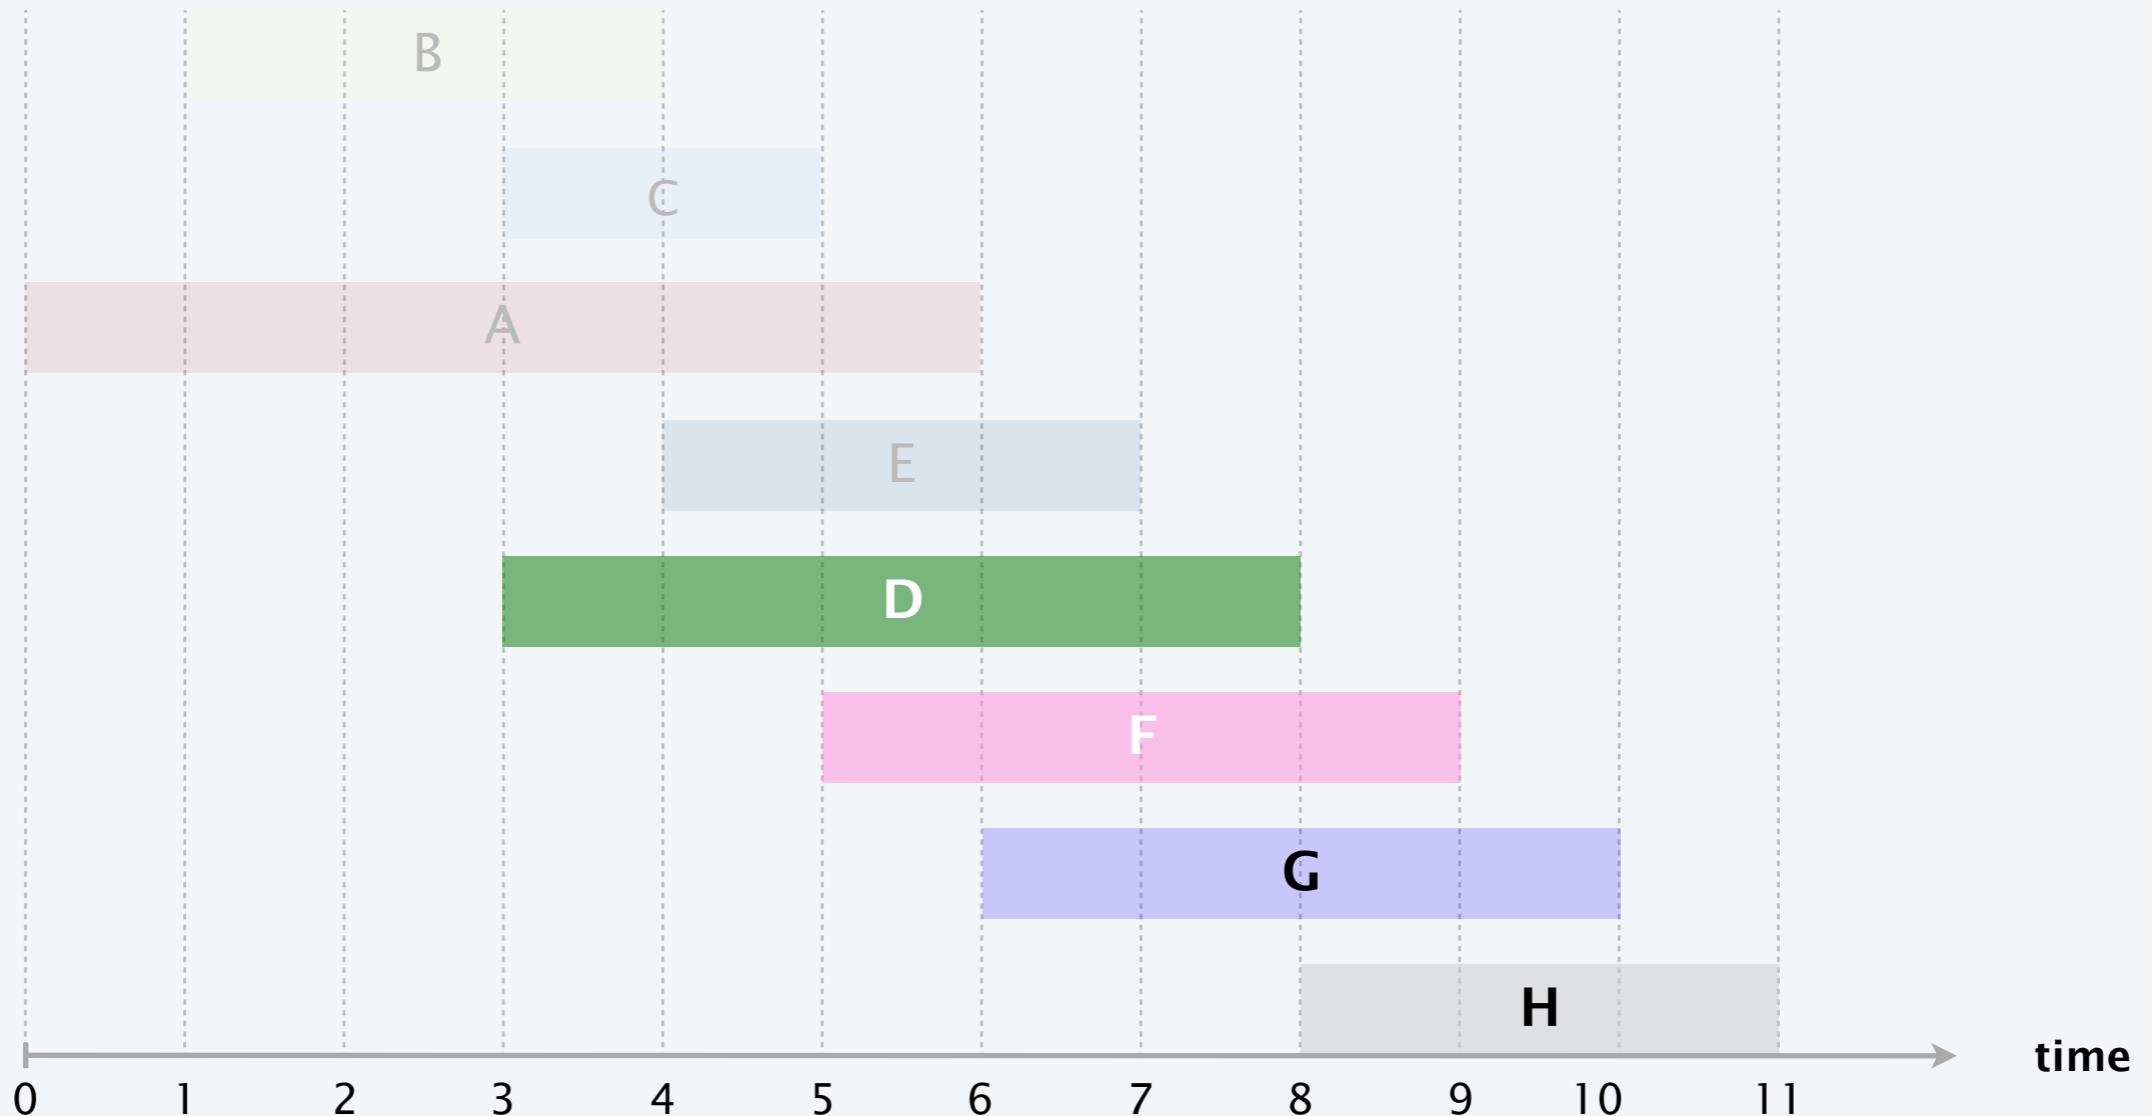
job E is compatible (add to schedule)



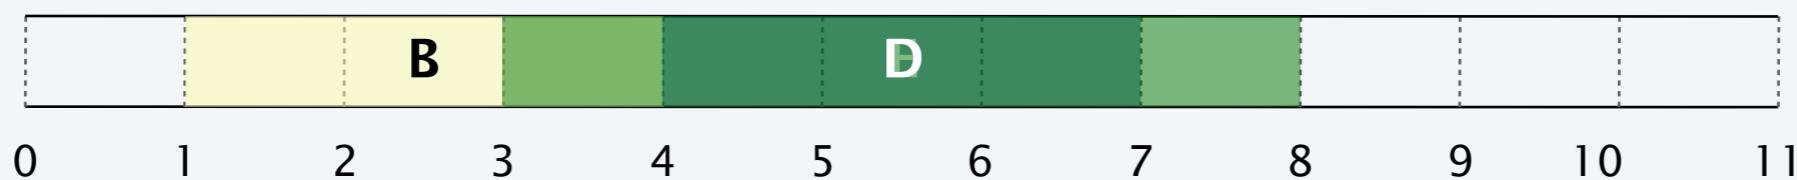
# Earliest-finish-time-first algorithm demo



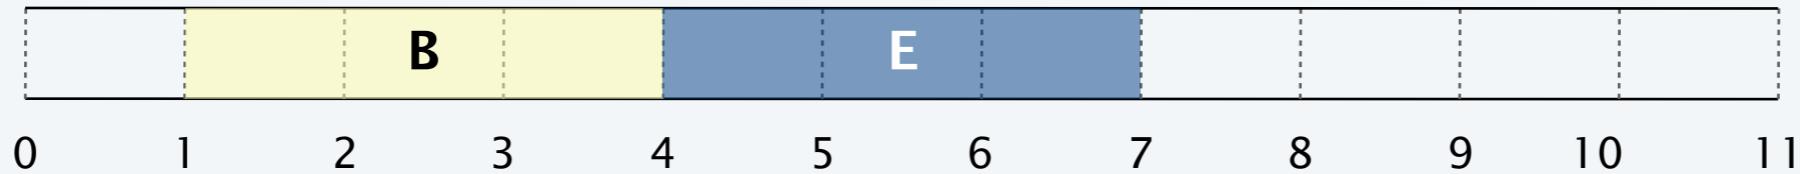
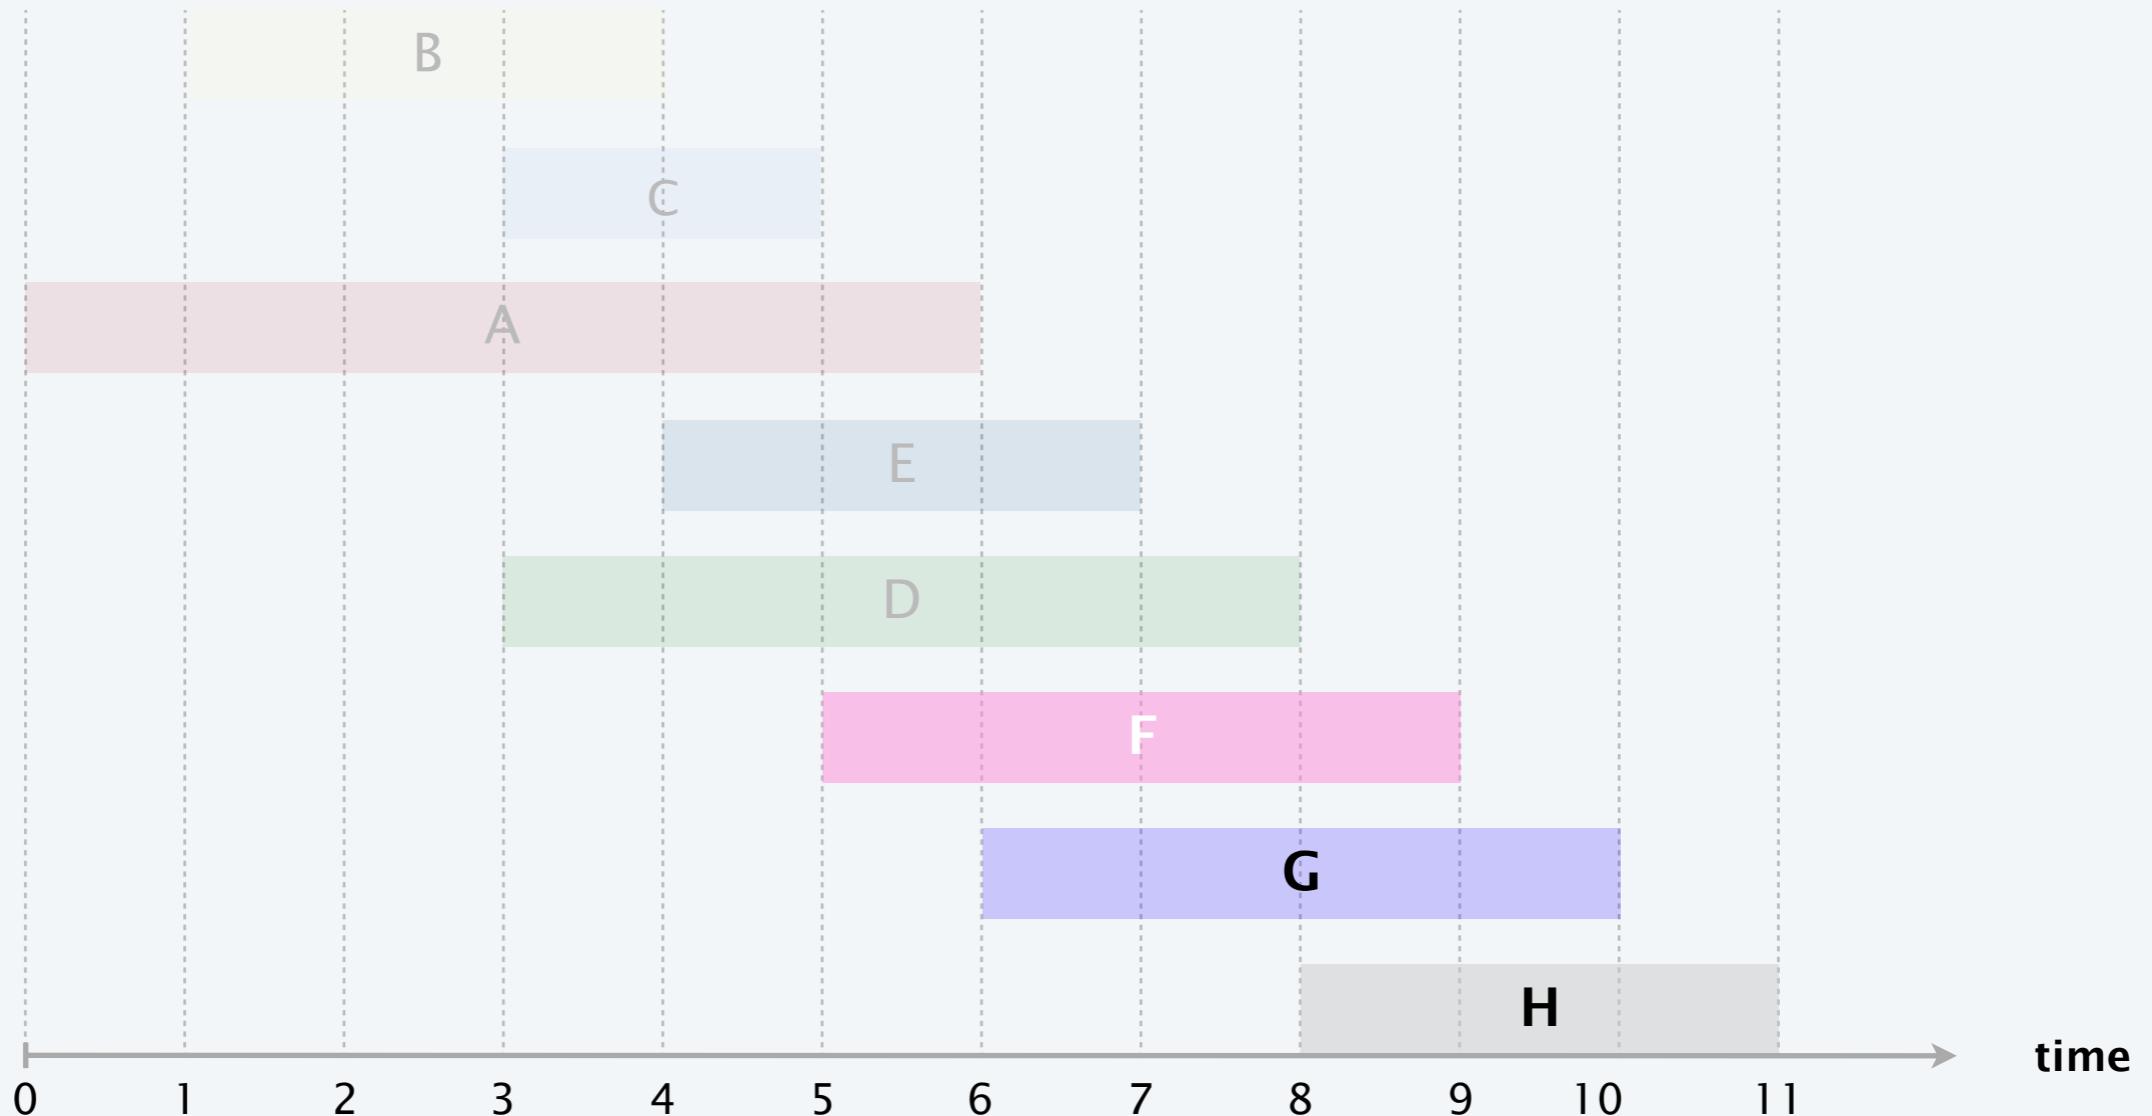
# Earliest-finish-time-first algorithm demo



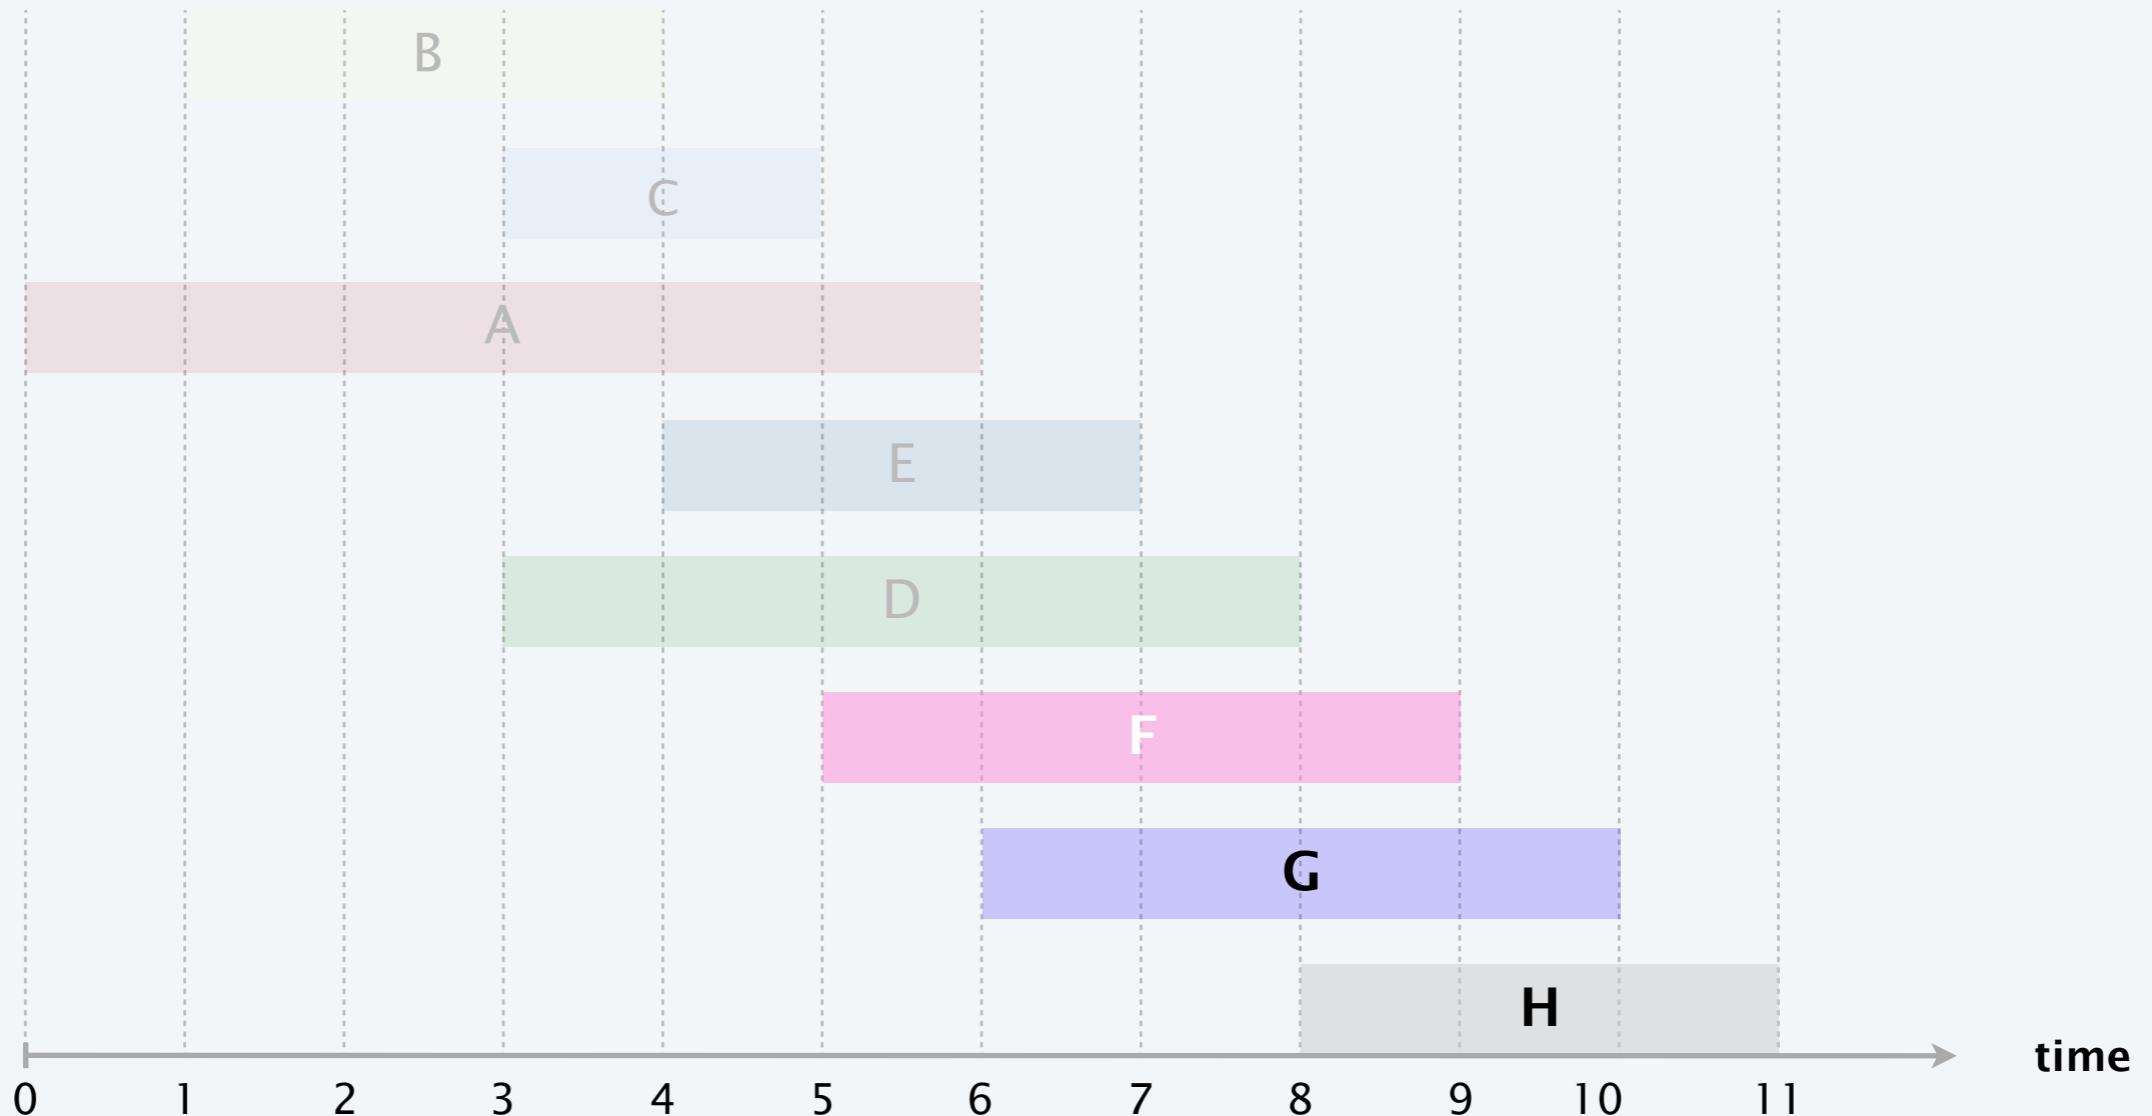
job D is incompatible (do not add to schedule)



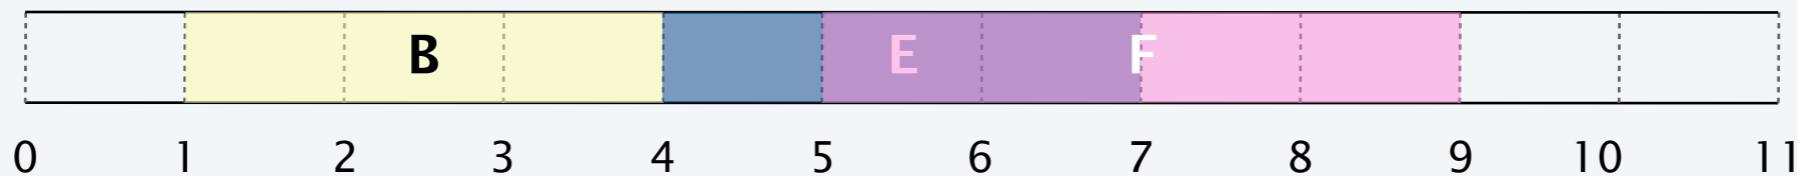
# Earliest-finish-time-first algorithm demo



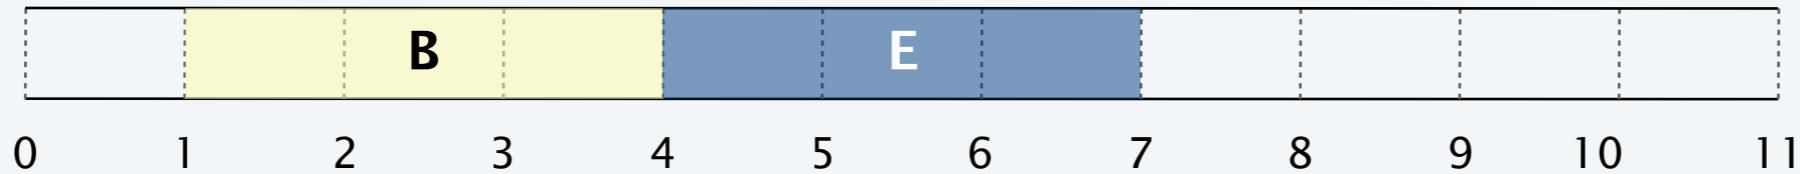
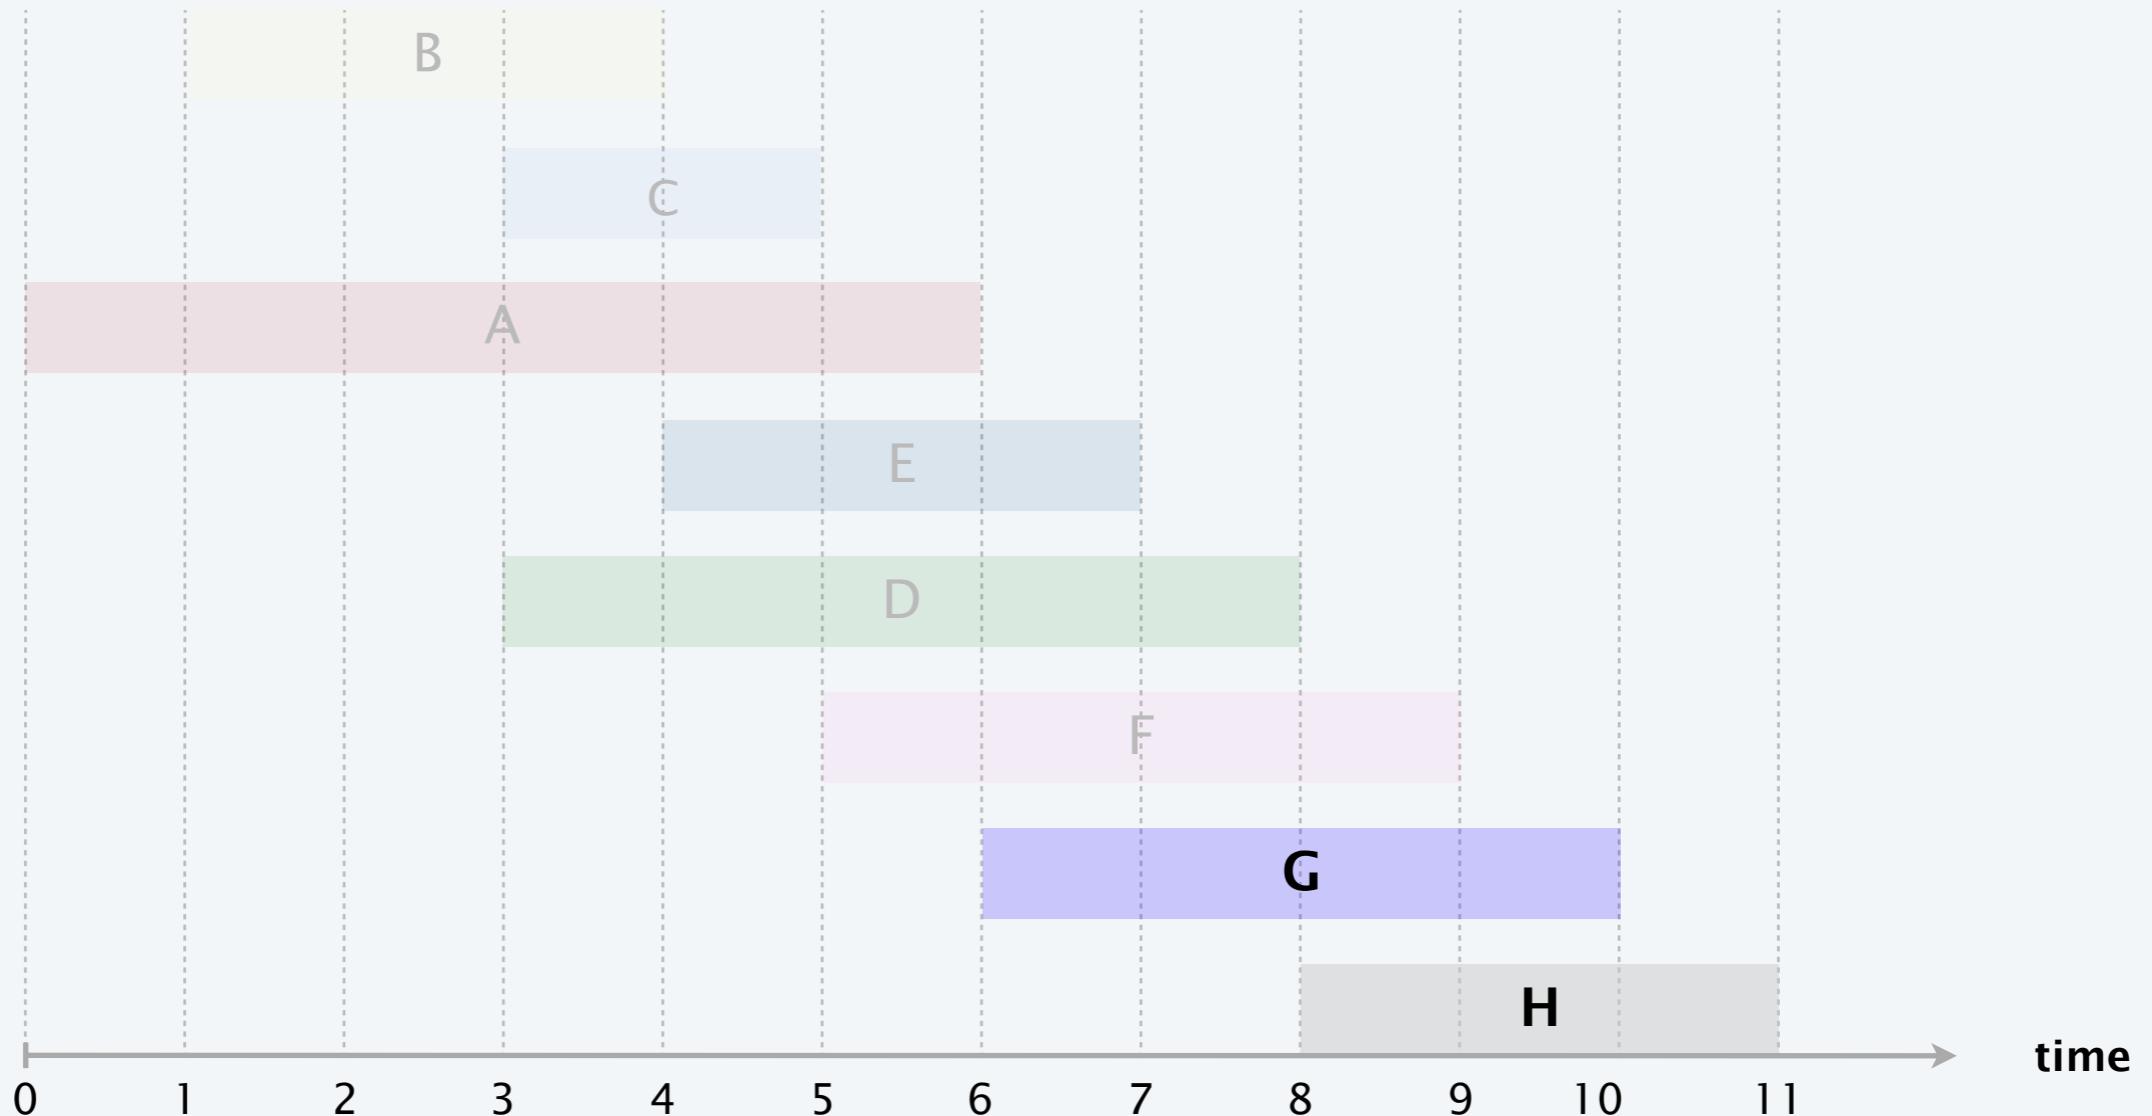
# Earliest-finish-time-first algorithm demo



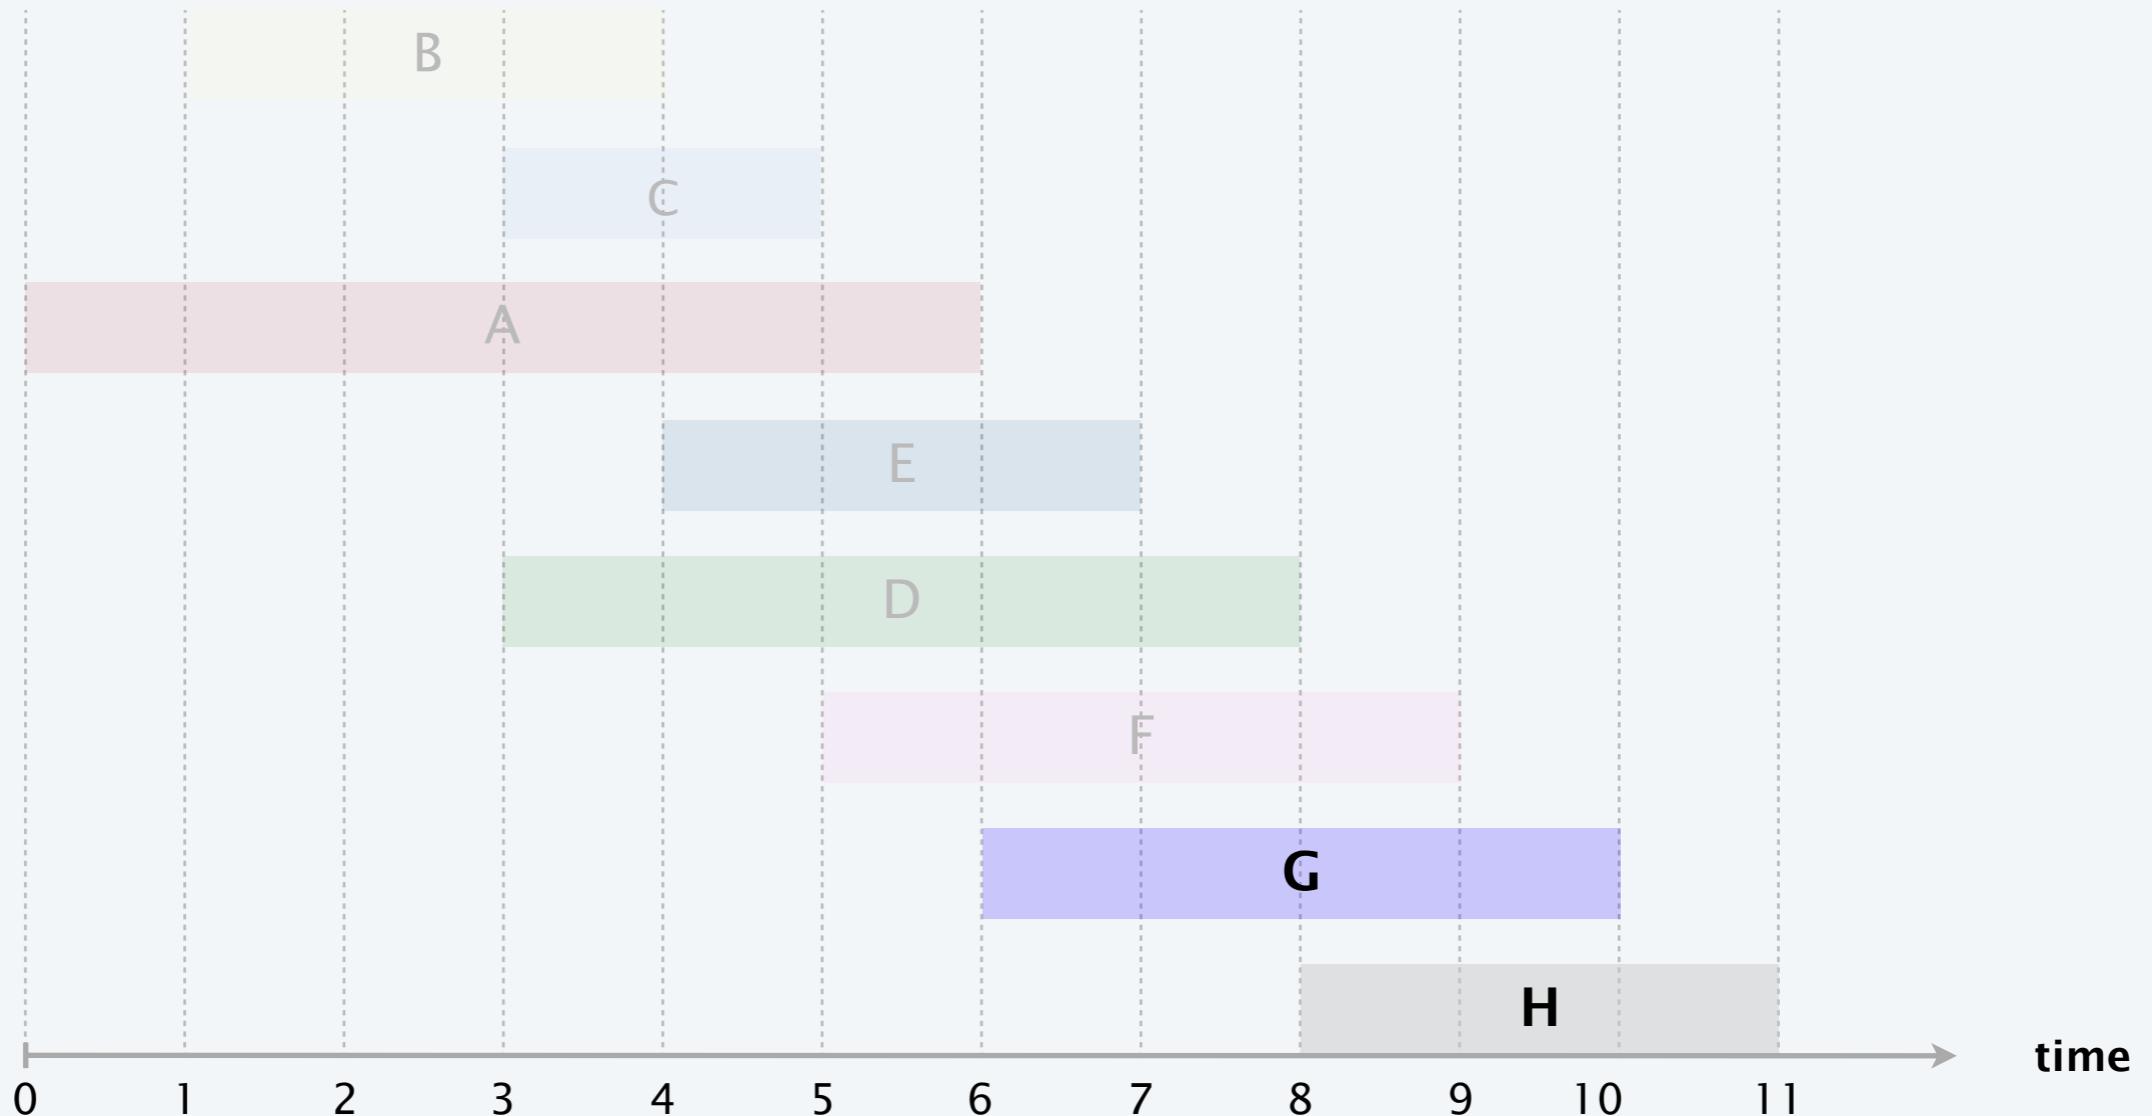
**job F is incompatible (do not add to schedule)**



# Earliest-finish-time-first algorithm demo



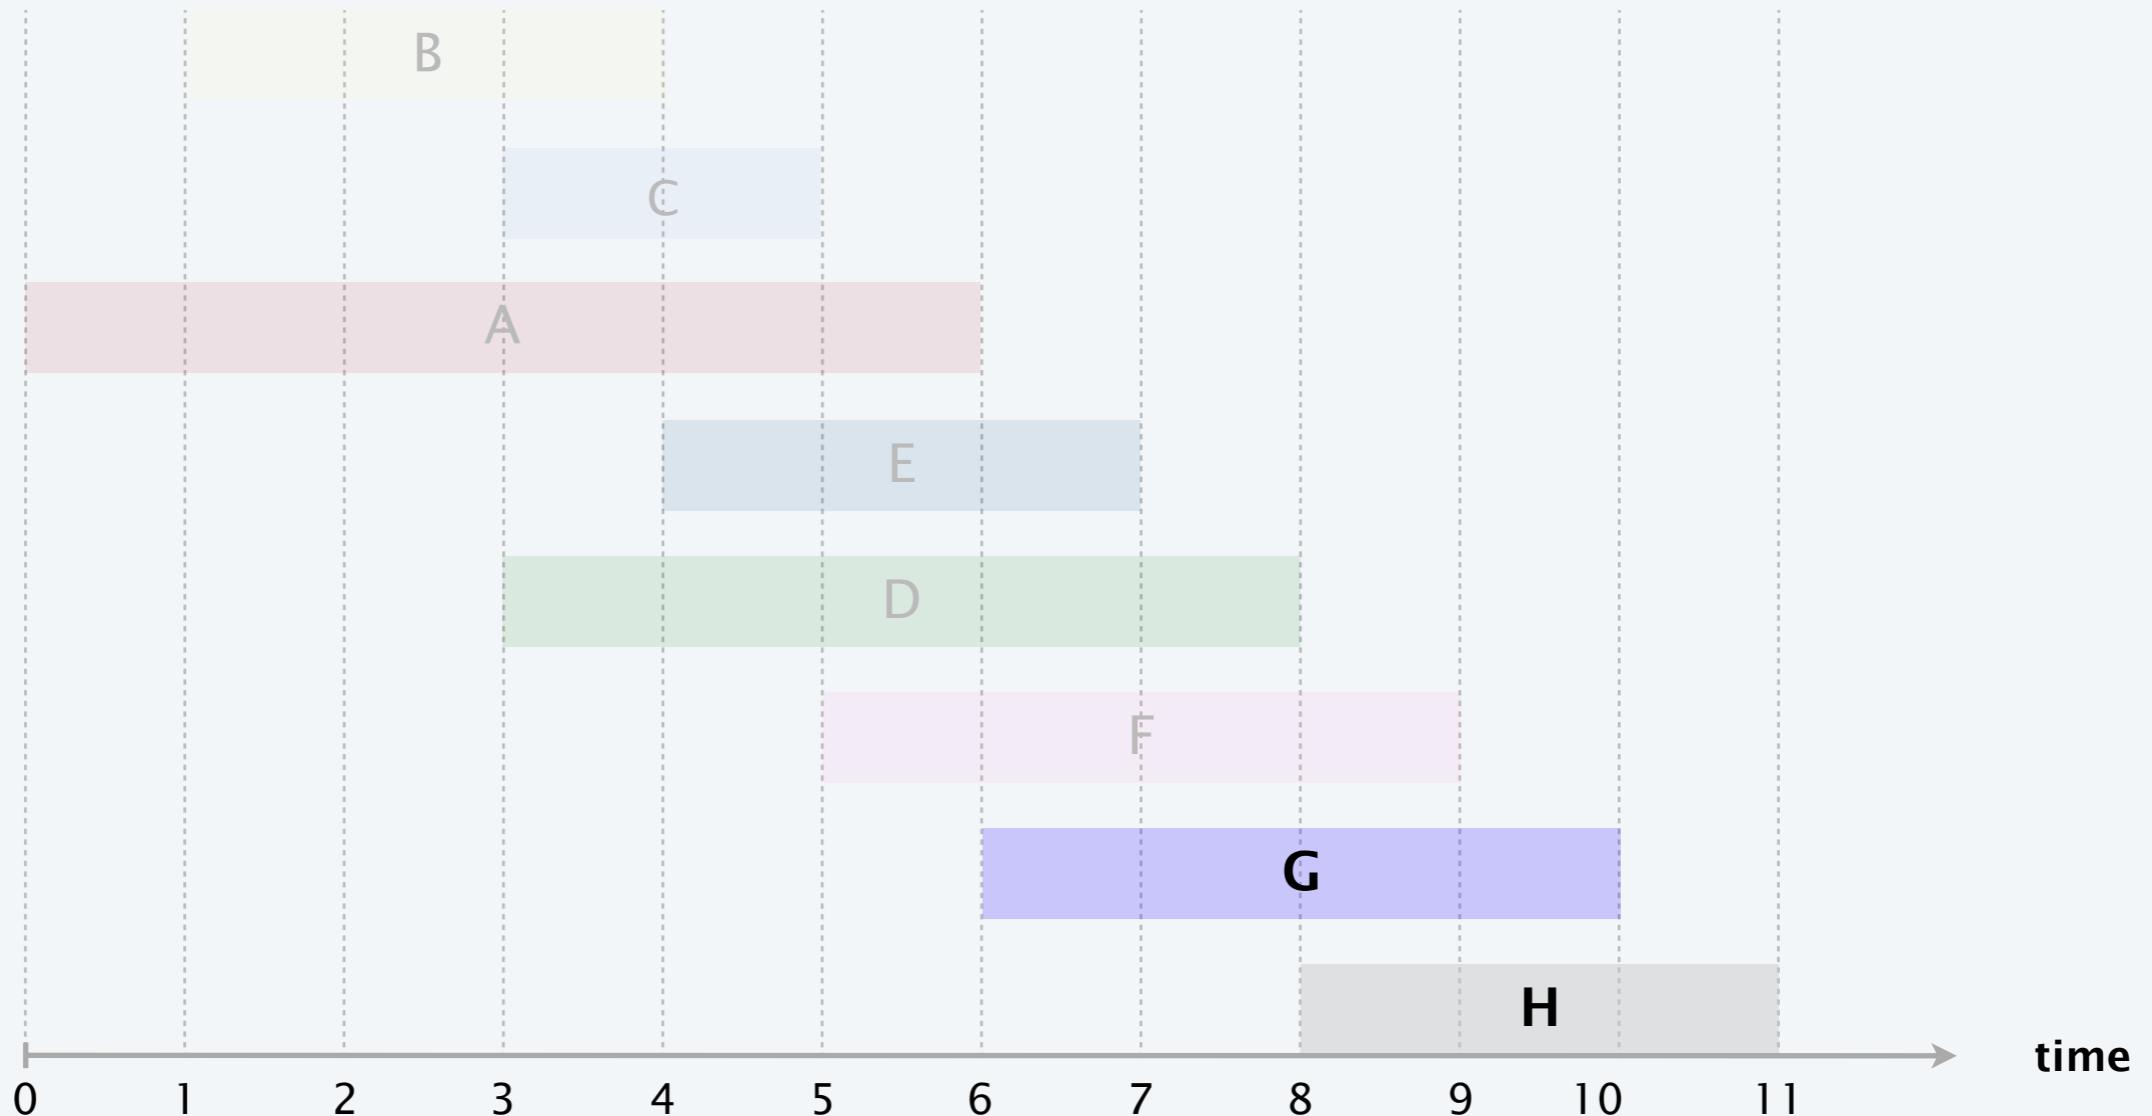
# Earliest-finish-time-first algorithm demo



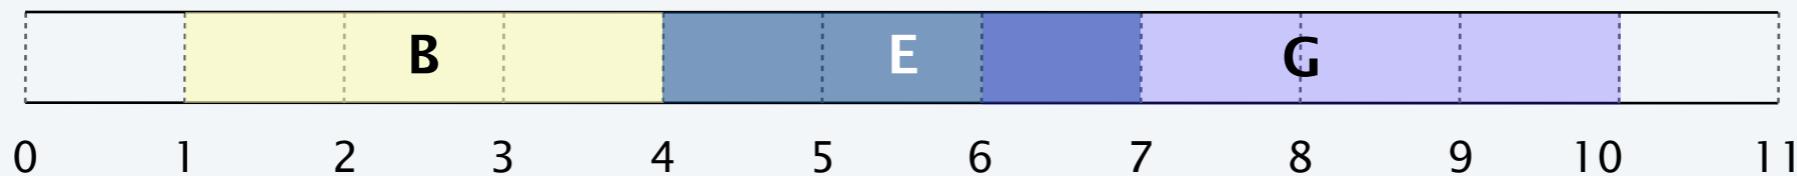
job G is incompatible (do not add to schedule)



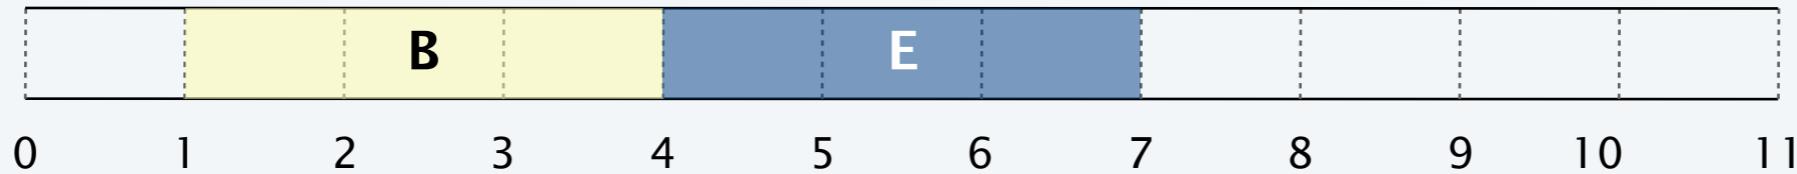
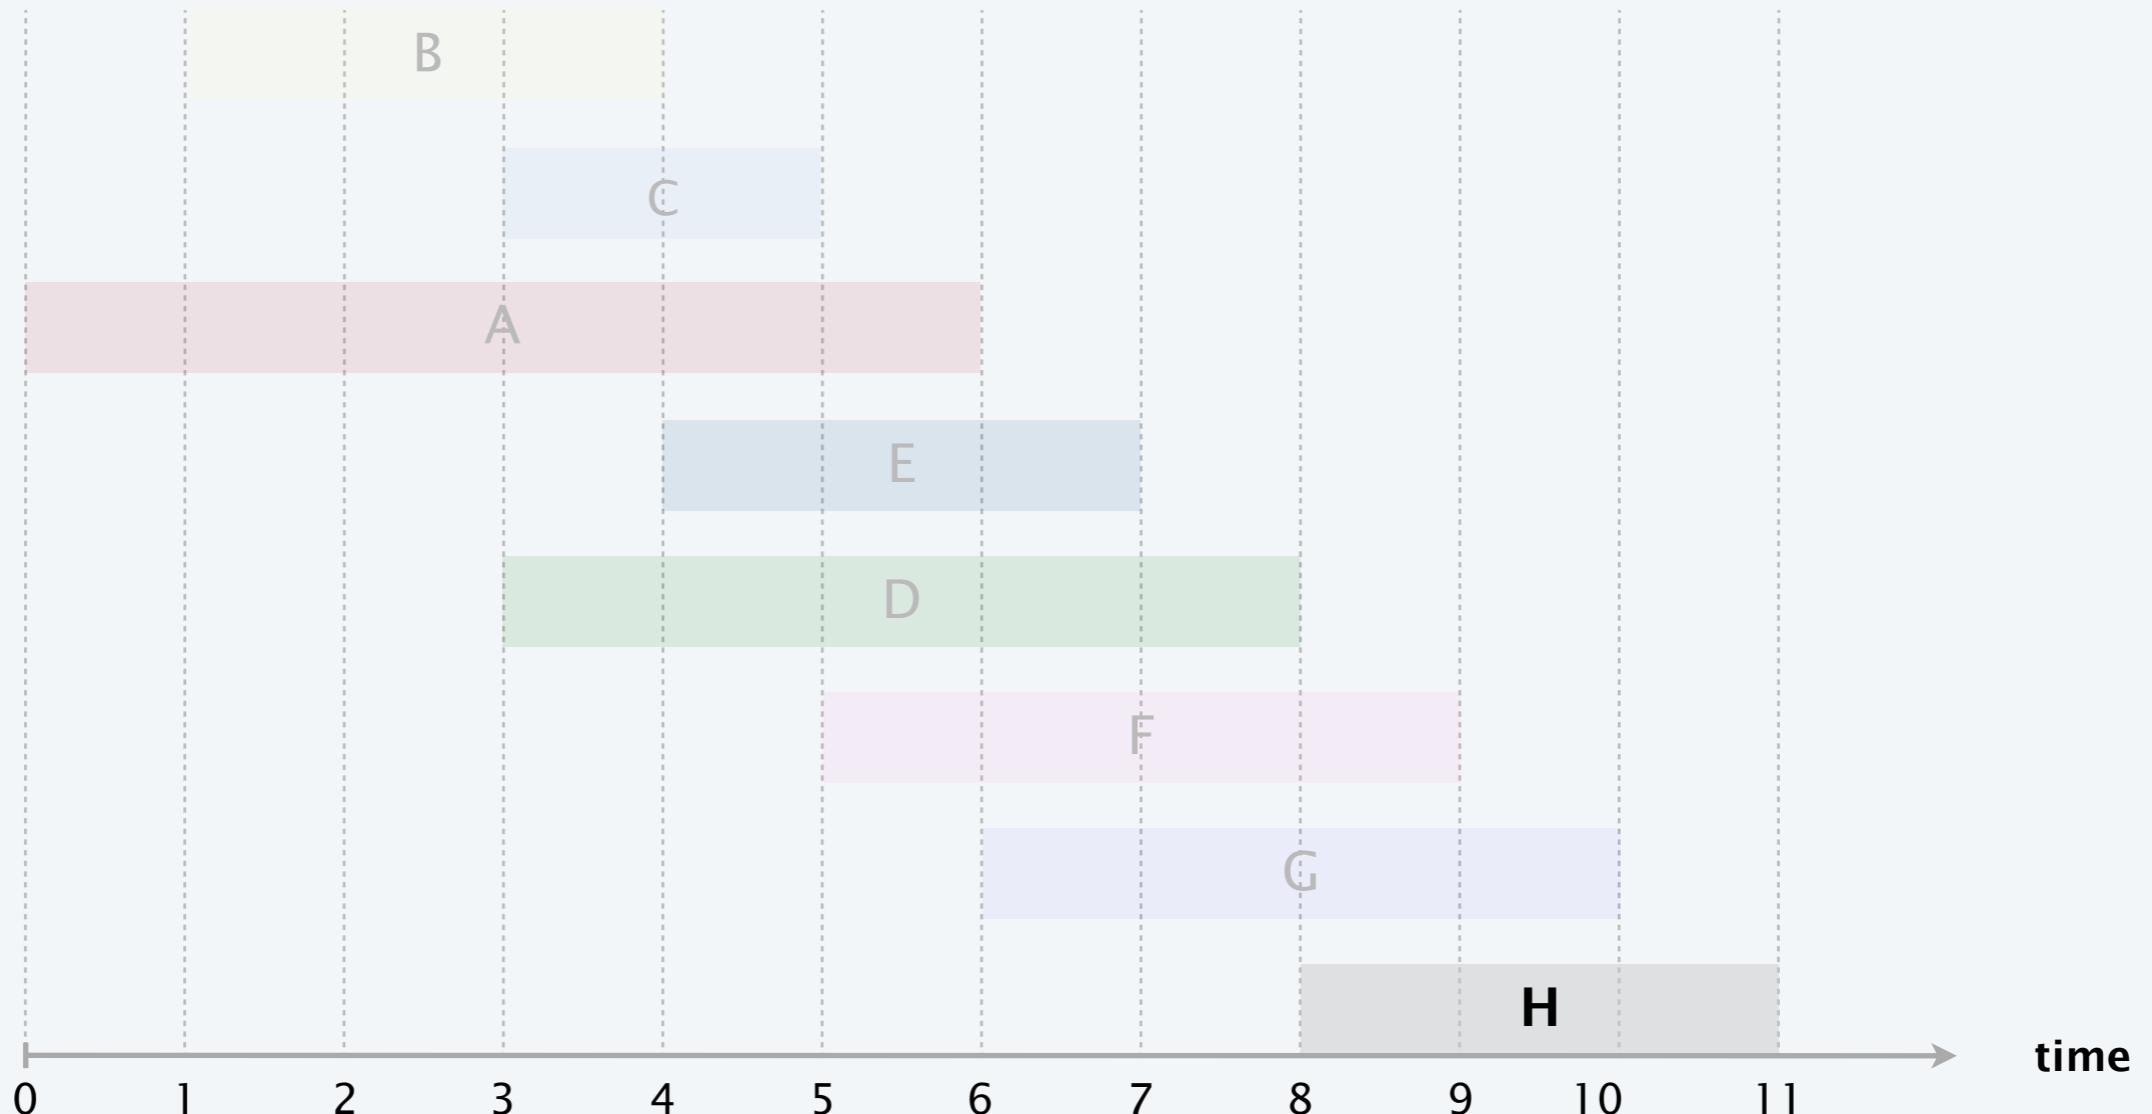
# Earliest-finish-time-first algorithm demo



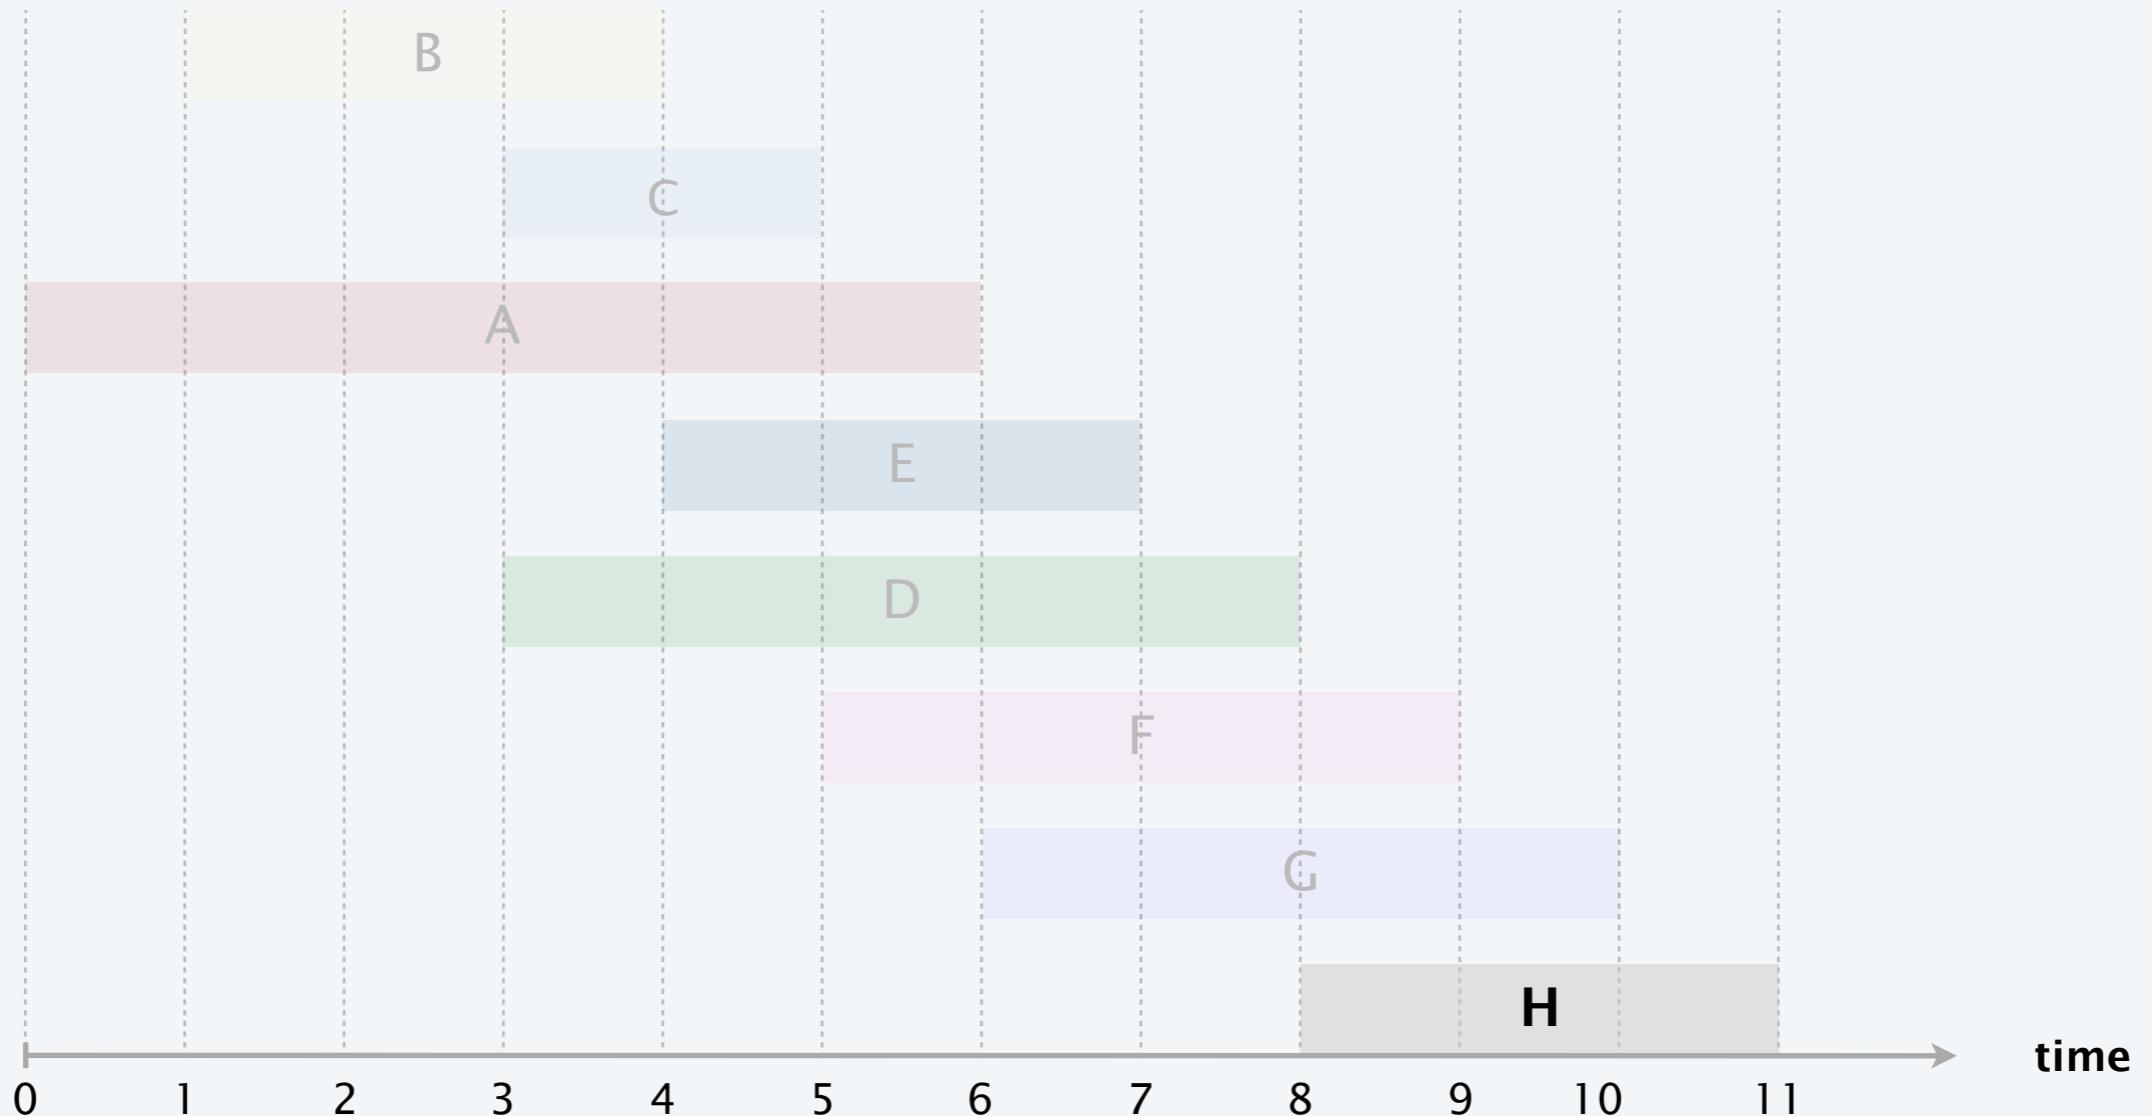
job G is incompatible (do not add to schedule)



# Earliest-finish-time-first algorithm demo



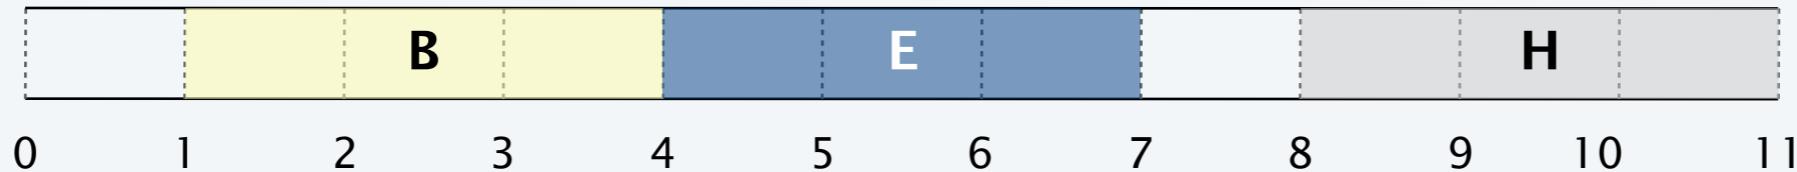
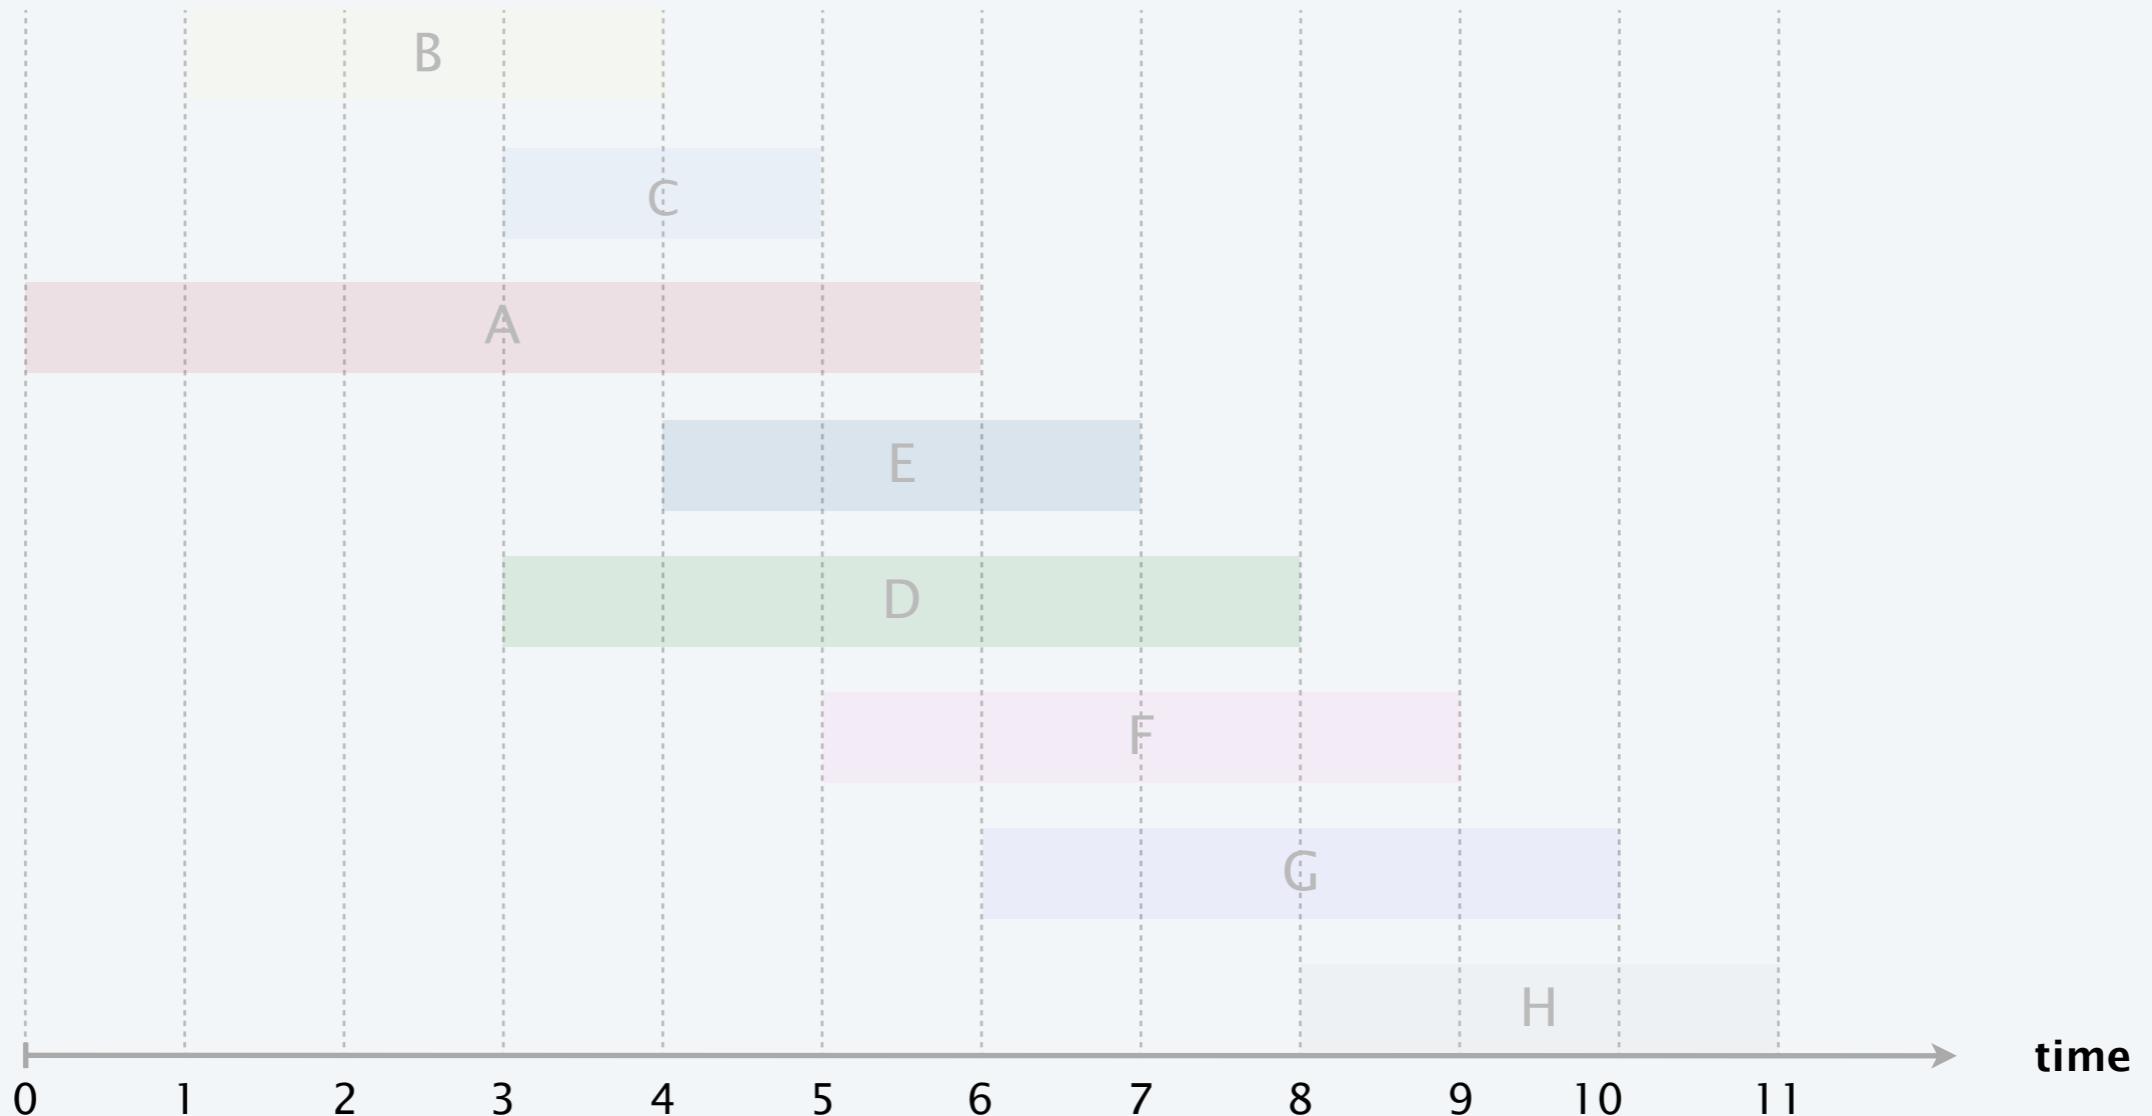
# Earliest-finish-time-first algorithm demo



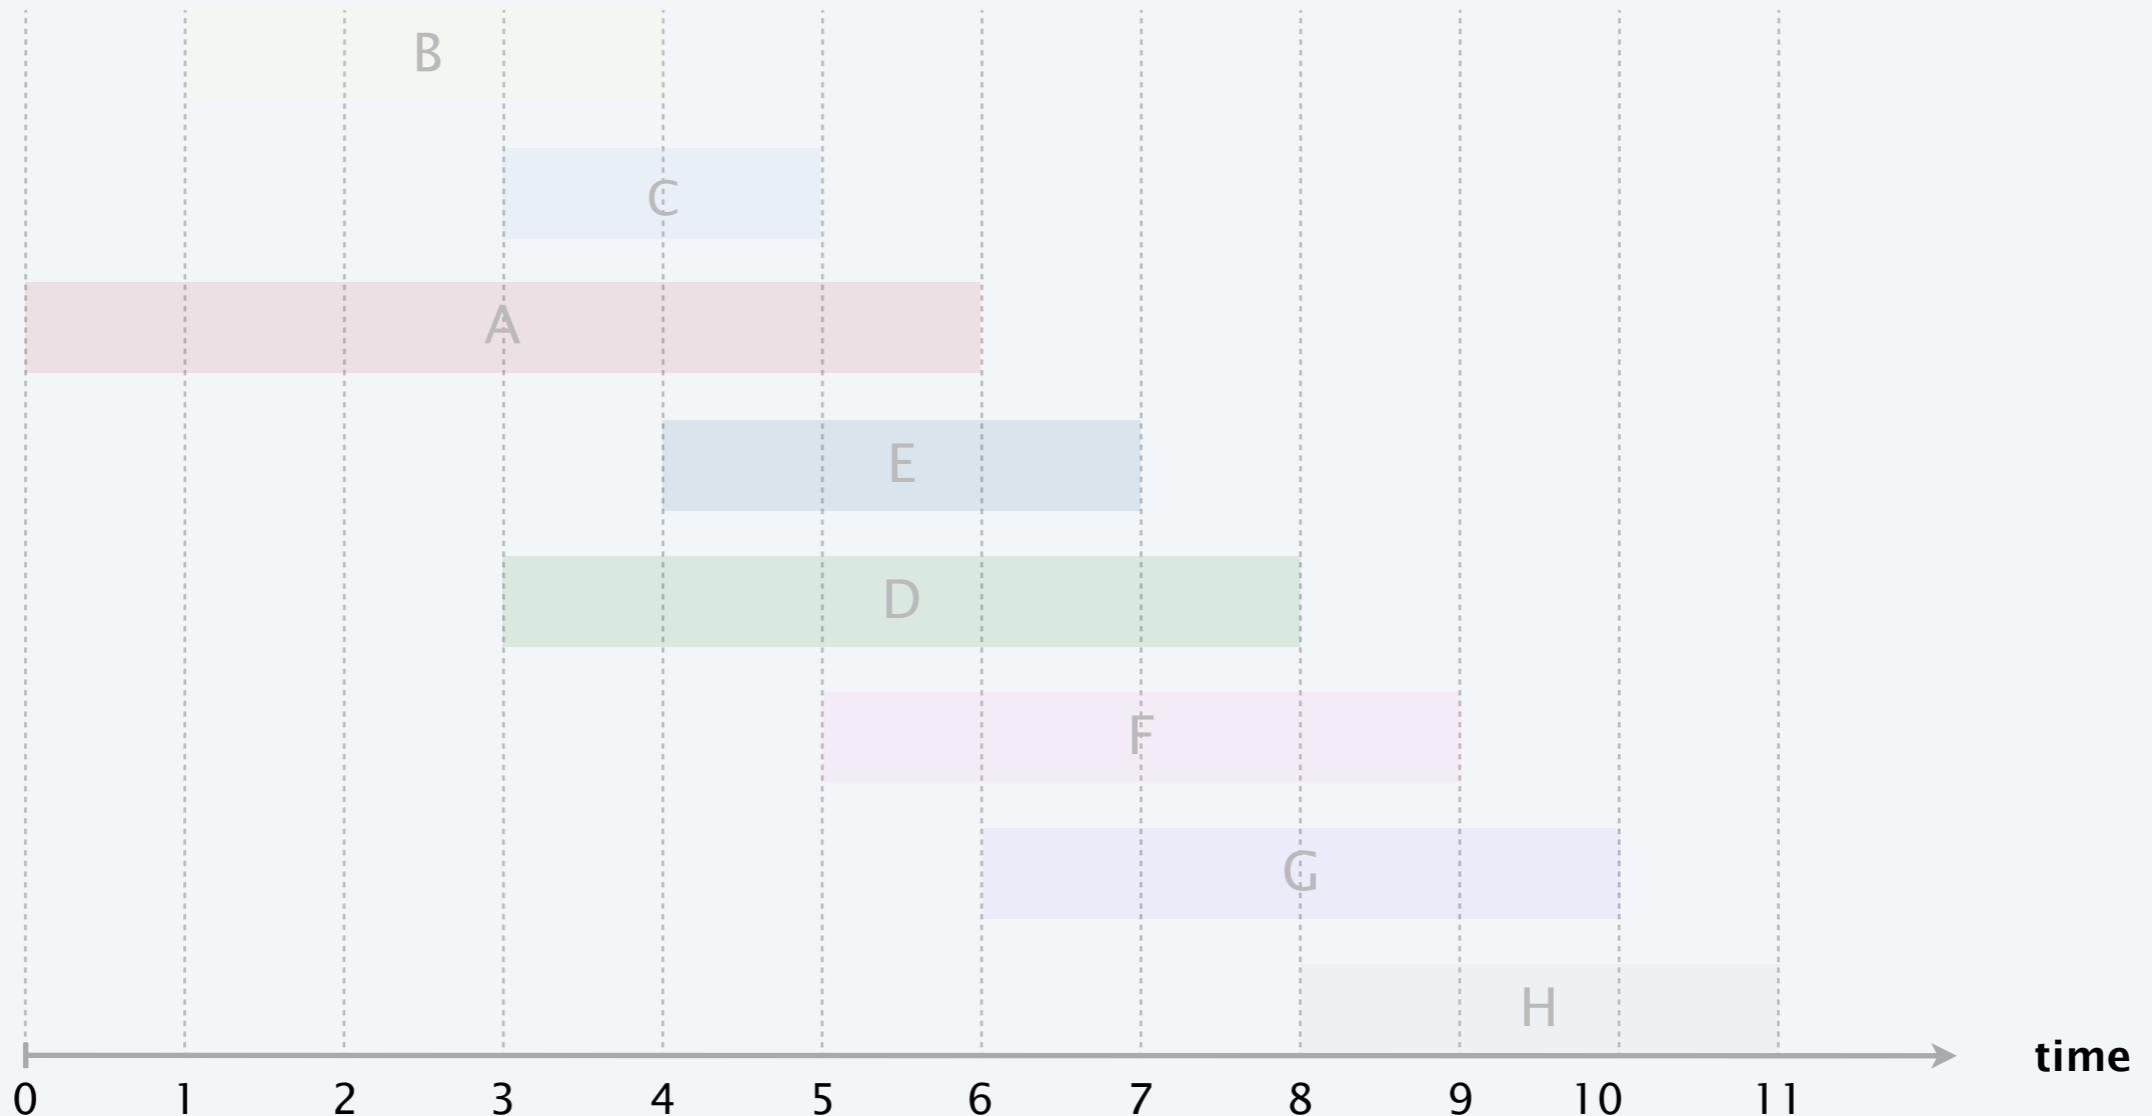
job H is compatible (add to schedule)



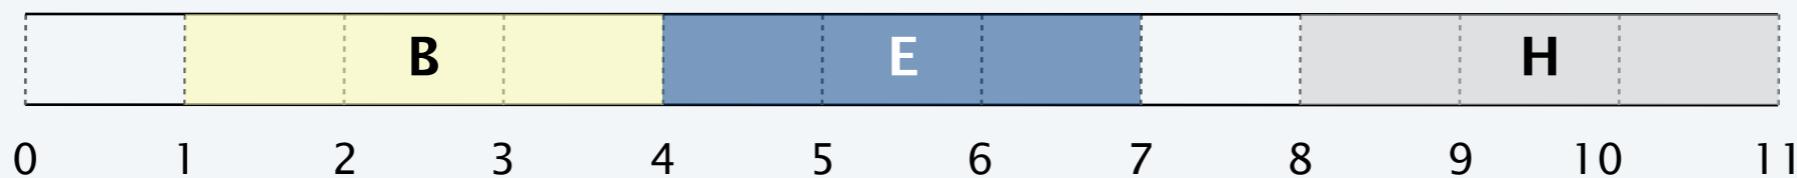
# Earliest-finish-time-first algorithm demo

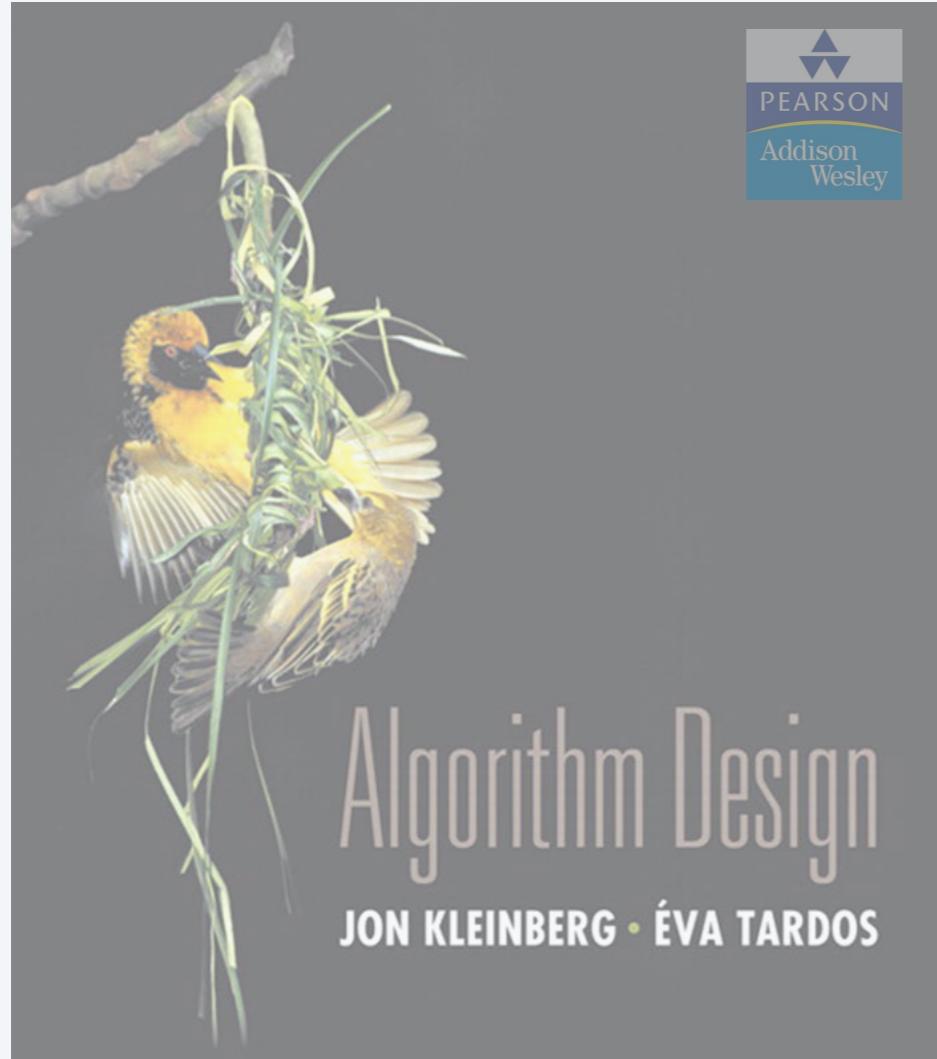


# Earliest-finish-time-first algorithm demo



done (an optimal set of jobs)





## SECTION 4.1

# 4. GREEDY ALGORITHMS I

---

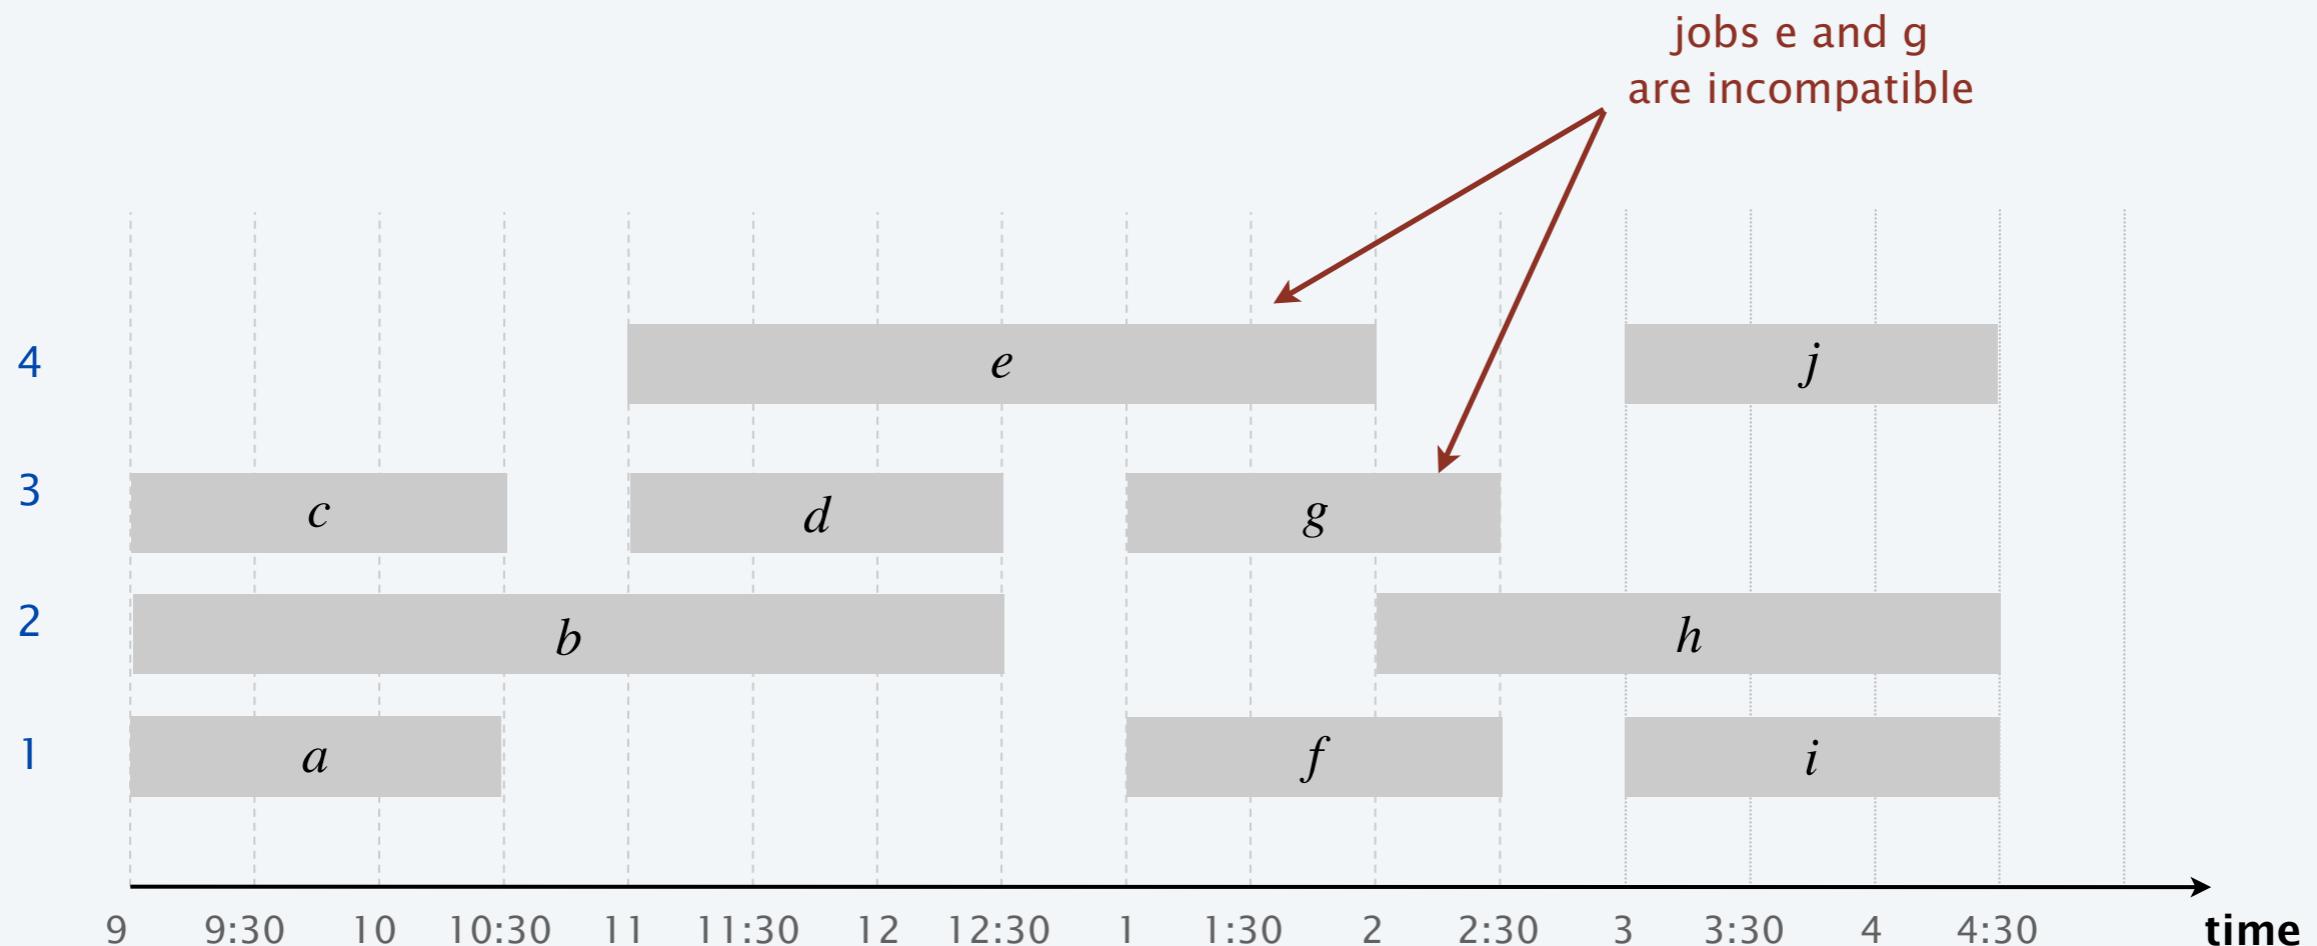
- ▶ *coin changing*
- ▶ *interval scheduling*
- ▶ *interval partitioning*
- ▶ *scheduling to minimize lateness*
- ▶ *optimal caching*

# Interval partitioning

---

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Ex. This schedule uses 4 classrooms to schedule 10 lectures.

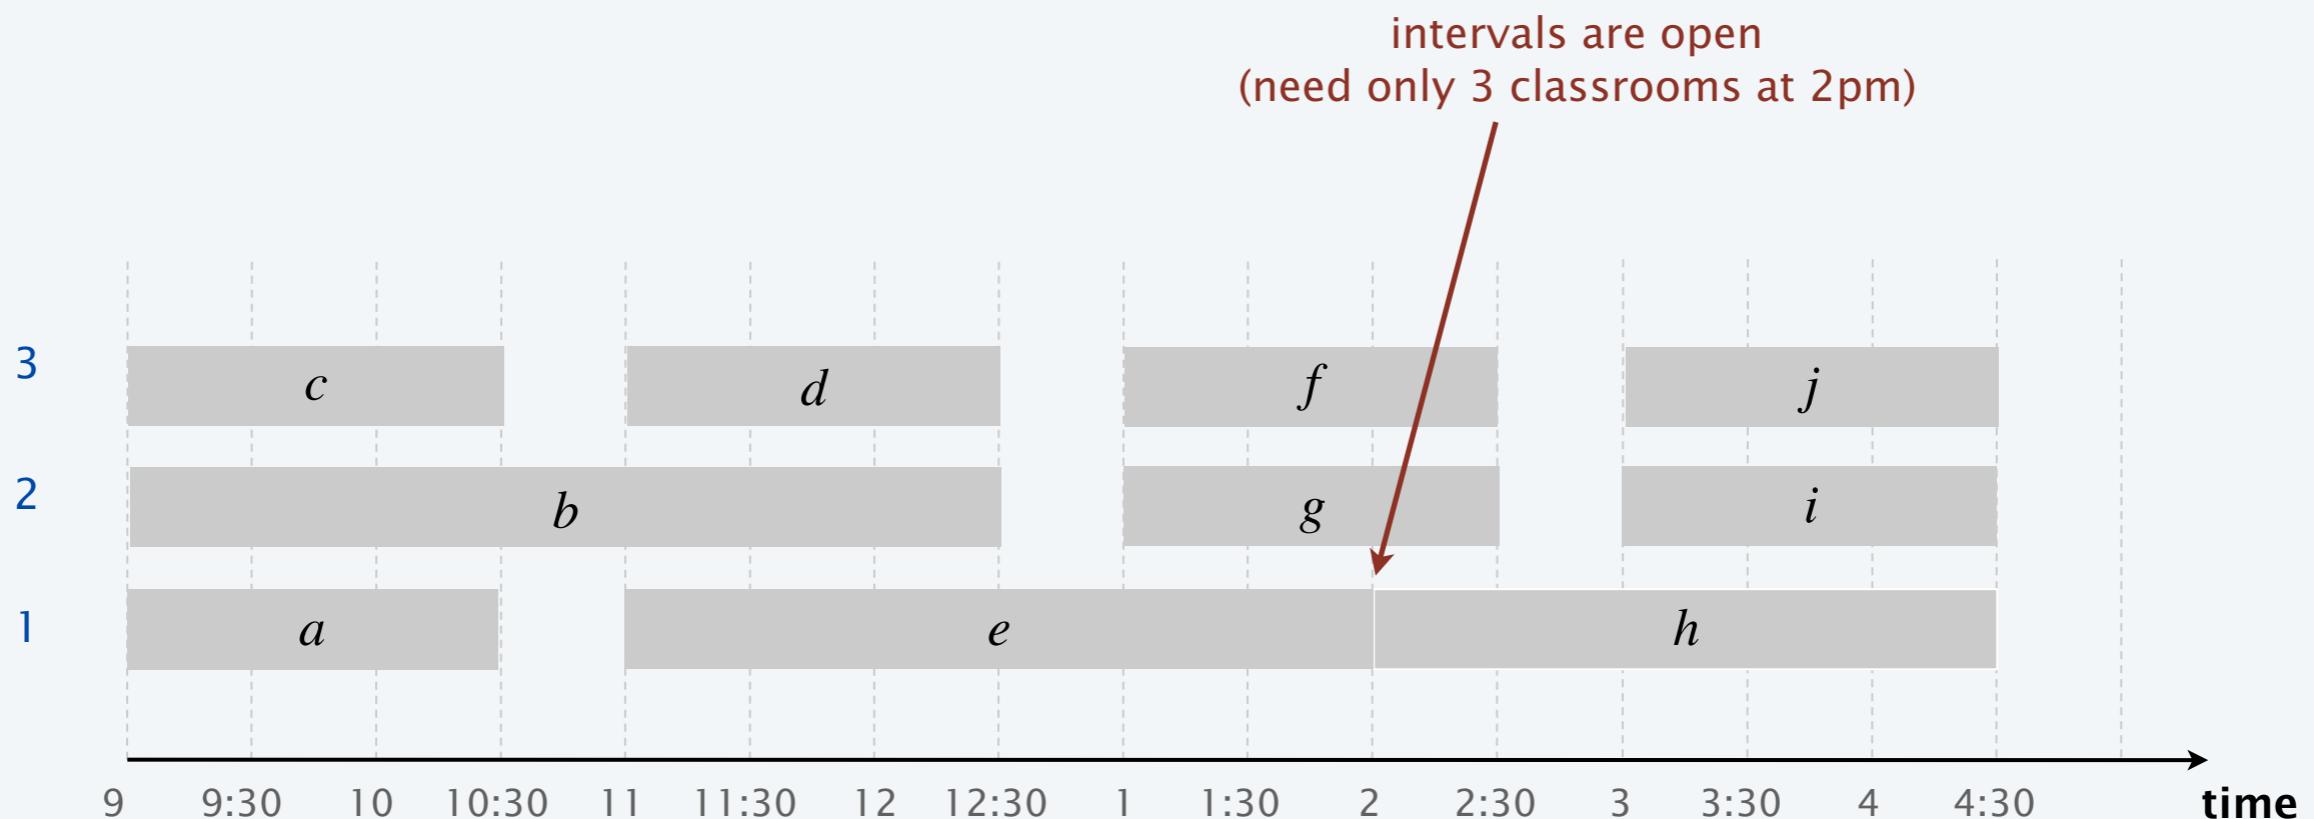


# Interval partitioning

---

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

Ex. This schedule uses 3 classrooms to schedule 10 lectures.





**Consider lectures in some order, assigning each lecture to first available classroom (opening a new classroom if none is available). Which rule is optimal?**

- A. [Earliest start time] Consider lectures in ascending order of  $s_j$ .
- B. [Earliest finish time] Consider lectures in ascending order of  $f_j$ .
- C. [Shortest interval] Consider lectures in ascending order of  $f_j - s_j$ .
- D. None of the above.

# Interval partitioning: earliest-start-time-first algorithm

---



**EARLIEST-START-TIME-FIRST** ( $n, s_1, s_2, \dots, s_n, f_1, f_2, \dots, f_n$ )

**SORT** lectures by start times and renumber so that  $s_1 \leq s_2 \leq \dots \leq s_n$ .

$d \leftarrow 0$ . ← number of allocated classrooms

**FOR**  $j = 1$  **TO**  $n$

**IF** (lecture  $j$  is compatible with some classroom)

        Schedule lecture  $j$  in any such classroom  $k$ .

**ELSE**

        Allocate a new classroom  $d + 1$ .

        Schedule lecture  $j$  in classroom  $d + 1$ .

$d \leftarrow d + 1$ .

**RETURN** schedule.

## Interval partitioning: earliest-start-time-first algorithm

---

**Proposition.** The earliest-start-time-first algorithm can be implemented in  $O(n \log n)$  time.

Pf.

- Sorting by start times takes  $O(n \log n)$  time.
- Store classrooms in a **priority queue** (key = finish time of its last lecture).
  - to allocate a new classroom, INSERT classroom onto priority queue.
  - to schedule lecture  $j$  in classroom  $k$ , INCREASE-KEY of classroom  $k$  to  $f_j$ .
  - to determine whether lecture  $j$  is compatible with any classroom, compare  $s_j$  to FIND-MIN
- Total # of priority queue operations is  $O(n)$ ; each takes  $O(\log n)$  time. ▀

**Remark.** This implementation chooses a classroom  $k$  whose finish time of its last lecture is the **earliest**.

## Interval partitioning: lower bound on optimal solution

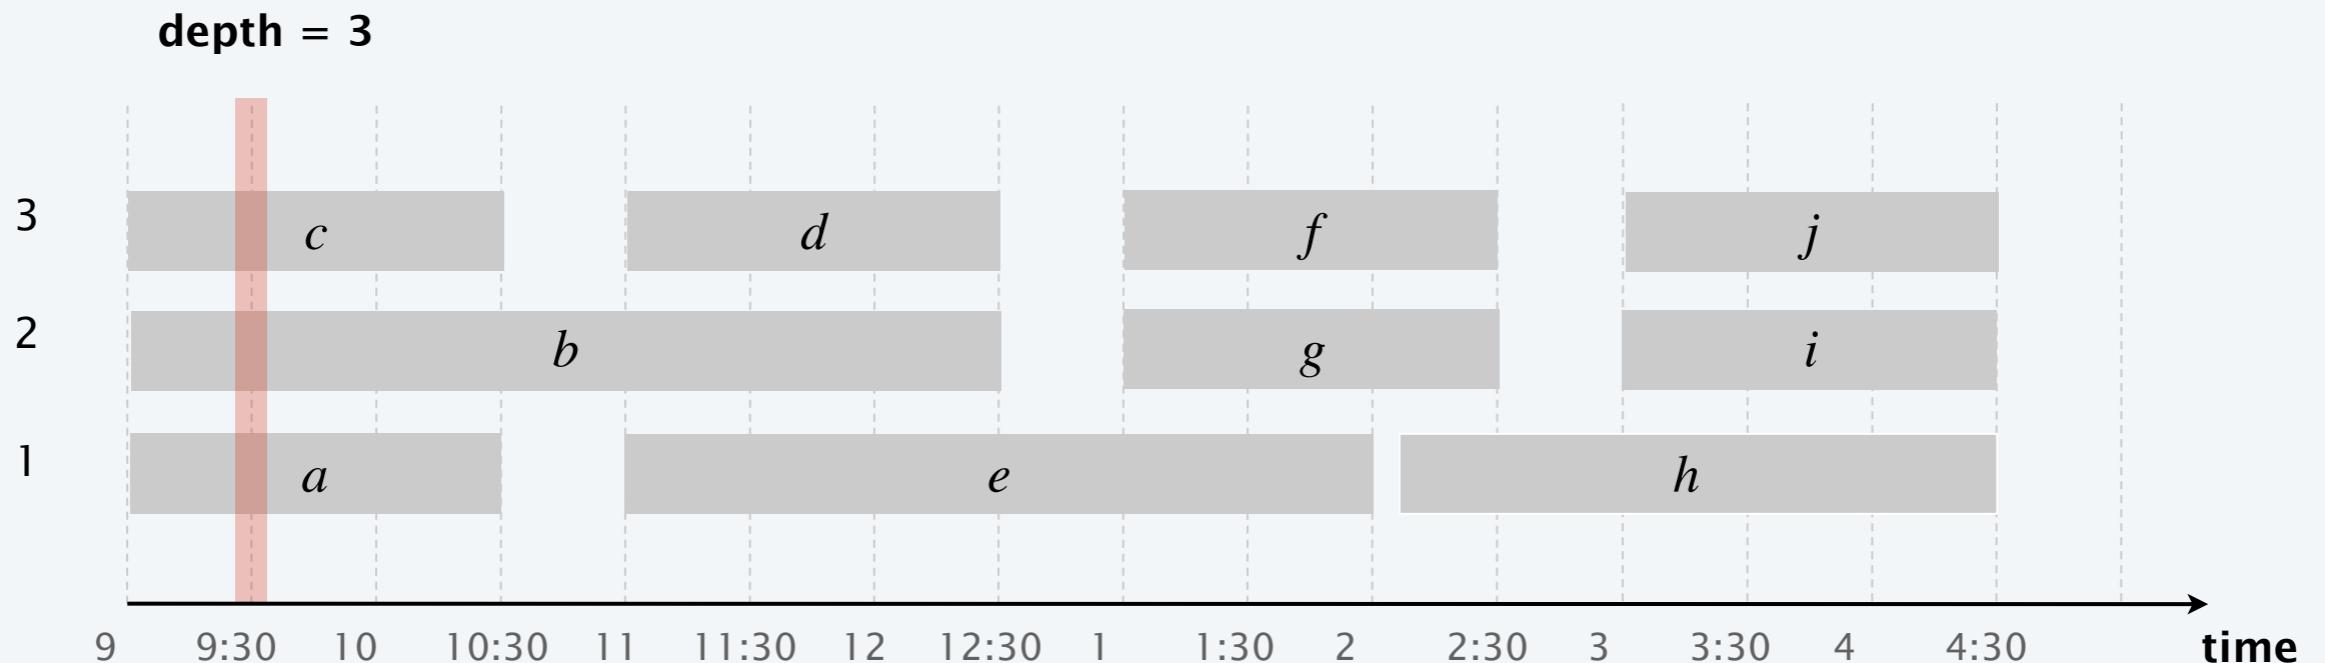
---

**Def.** The **depth** of a set of open intervals is the maximum number of intervals that contain any given point.

**Key observation.** Number of classrooms needed  $\geq$  depth.

**Q.** Does minimum number of classrooms needed always equal depth?

**A.** Yes! Moreover, earliest-start-time-first algorithm finds a schedule whose number of classrooms equals the depth.



## Interval partitioning: analysis of earliest-start-time-first algorithm

---

**Observation.** The earliest-start-time first algorithm never schedules two incompatible lectures in the same classroom.

**Theorem.** Earliest-start-time-first algorithm is optimal.

Pf.

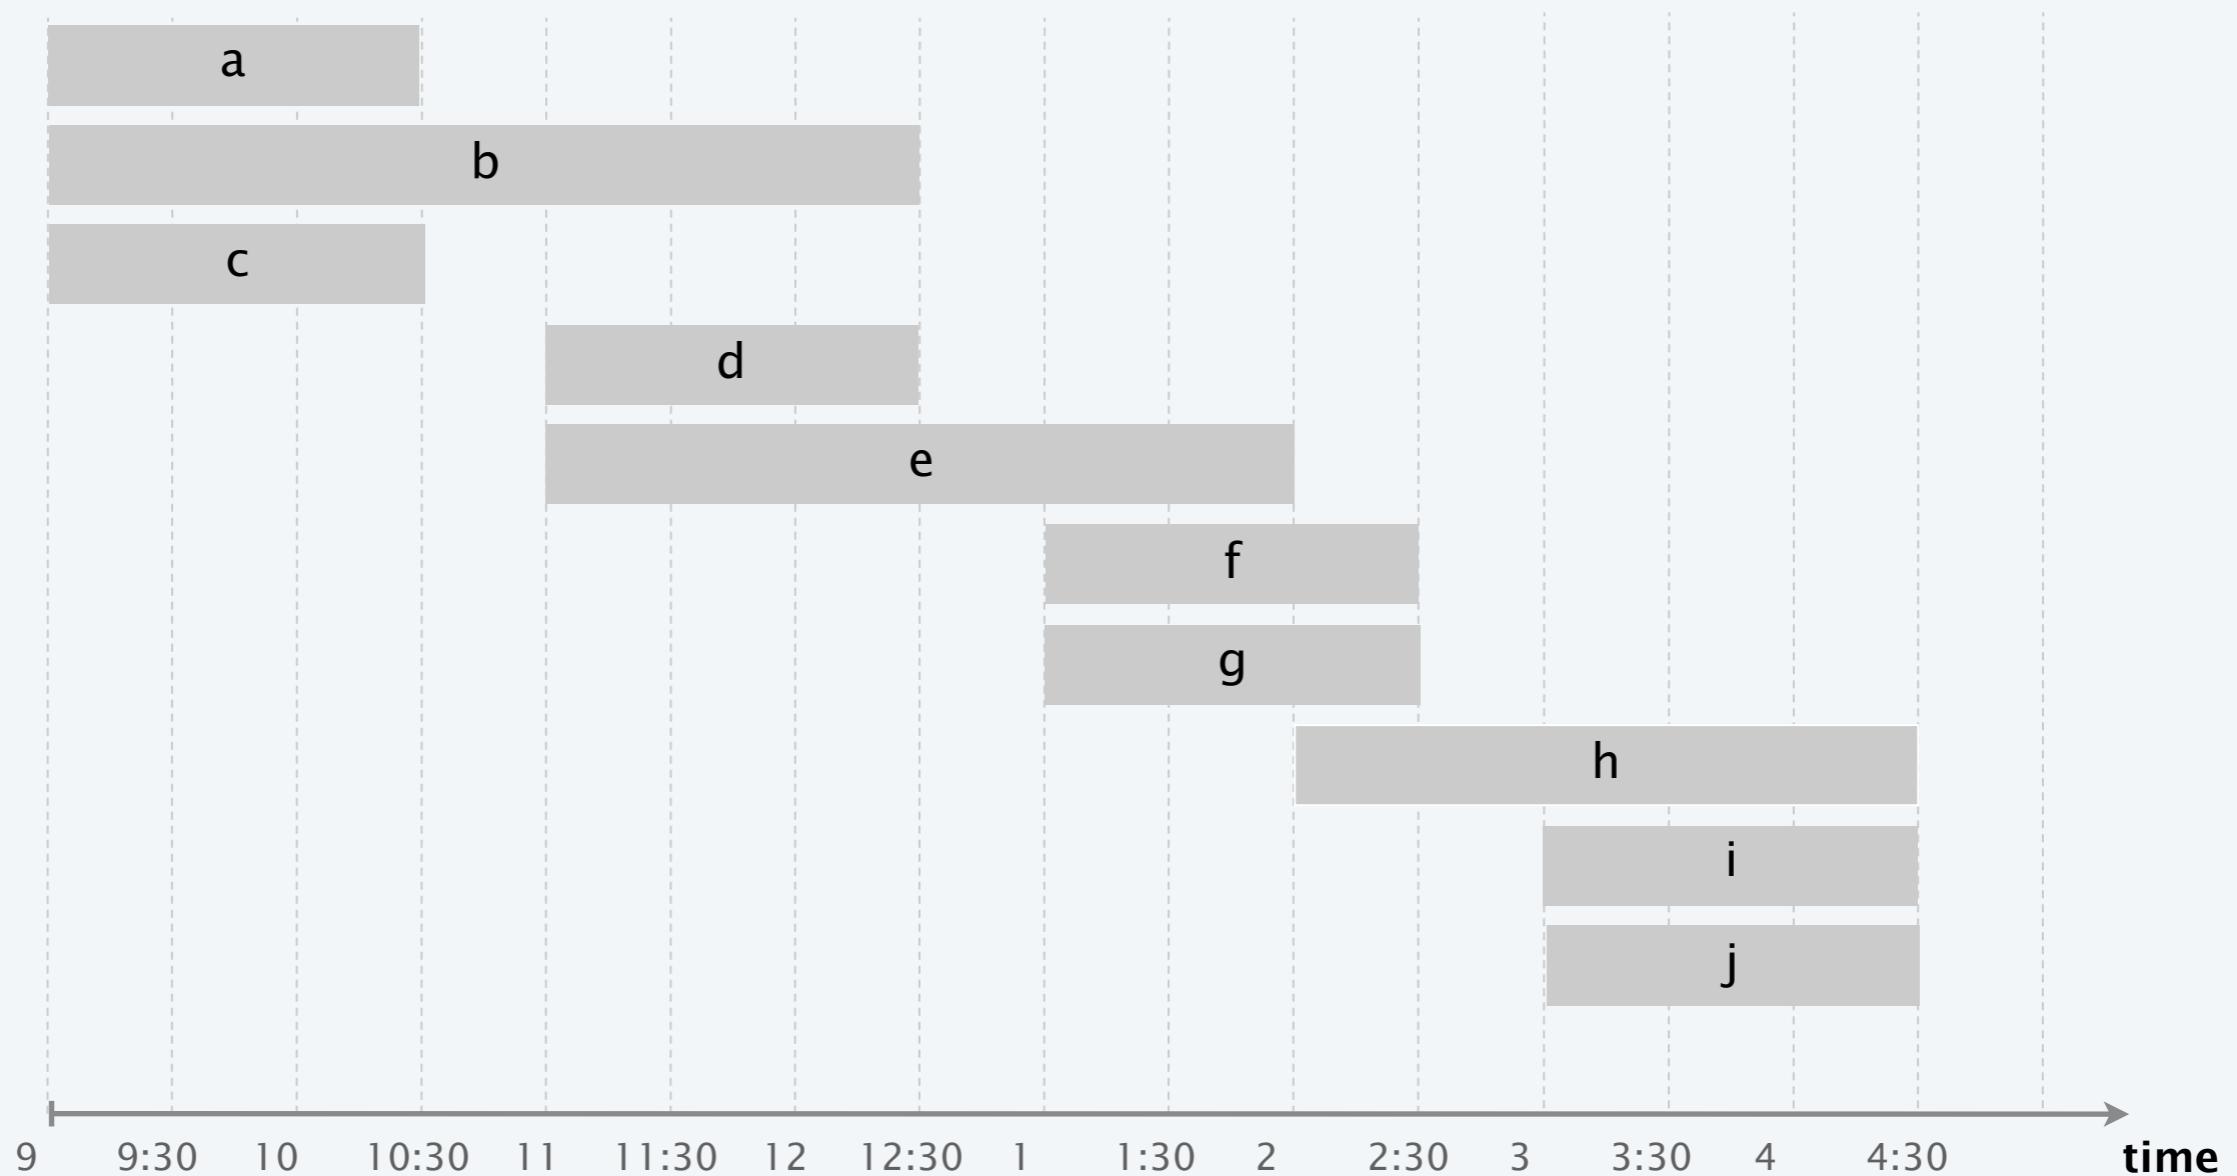
- Let  $d = \text{number of classrooms that the algorithm allocates.}$
- Classroom  $d$  is opened because we needed to schedule a lecture, say  $j$ , that is incompatible with a lecture in each of  $d - 1$  other classrooms.
- Thus, these  $d$  lectures each end after  $s_j$ .
- Since we sorted by start time, each of these incompatible lectures start no later than  $s_j$ .
- Thus, we have  $d$  lectures overlapping at time  $s_j + \varepsilon$ .
- Key observation  $\Rightarrow$  all schedules use  $\geq d$  classrooms. ■

# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.



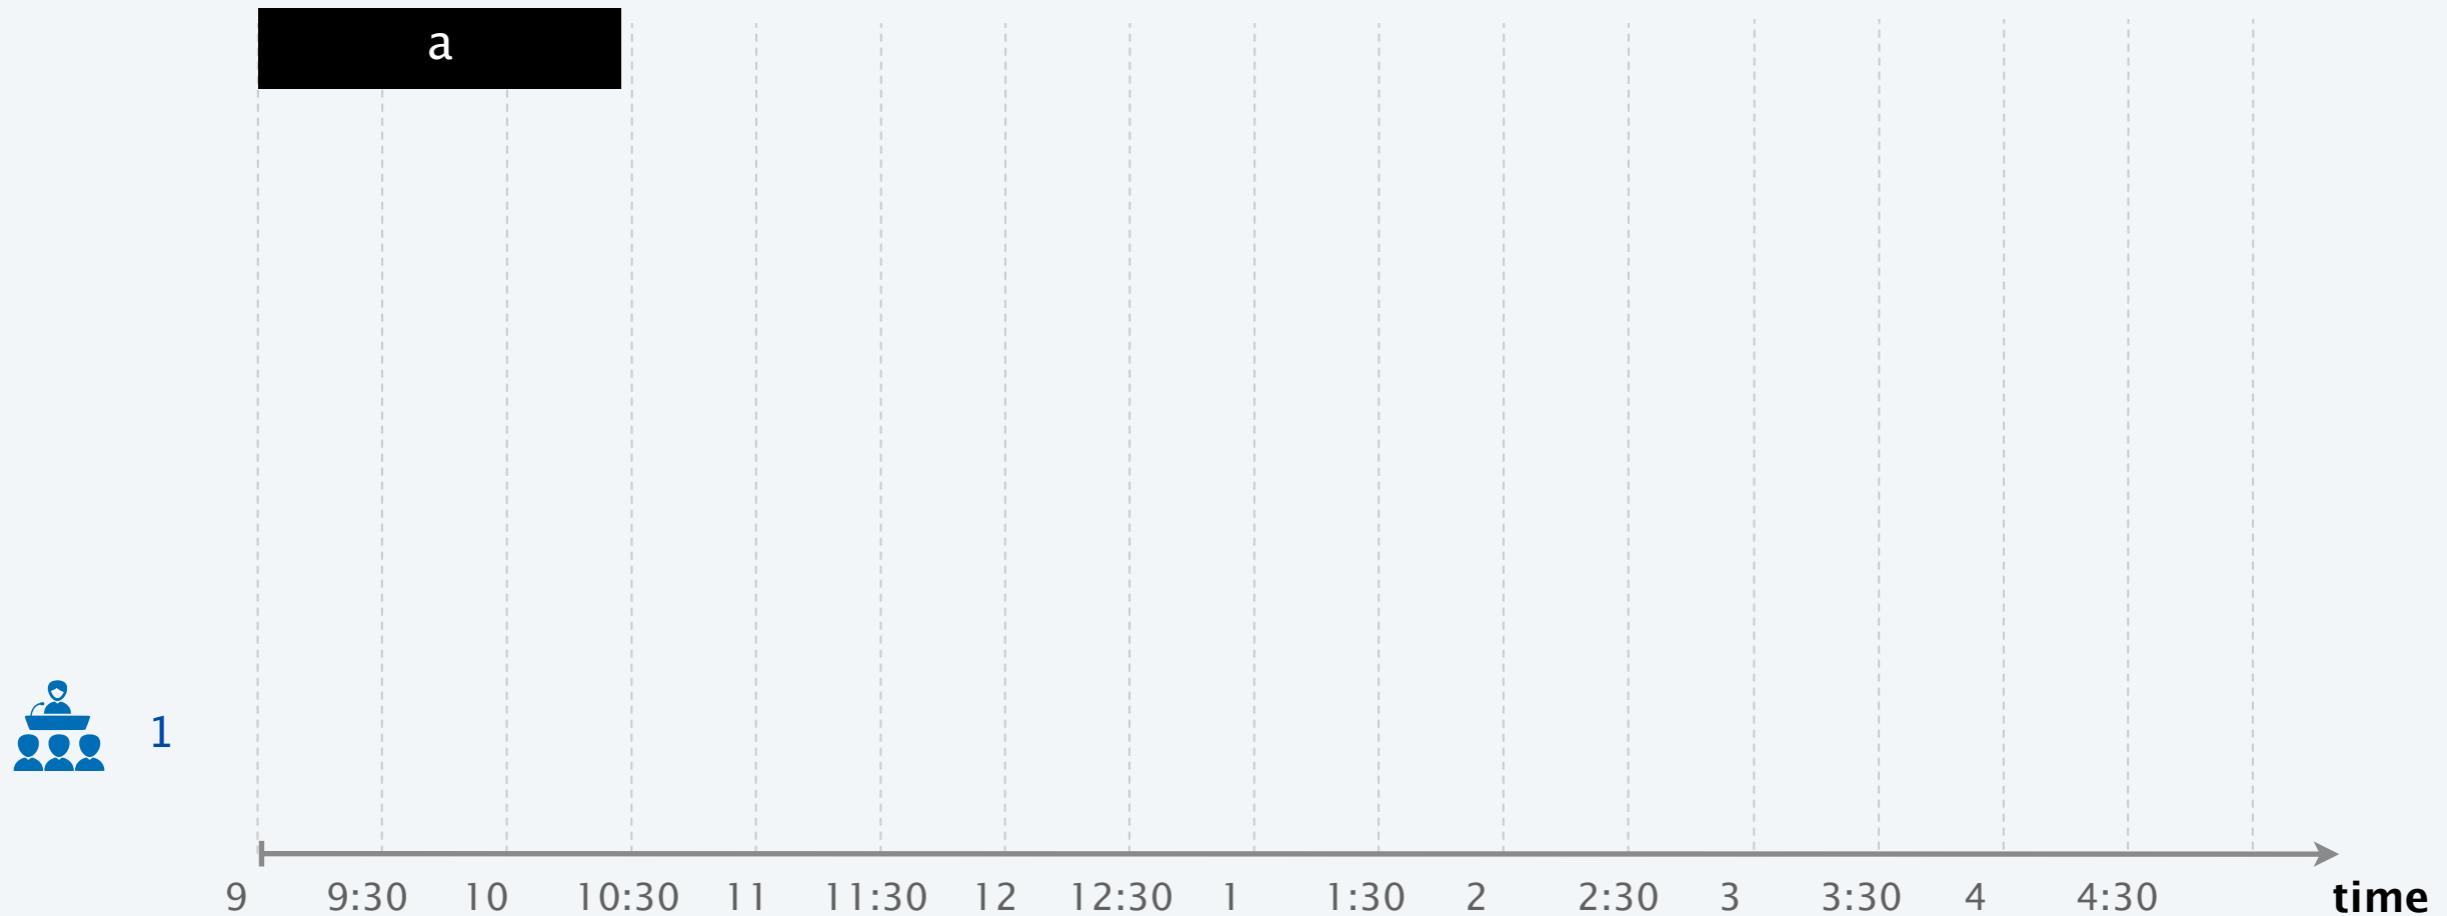
# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**no compatible classroom: open up a new classroom and assign lecture to it**



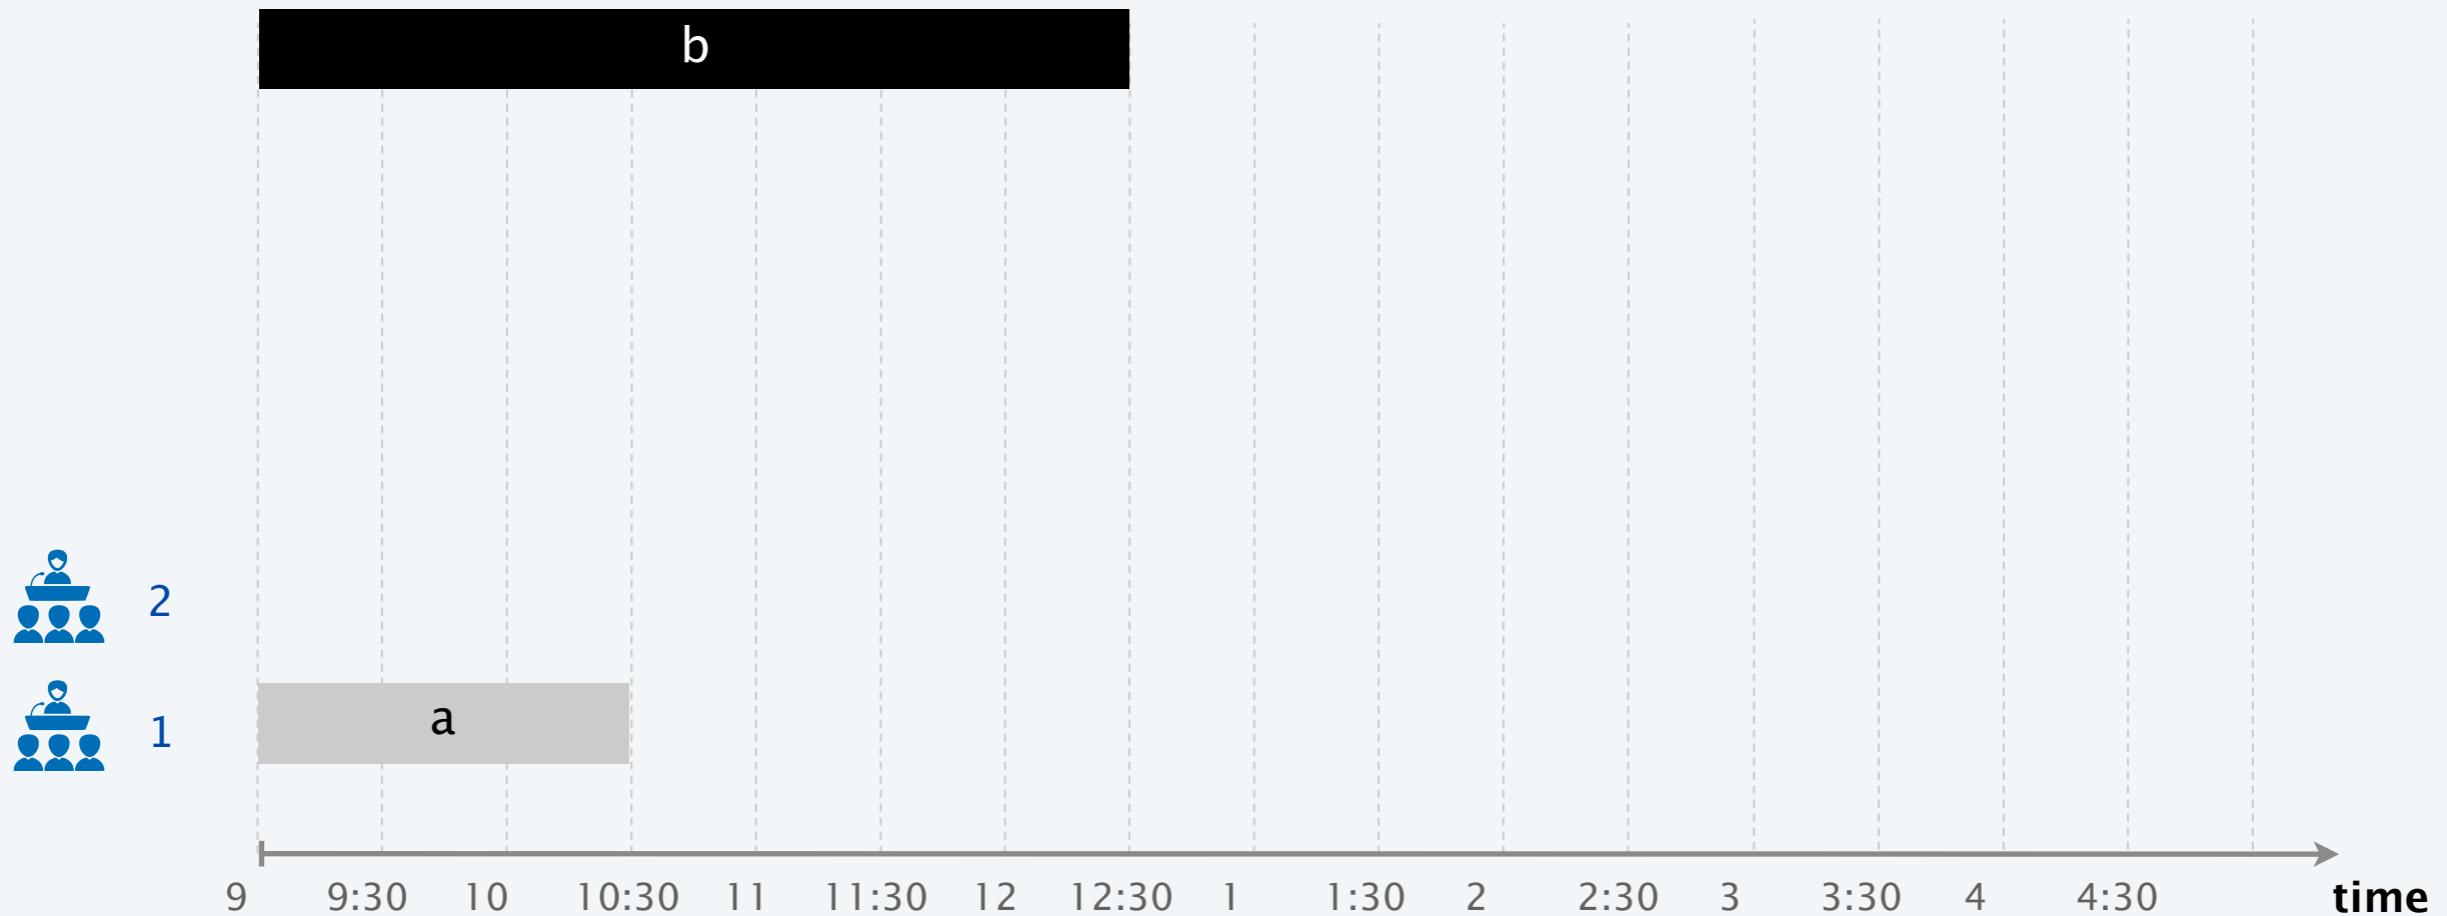
# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**no compatible classroom: open up a new classroom and assign lecture to it**



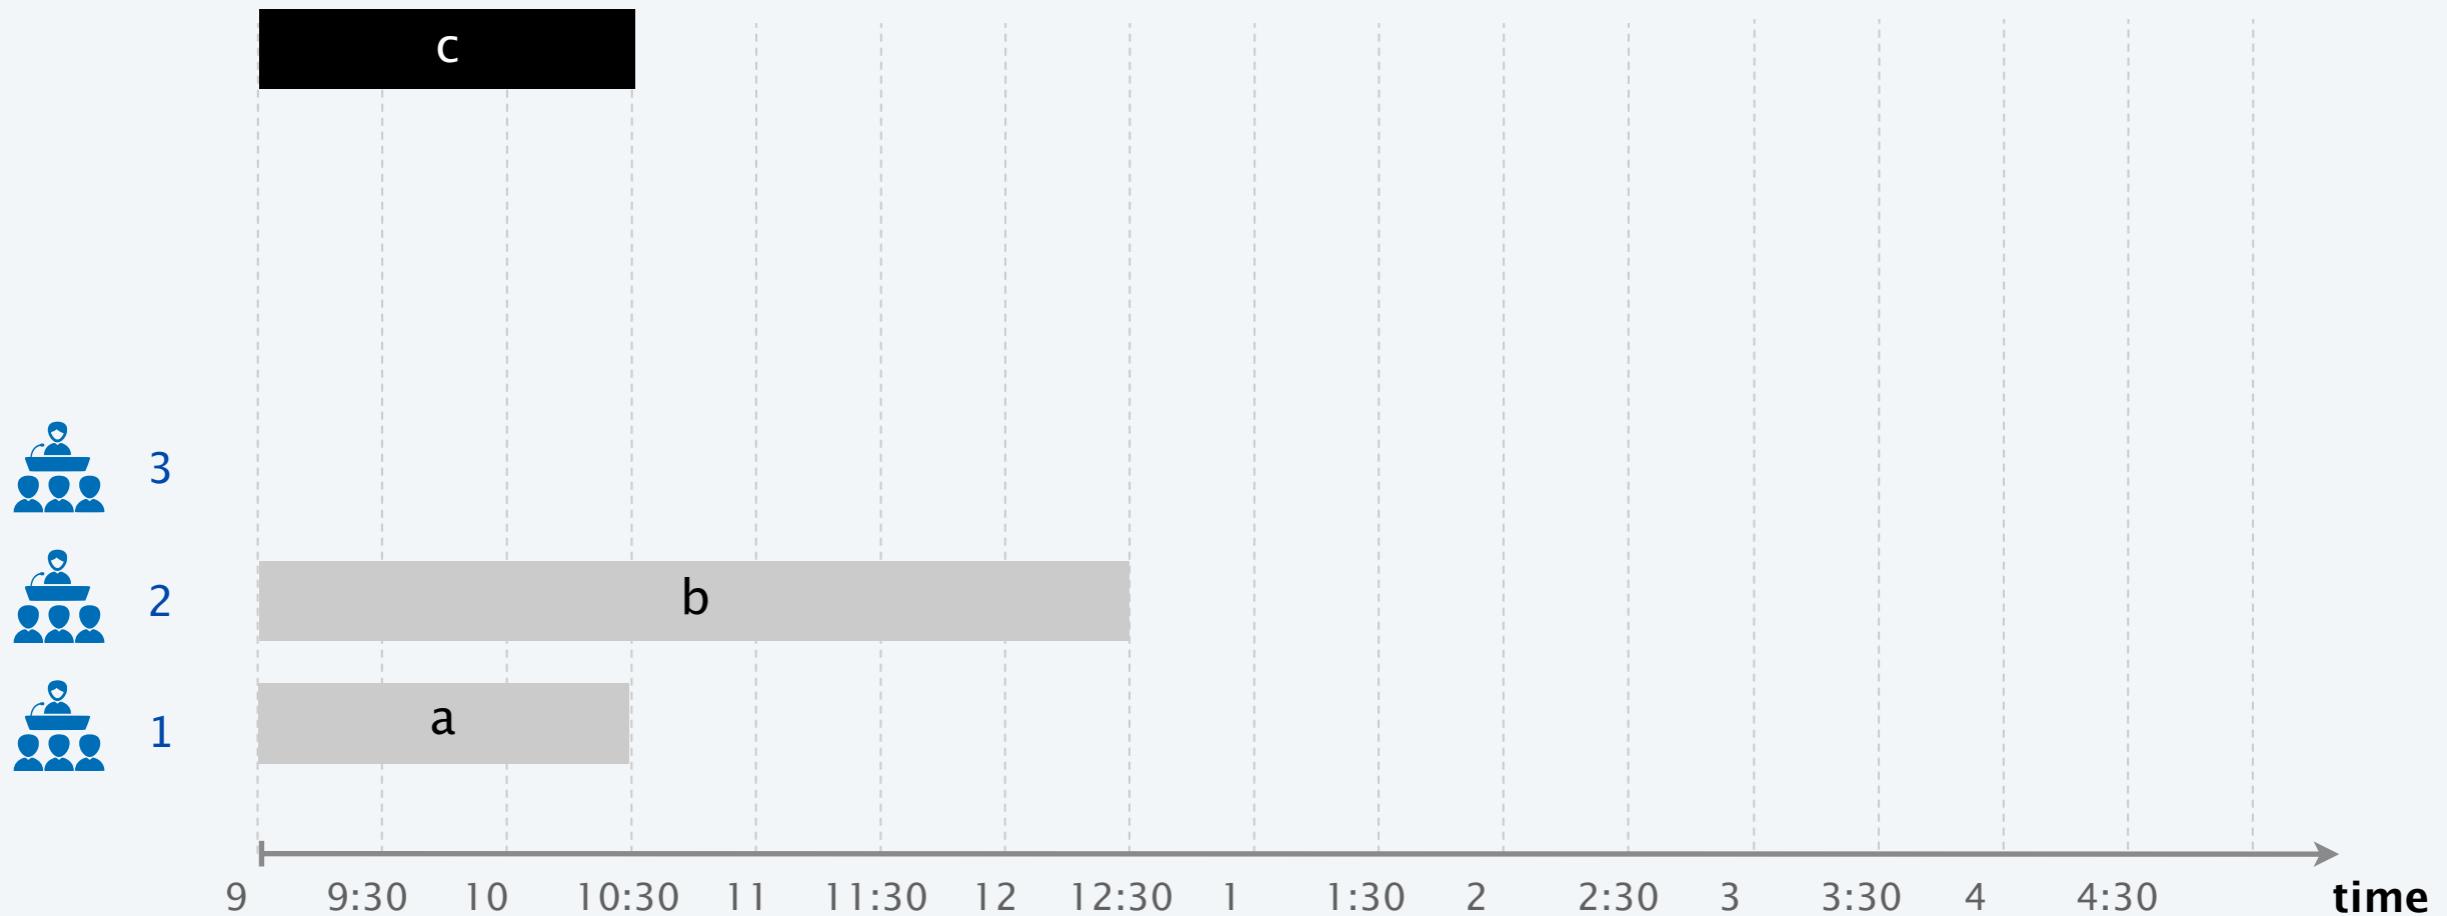
# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**no compatible classroom: open up a new classroom and assign lecture to it**



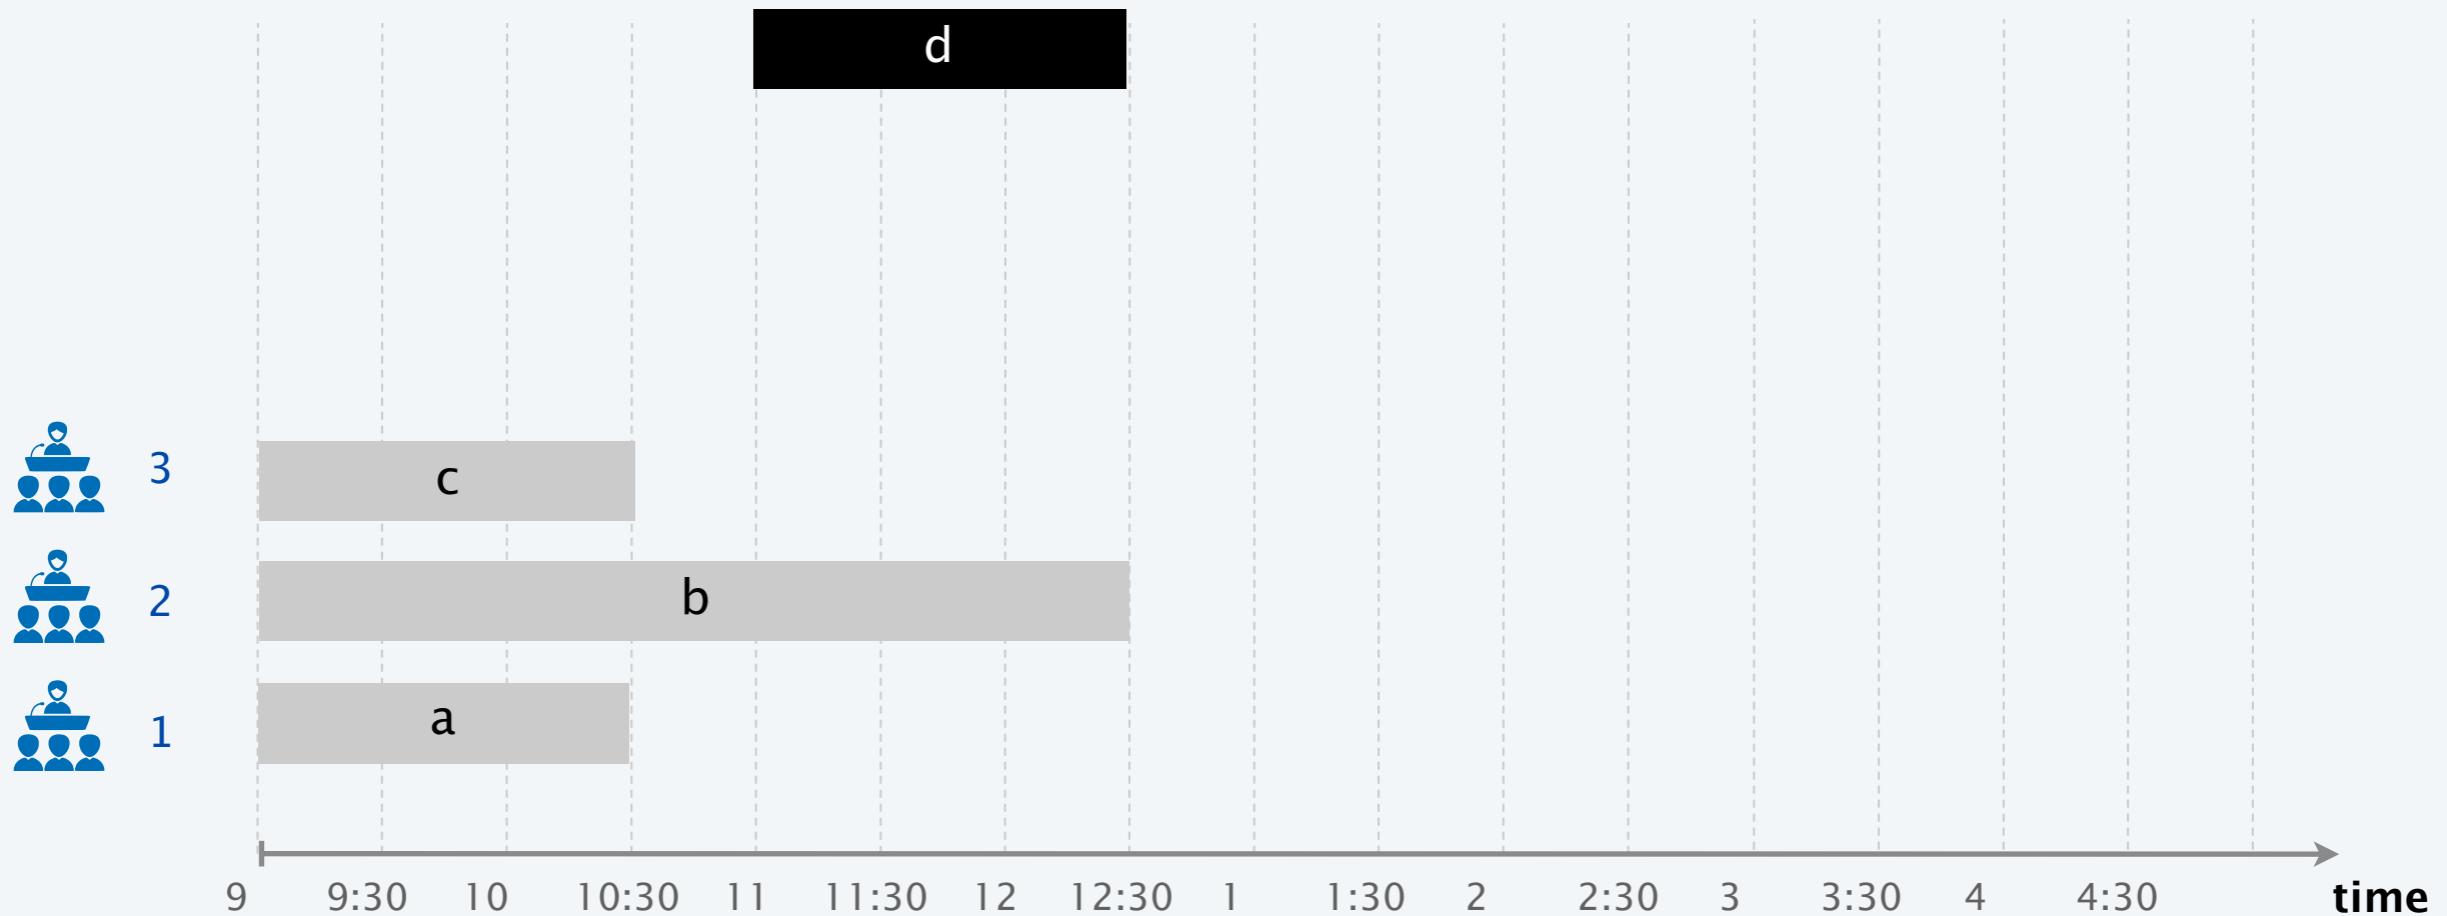
# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture d is compatible with classrooms 1 and 3**



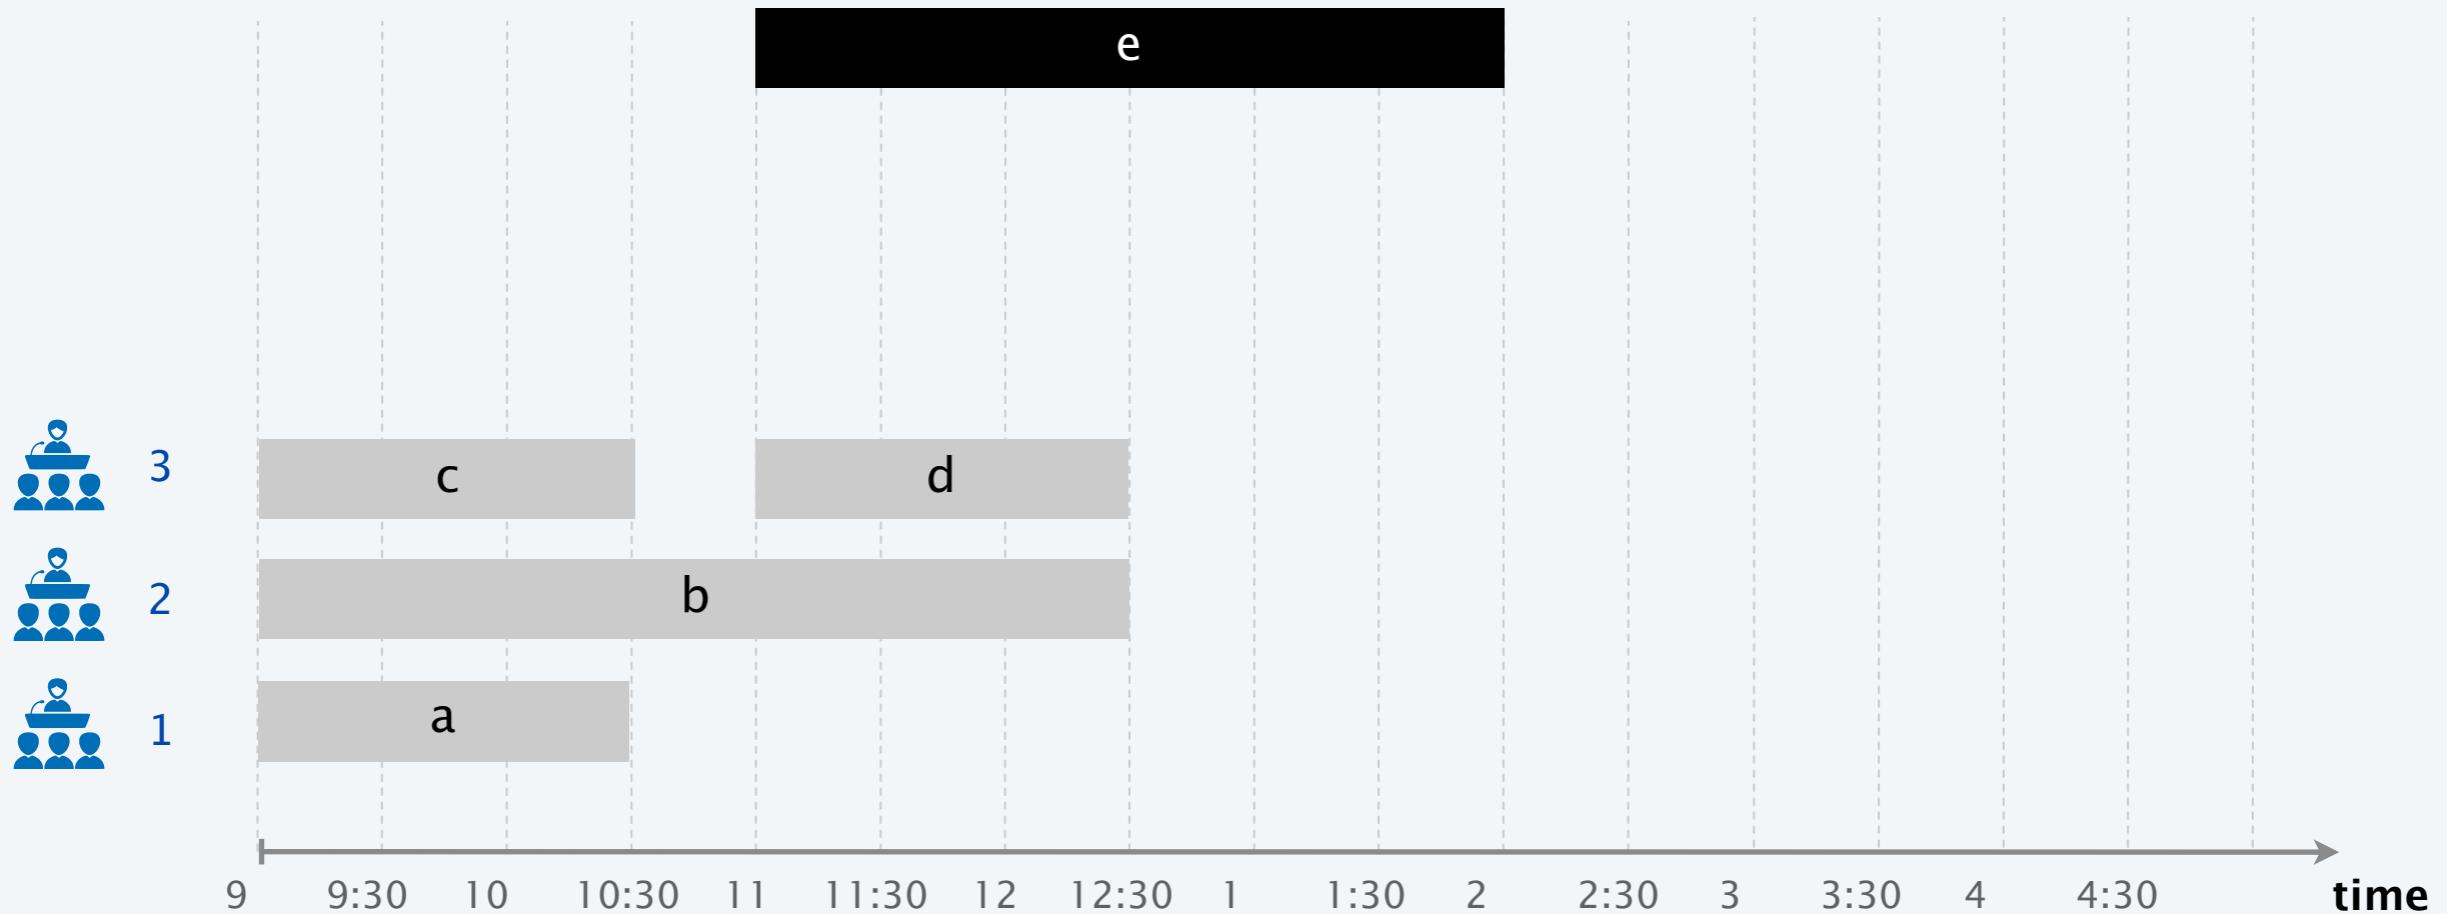
# Earliest-start-time-first algorithm demo

---

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture e is compatible with classroom 1**

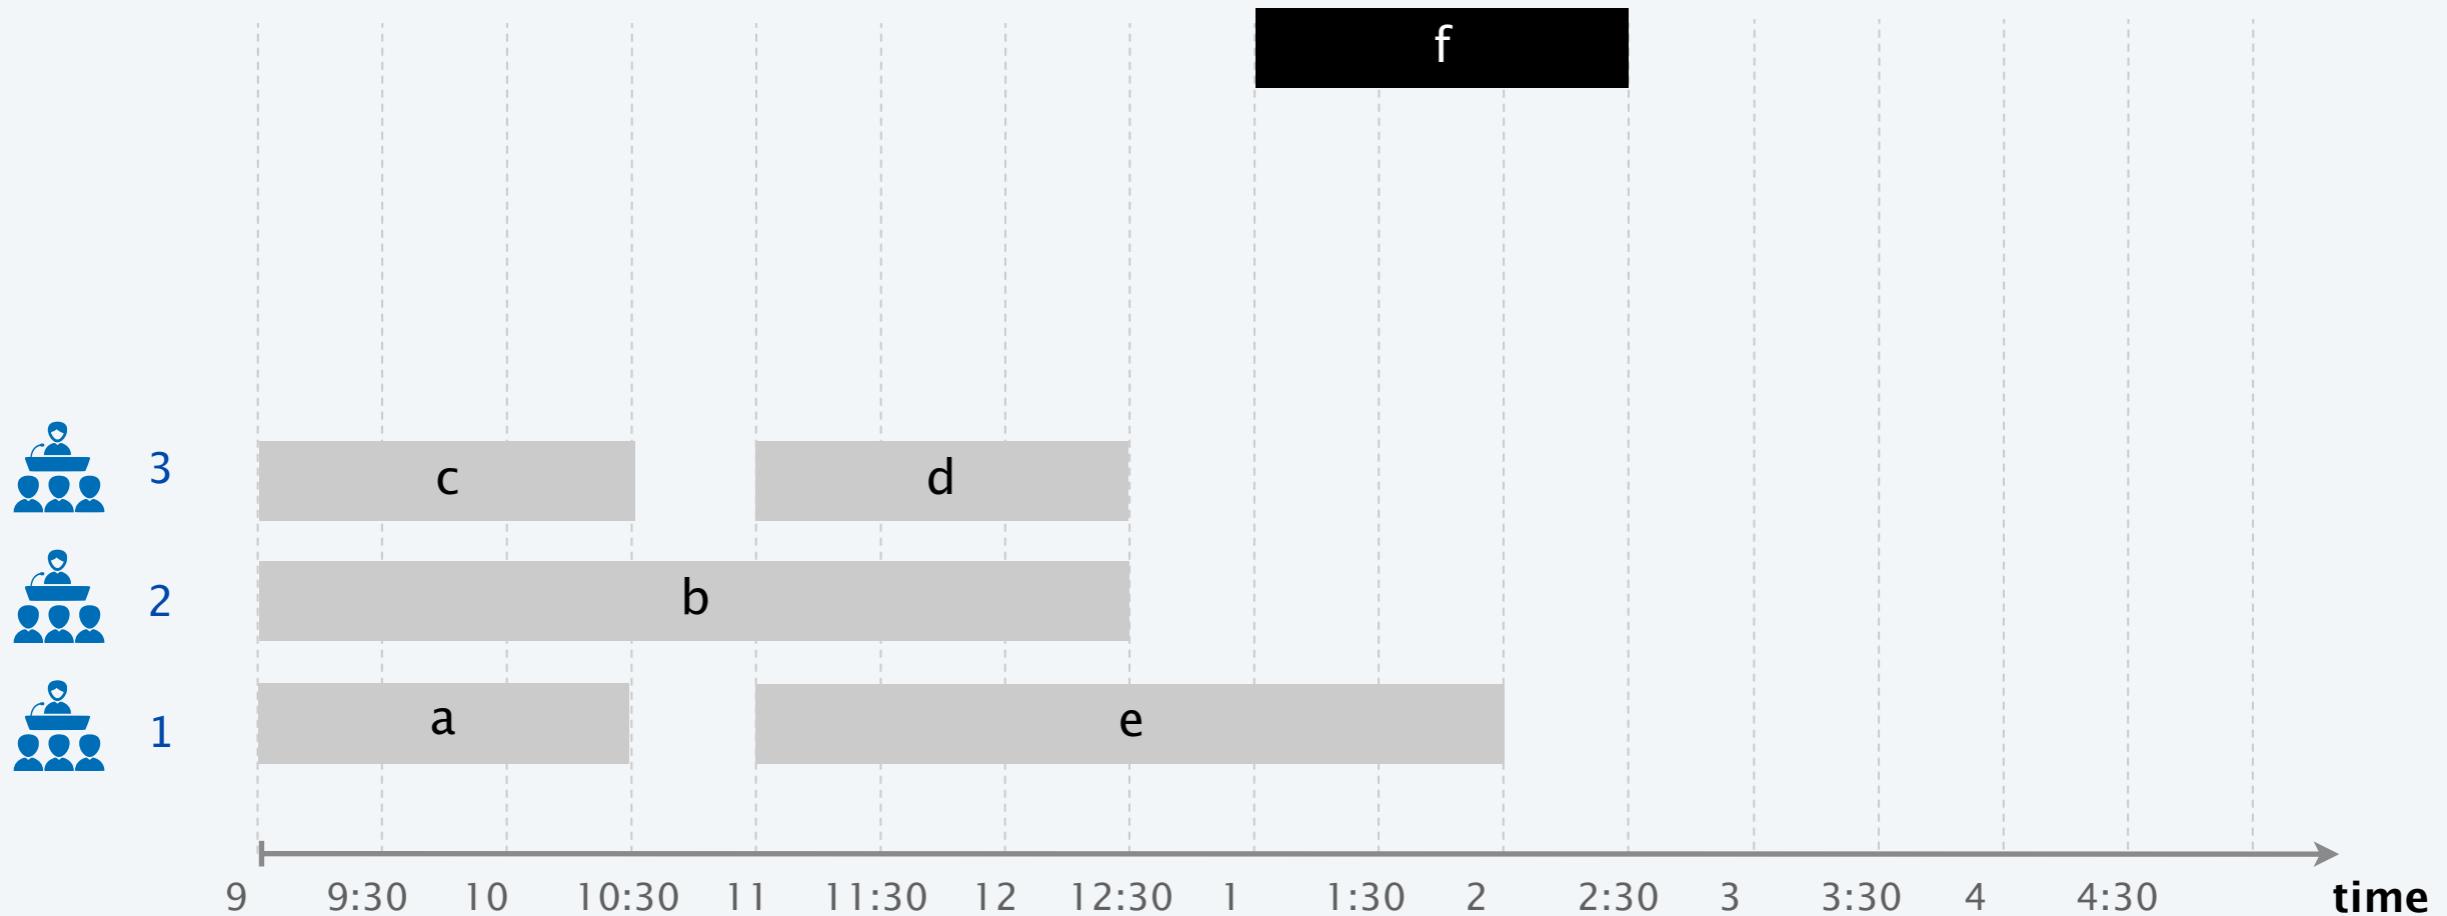


# Earliest-start-time-first algorithm demo

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

lecture f is compatible with classroom 2 and 3

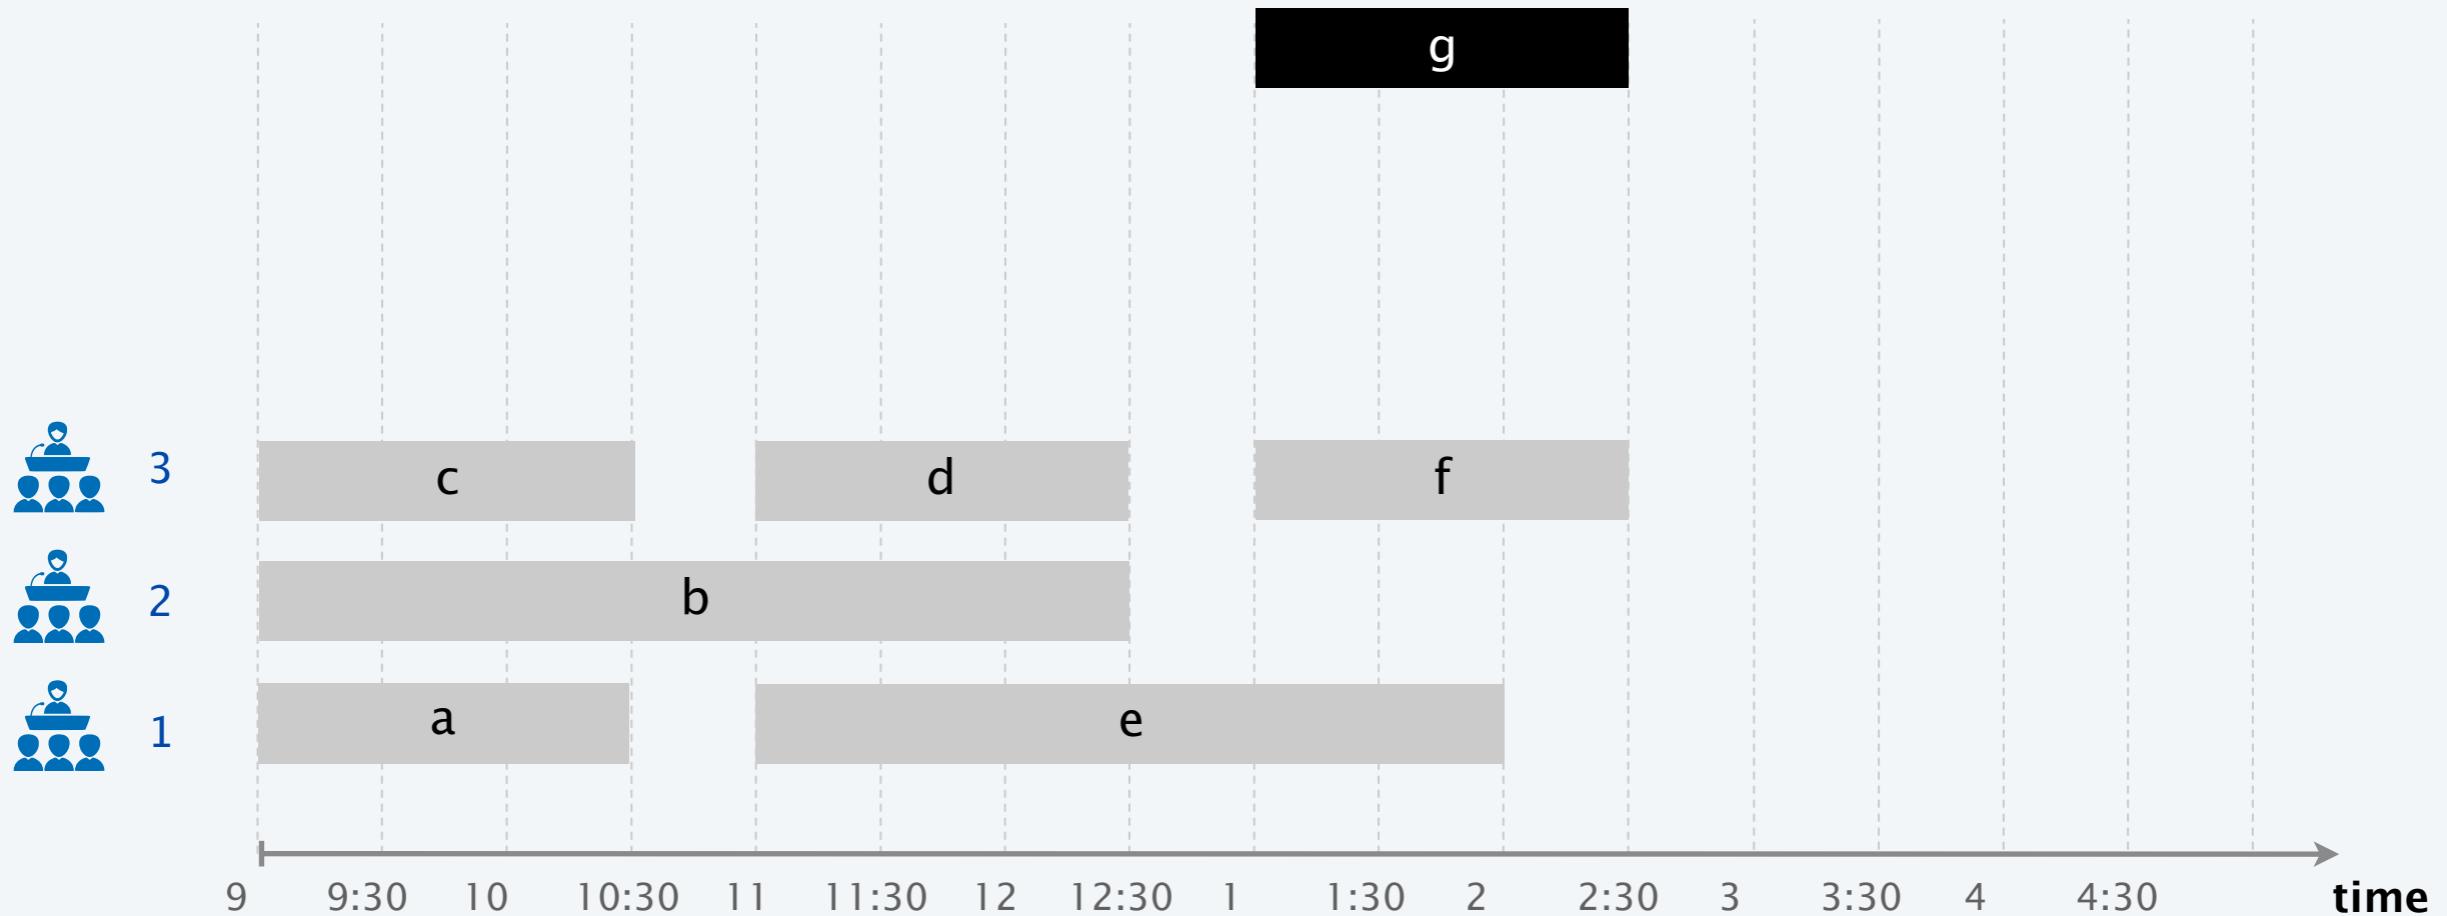


# Earliest-start-time-first algorithm demo

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

lecture g is compatible with classroom 2

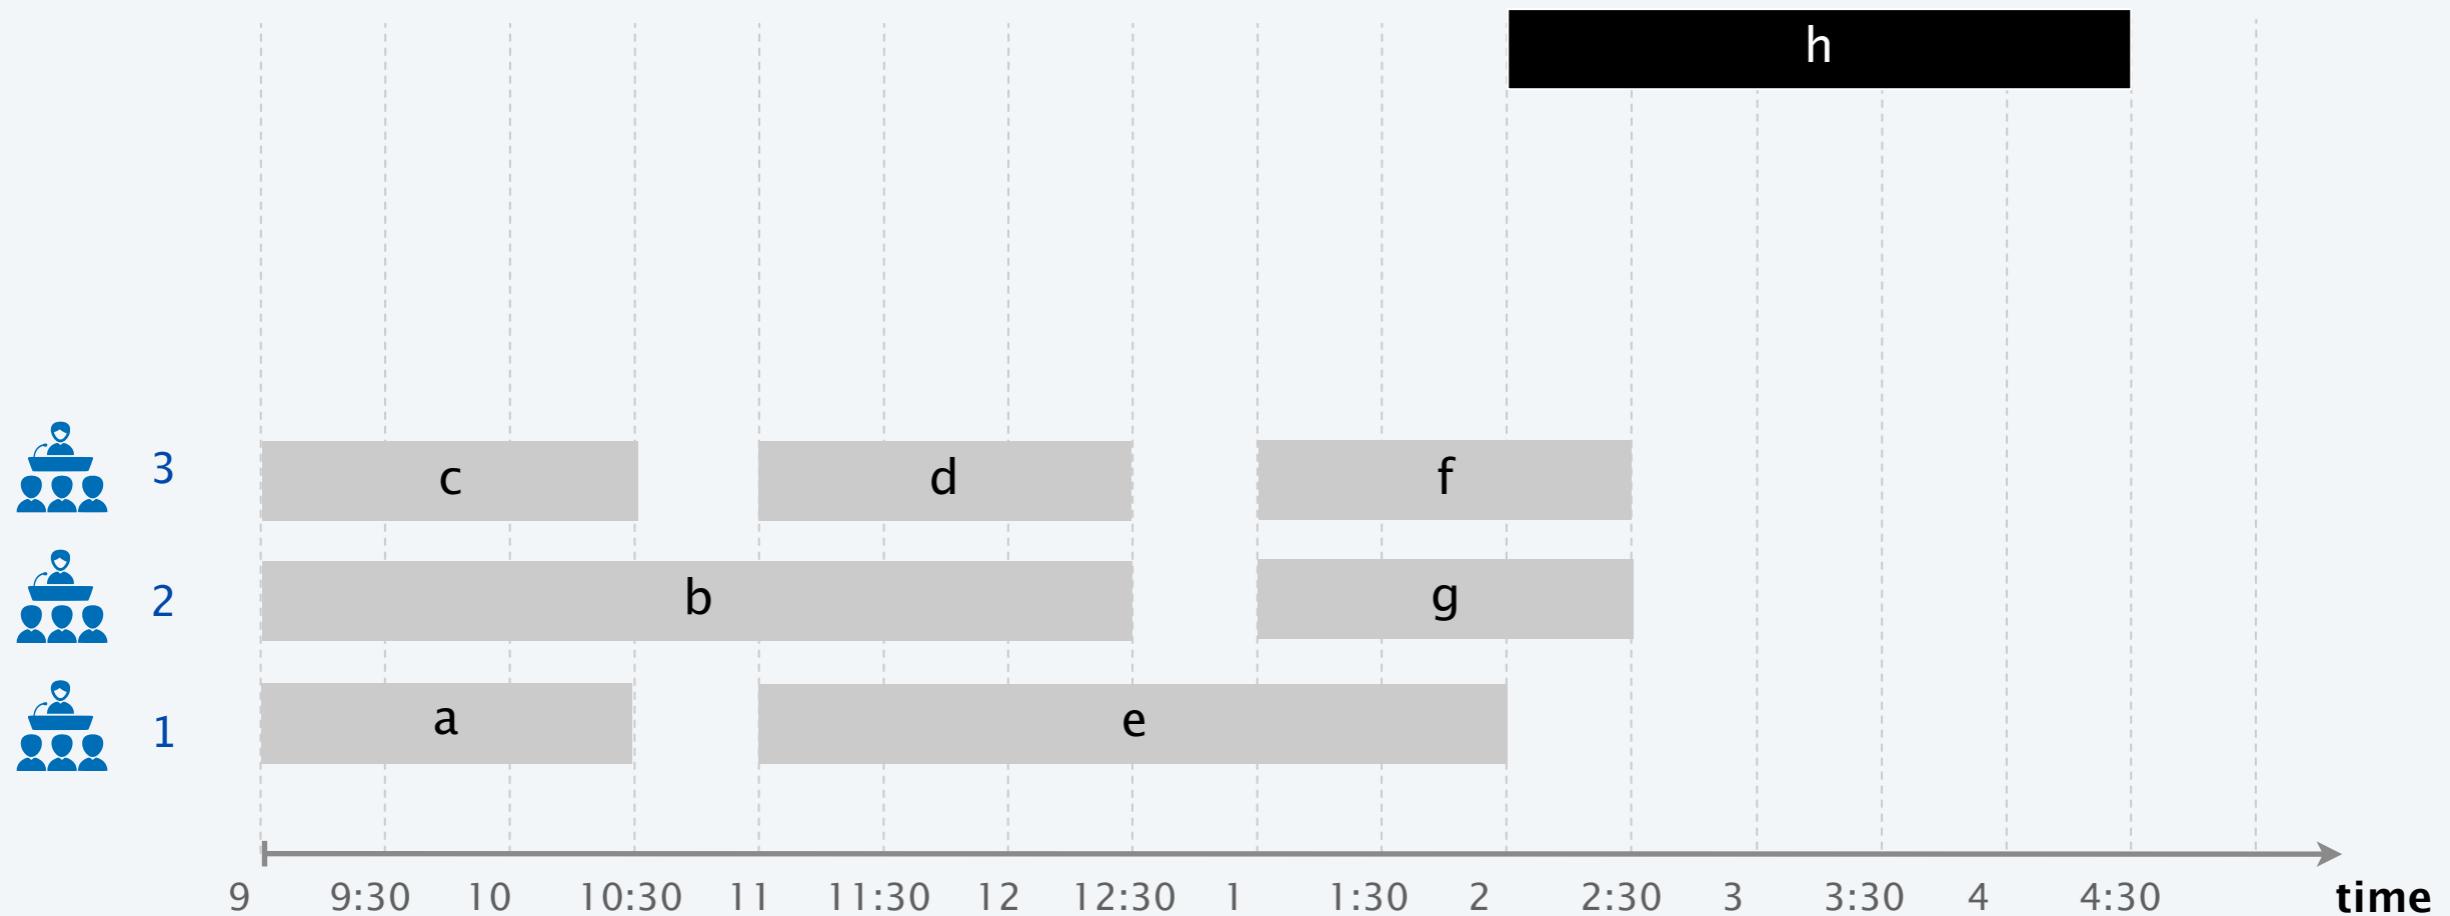


# Earliest-start-time-first algorithm demo

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture h is compatible with classroom 1**

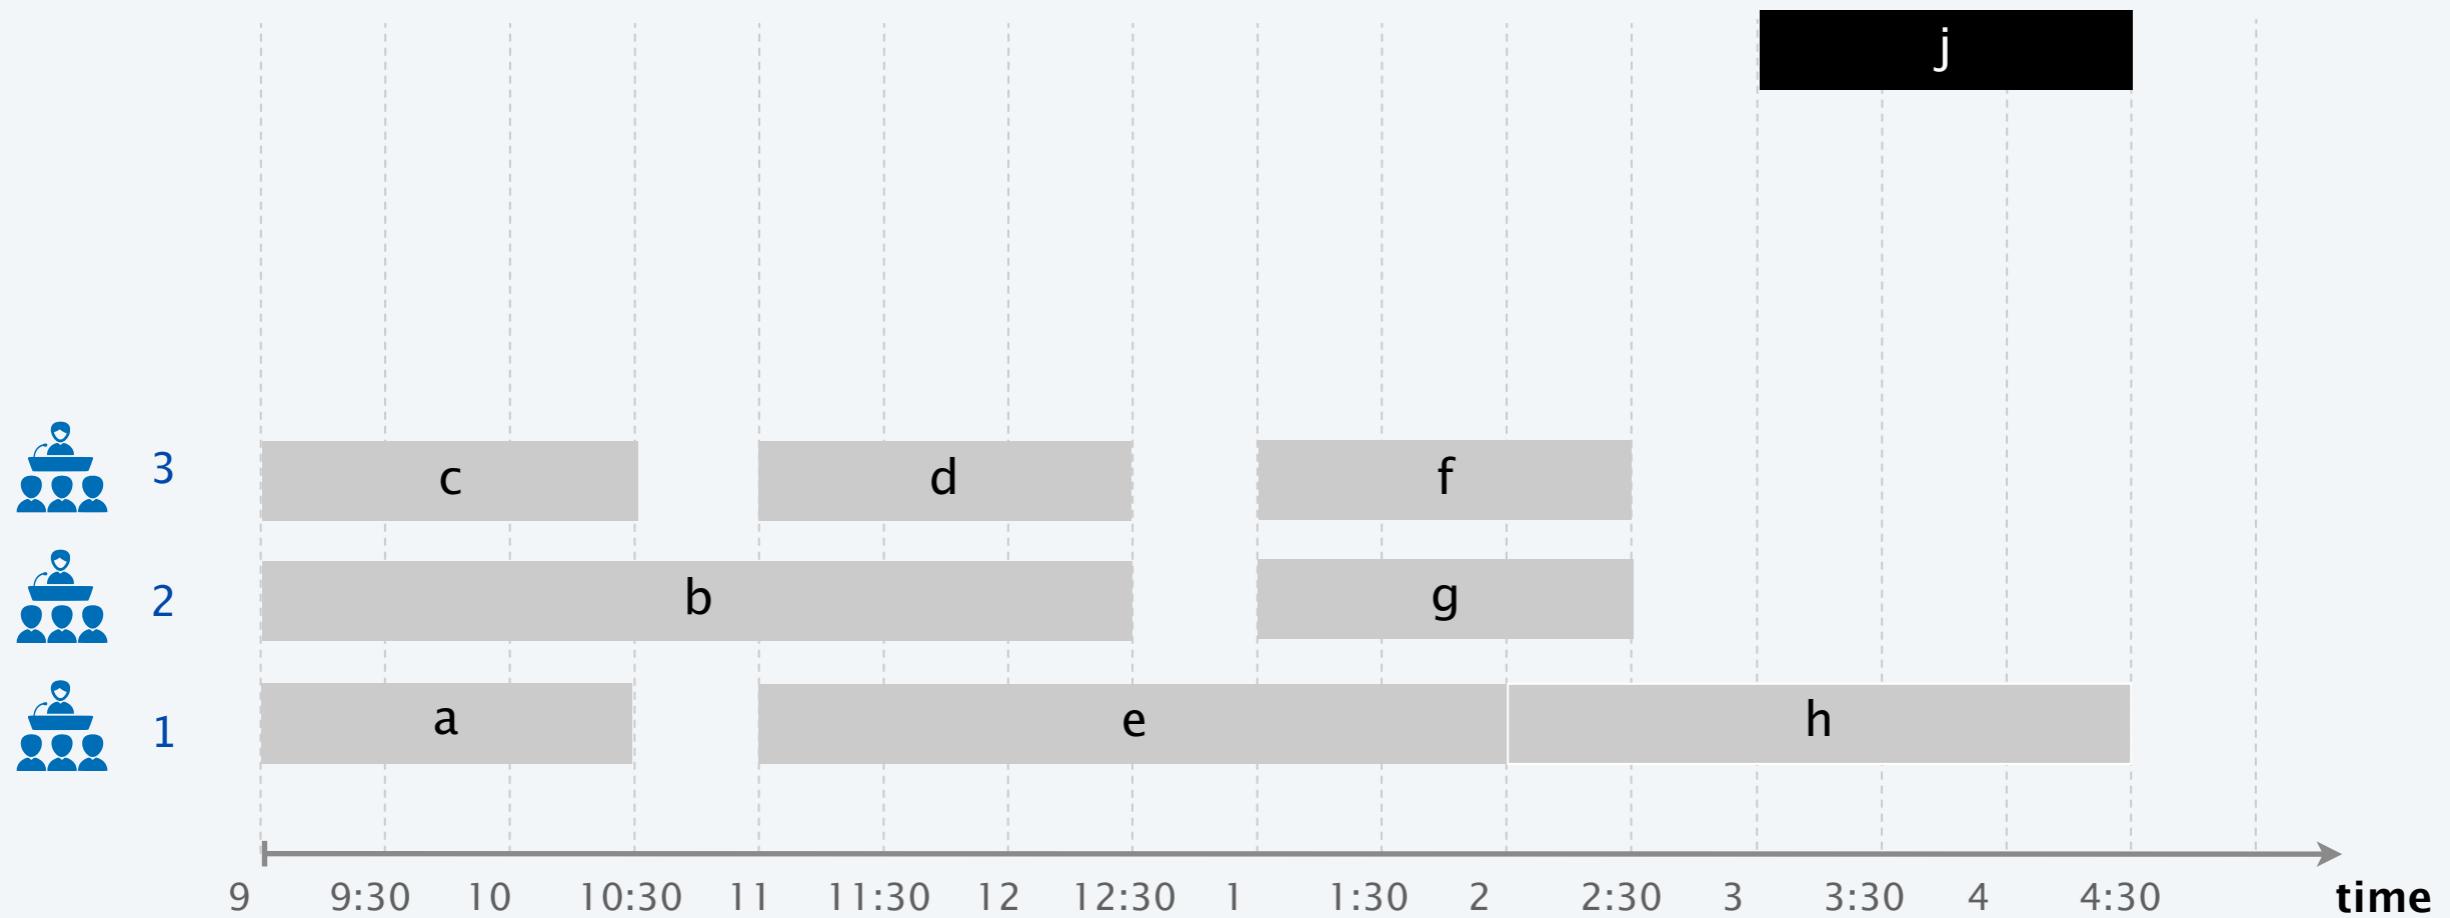


# Earliest-start-time-first algorithm demo

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture j is compatible with classrooms 2 and 3**

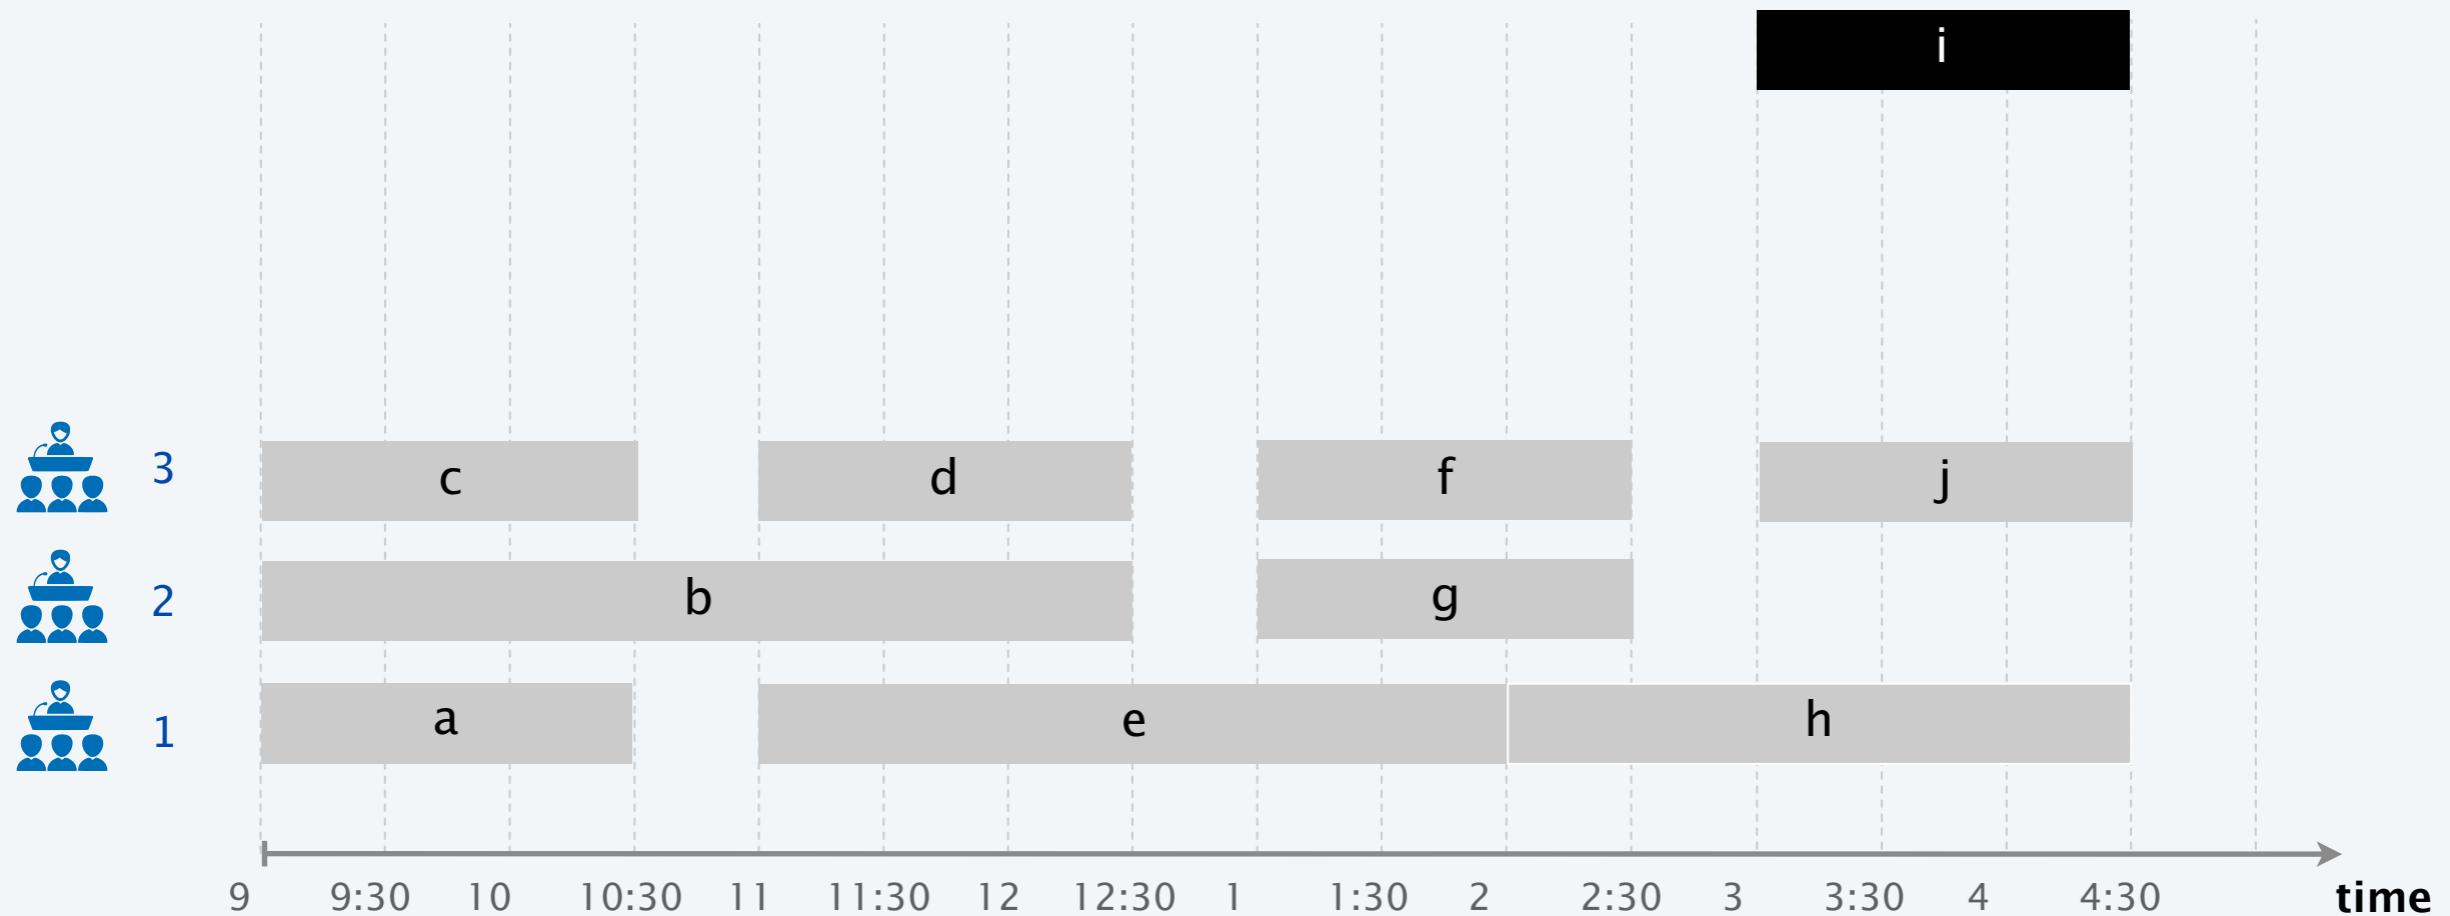


# Earliest-start-time-first algorithm demo

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

**lecture i is compatible with classroom 2**

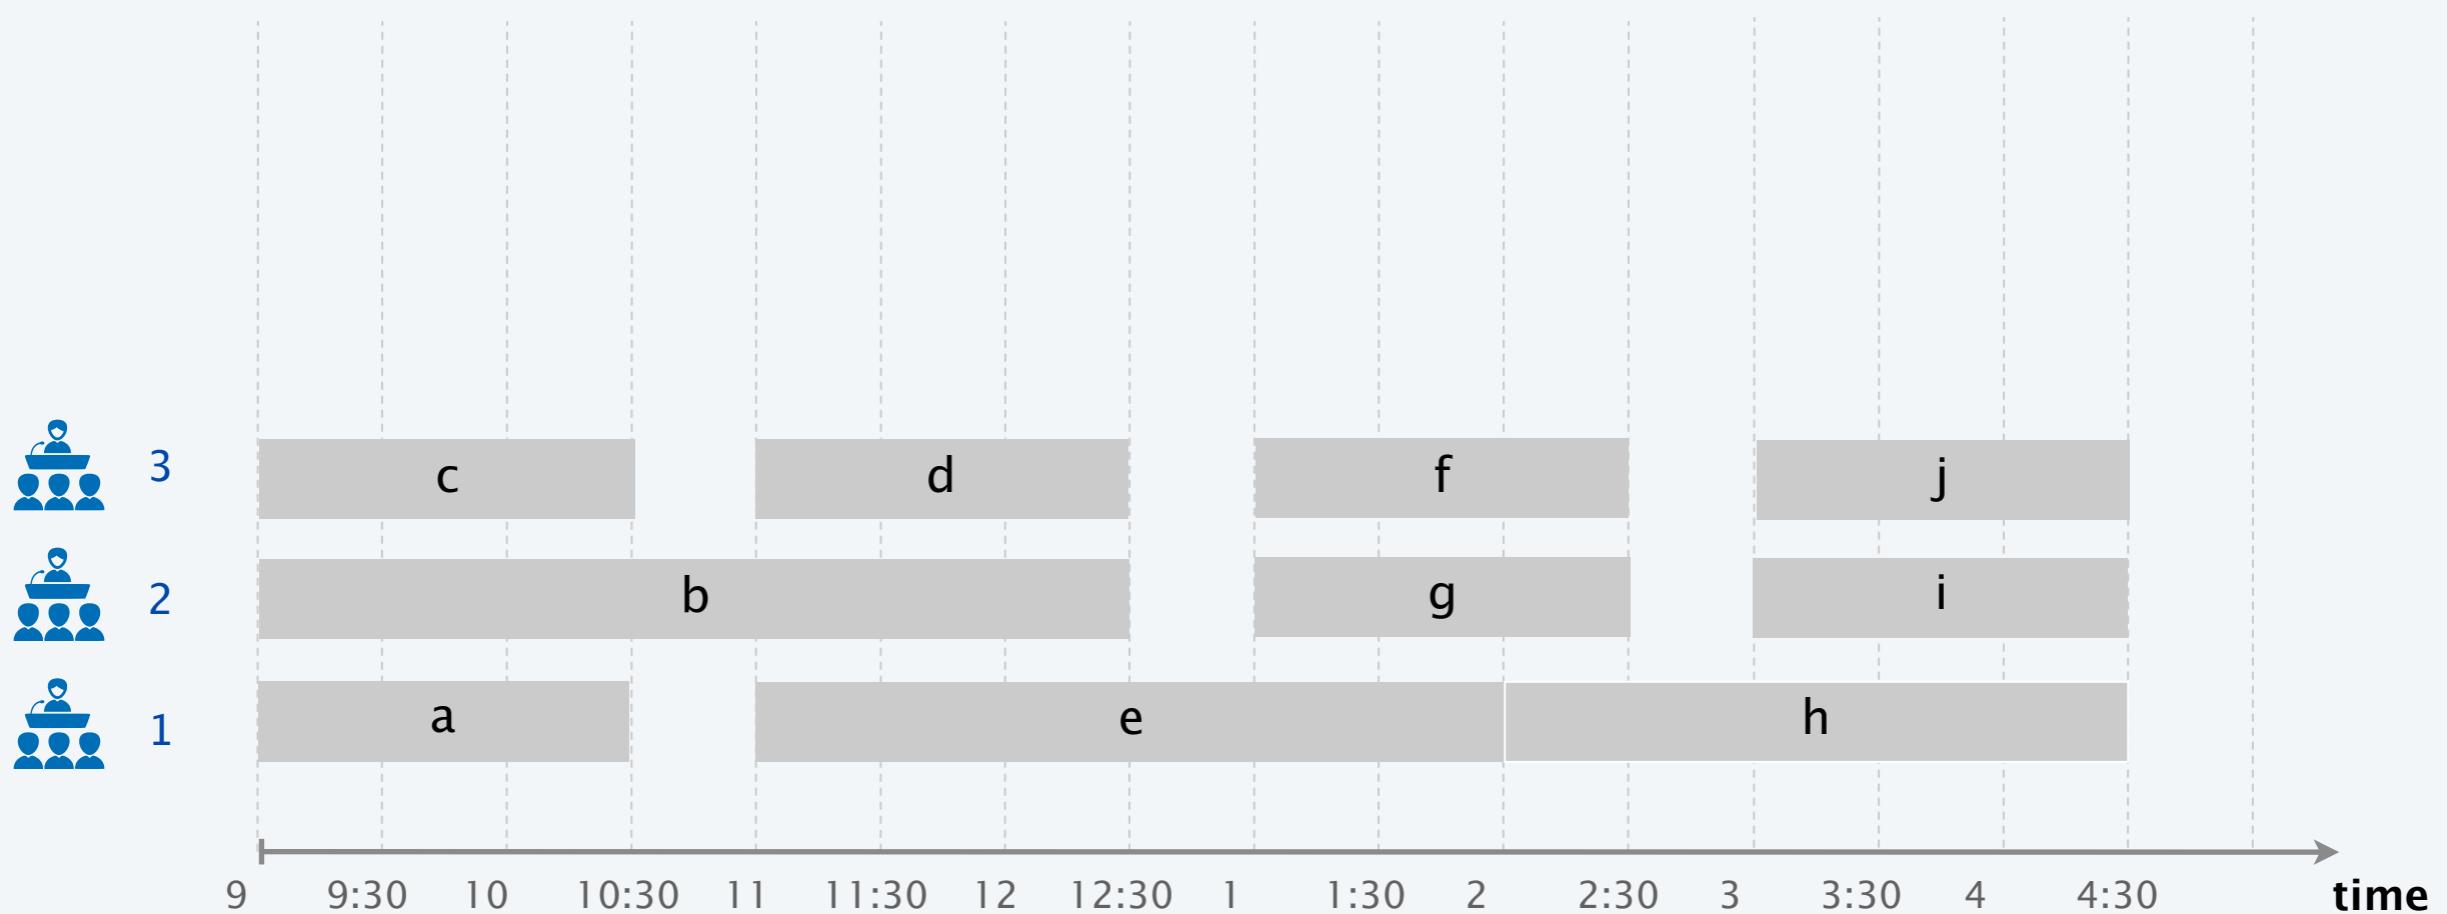


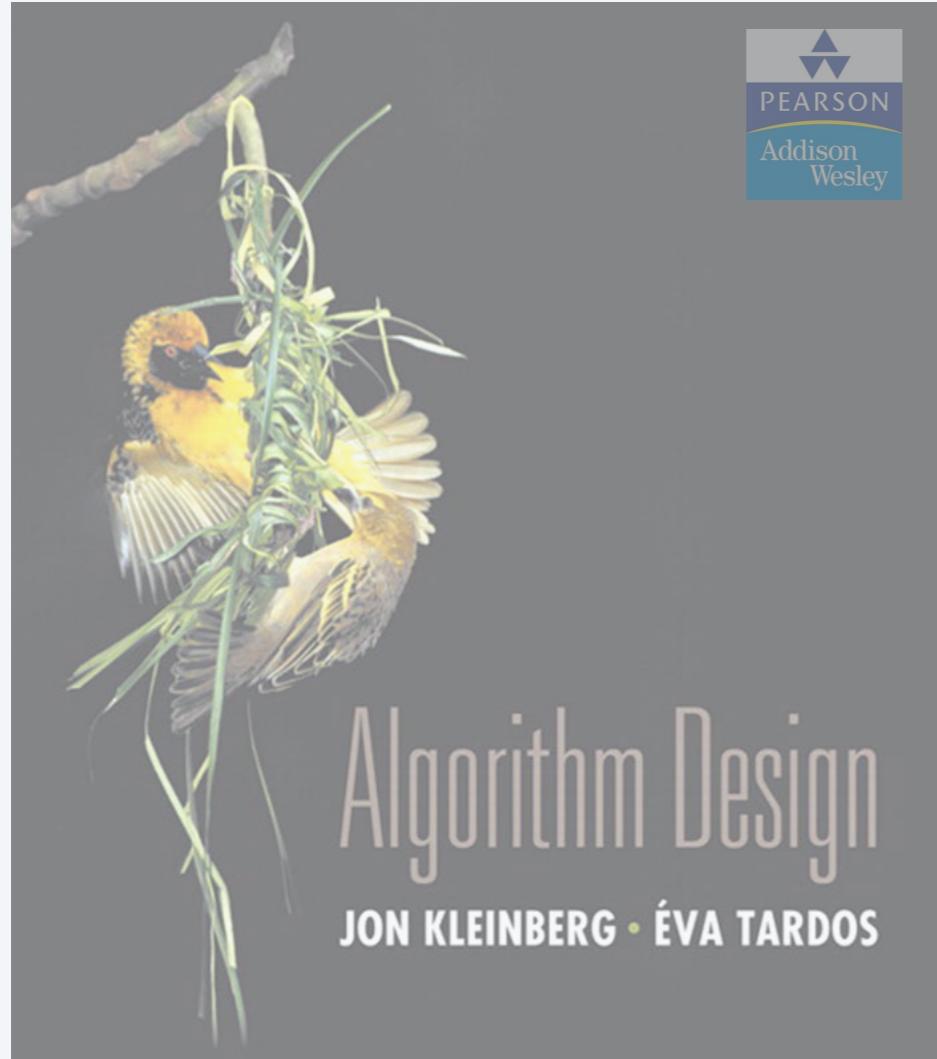
# Earliest-start-time-first algorithm demo

Consider lectures in order of start time:

- Assign next lecture to any compatible classroom (if one exists).
- Otherwise, open up a new classroom.

done





## SECTION 4.2

# 4. GREEDY ALGORITHMS I

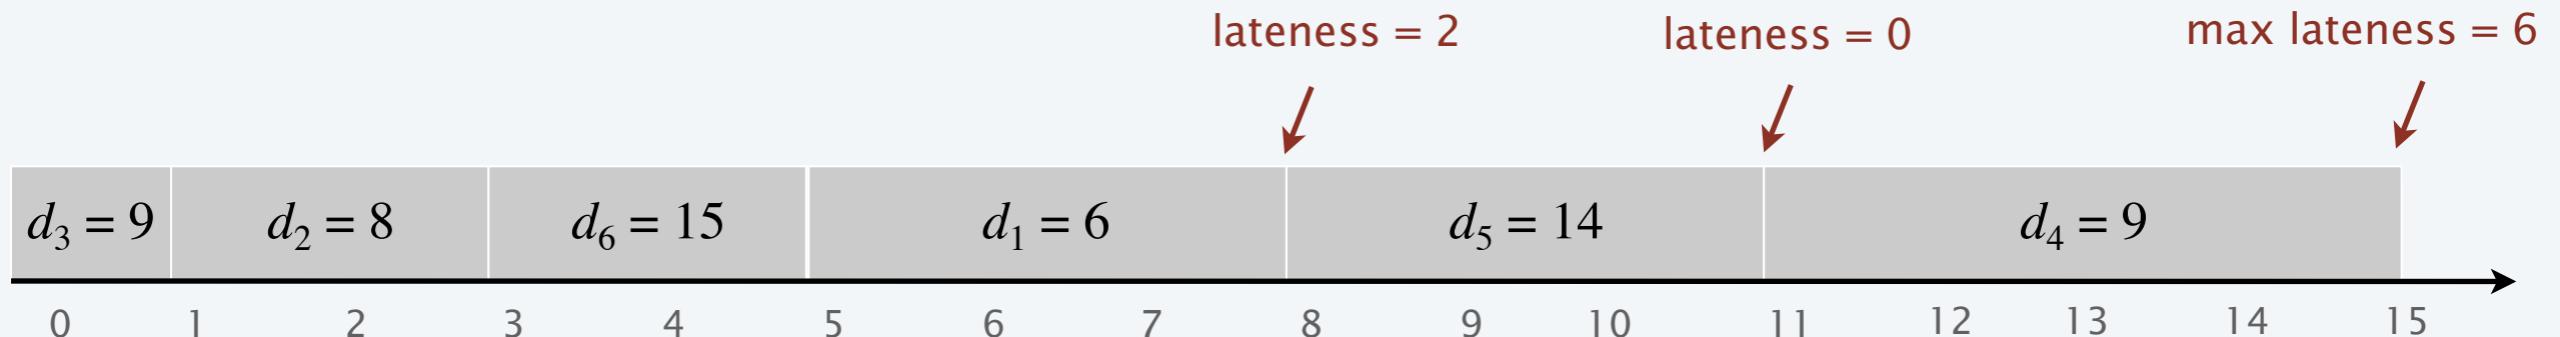
---

- ▶ *coin changing*
- ▶ *interval scheduling*
- ▶ *interval partitioning*
- ▶ ***scheduling to minimize lateness***
- ▶ *optimal caching*

# Scheduling to minimizing lateness

- Single resource processes one job at a time.
- Job  $j$  requires  $t_j$  units of processing time and is due at time  $d_j$ .
- If  $j$  starts at time  $s_j$ , it finishes at time  $f_j = s_j + t_j$ .
- Lateness:  $\ell_j = \max \{ 0, f_j - d_j \}$ .
- Goal: schedule all jobs to minimize **maximum** lateness  $L = \max_j \ell_j$ .

	1	2	3	4	5	6
$t_j$	3	2	1	4	3	2
$d_j$	6	8	9	9	14	15





Schedule jobs according to some natural order. Which order minimizes the maximum lateness?

- A. [shortest processing time] Ascending order of processing time  $t_j$ .
- B. [earliest deadline first] Ascending order of deadline  $d_j$ .
- C. [smallest slack] Ascending order of slack:  $d_j - t_j$ .
- D. None of the above.

# Minimizing lateness: earliest deadline first

EARLIEST-DEADLINE-FIRST  $(n, t_1, t_2, \dots, t_n, d_1, d_2, \dots, d_n)$

SORT jobs by due times and renumber so that  $d_1 \leq d_2 \leq \dots \leq d_n$ .

$t \leftarrow 0$ .

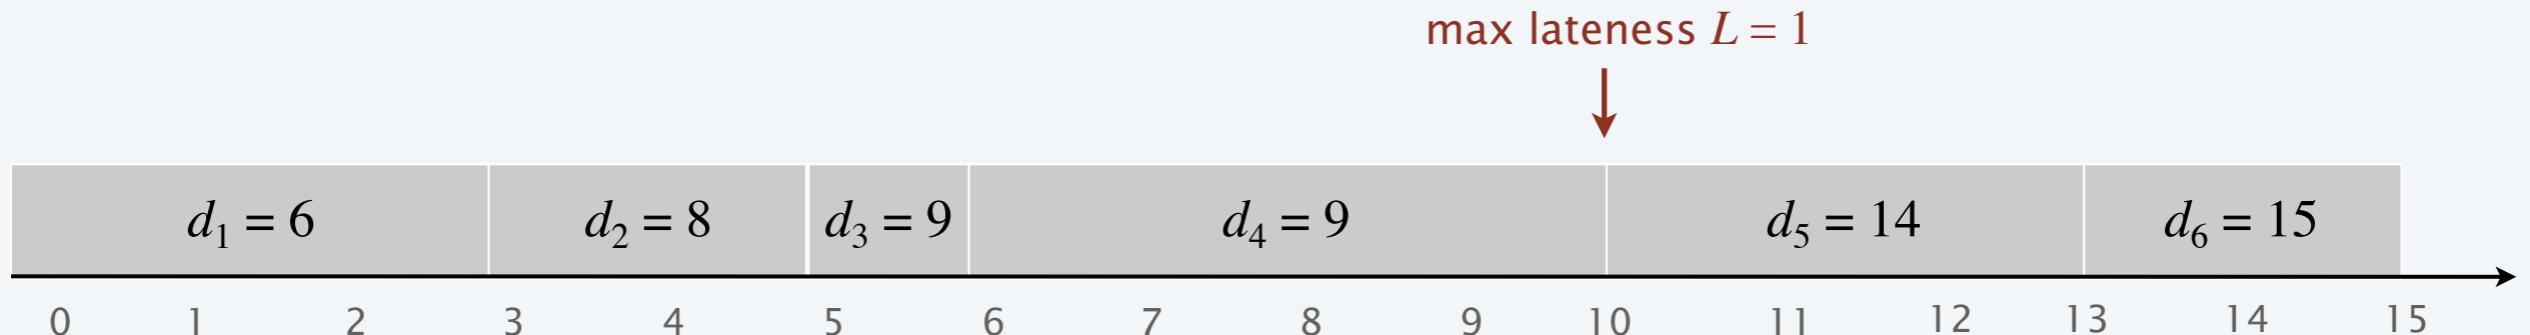
FOR  $j = 1$  TO  $n$

    Assign job  $j$  to interval  $[t, t + t_j]$ .

$s_j \leftarrow t$ ;  $f_j \leftarrow t + t_j$ .

$t \leftarrow t + t_j$ .

RETURN intervals  $[s_1, f_1], [s_2, f_2], \dots, [s_n, f_n]$ .



## Minimizing lateness: no idle time

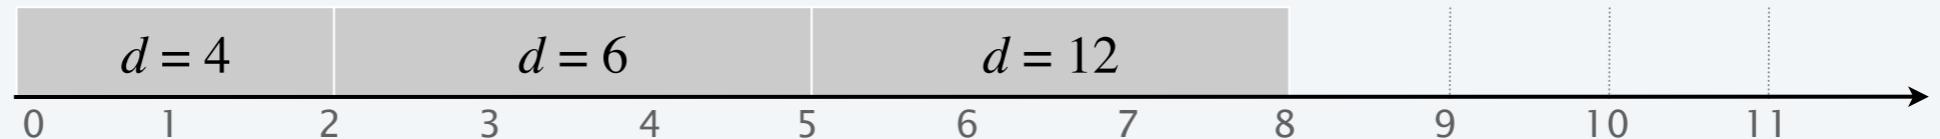
---

Observation 1. There exists an optimal schedule with no **idle time**.

an optimal schedule



an optimal schedule  
with no idle time

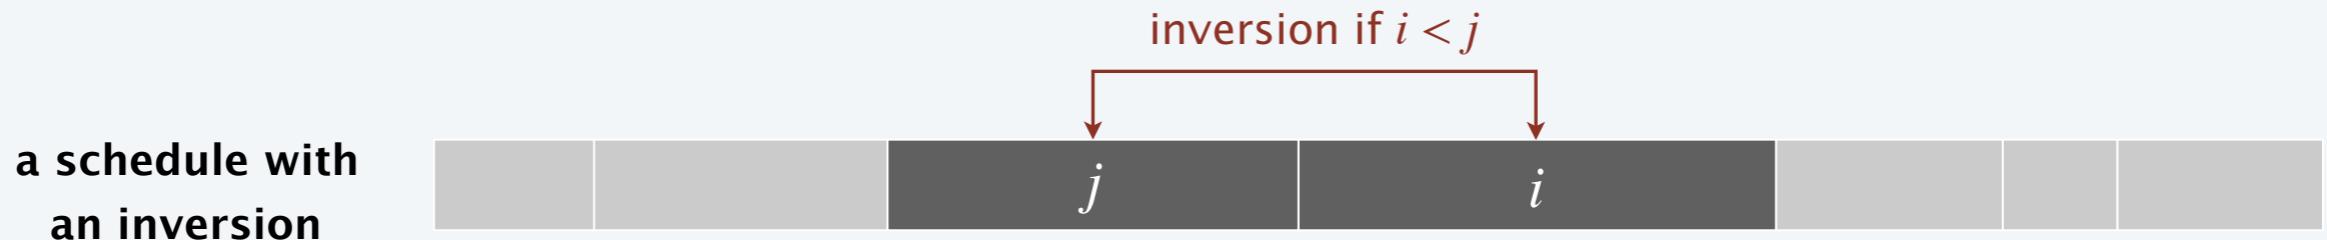


Observation 2. The earliest-deadline-first schedule has no idle time.

## Minimizing lateness: inversions

---

**Def.** Given a schedule  $S$ , an **inversion** is a pair of jobs  $i$  and  $j$  such that:  $i < j$  but  $j$  is scheduled before  $i$ .



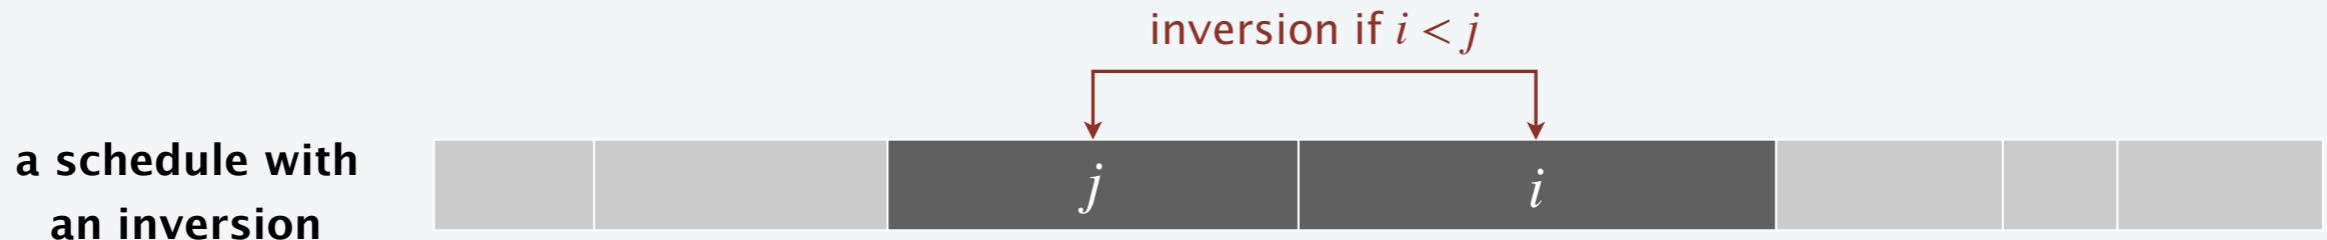
recall: we assume the jobs are numbered so that  $d_1 \leq d_2 \leq \dots \leq d_n$

**Observation 3.** The earliest-deadline-first schedule is the unique idle-free schedule with no inversions.



## Minimizing lateness: inversions

**Def.** Given a schedule  $S$ , an **inversion** is a pair of jobs  $i$  and  $j$  such that:  $i < j$  but  $j$  is scheduled before  $i$ .

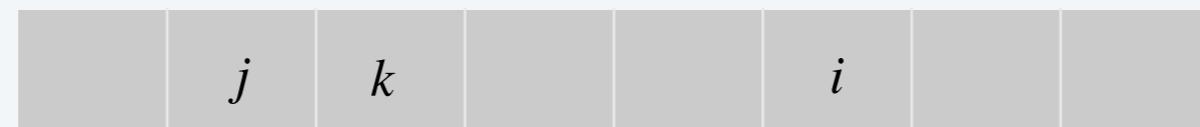


recall: we assume the jobs are numbered so that  $d_1 \leq d_2 \leq \dots \leq d_n$

**Observation 4.** If an idle-free schedule has an inversion, then it has an adjacent inversion.

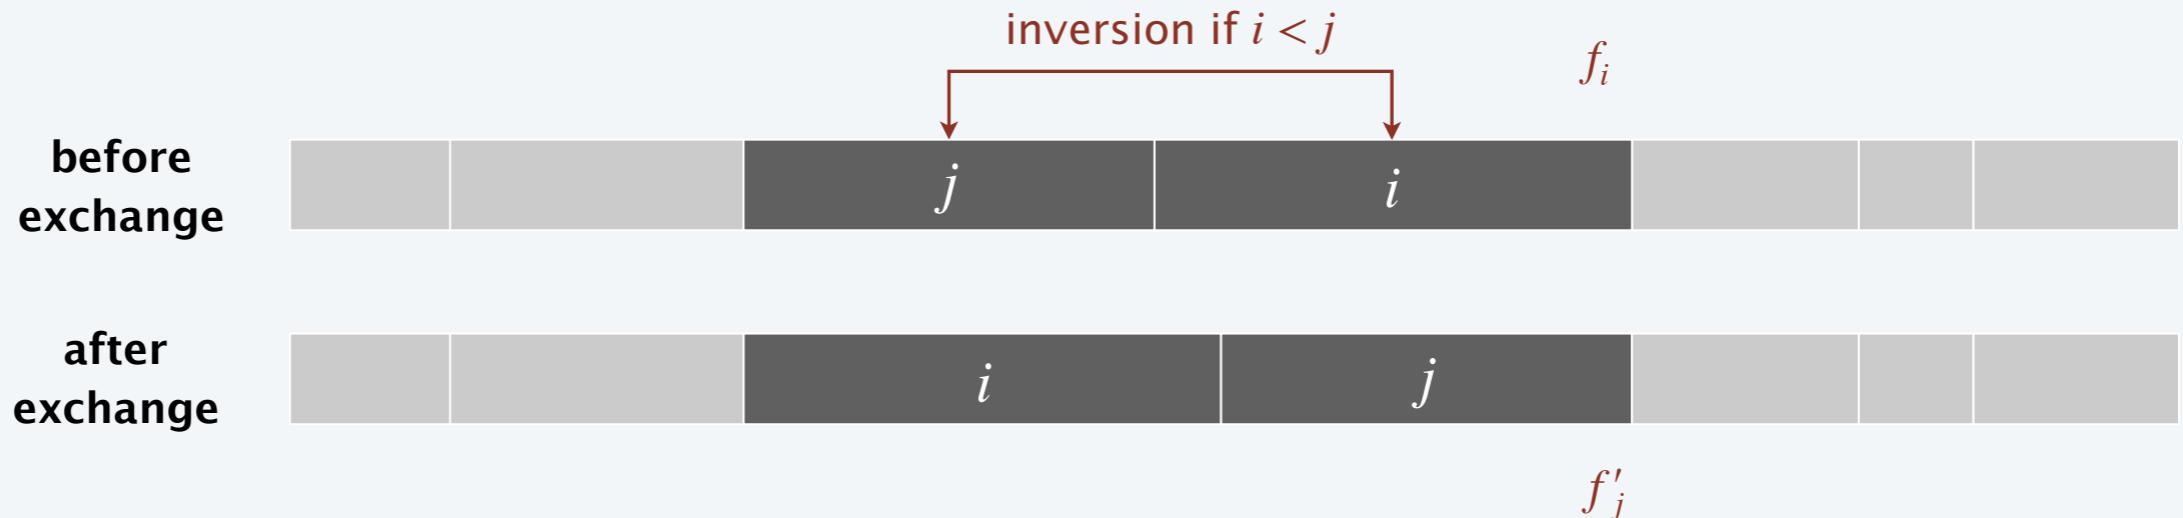
Pf. two inverted jobs scheduled consecutively

- Let  $i-j$  be a closest inversion.
- Let  $k$  be element immediately to the right of  $j$ .
- Case 1.  $[j > k]$  Then  $j-k$  is an adjacent inversion.
- Case 2.  $[j < k]$  Then  $i-k$  is a closer inversion since  $i < j < k$ .  $\diamond$



## Minimizing lateness: inversions

**Def.** Given a schedule  $S$ , an **inversion** is a pair of jobs  $i$  and  $j$  such that:  $i < j$  but  $j$  is scheduled before  $i$ .



**Key claim.** Exchanging two adjacent, inverted jobs  $i$  and  $j$  reduces the number of inversions by 1 and does not increase the max lateness.

**Pf.** Let  $\ell$  be the lateness before the swap, and let  $\ell'$  be it afterwards.

- $\ell'_k = \ell_k$  for all  $k \neq i, j$ .
- $\ell'_i \leq \ell_i$ .
- If job  $j$  is late,  $\ell'_j = f'_j - d_j \leftarrow \text{definition}$   
 $= f_i - d_j \leftarrow j \text{ now finishes at time } f_i$   
 $\leq f_i - d_i \leftarrow i < j \Rightarrow d_i \leq d_j$   
 $\leq \ell_i . \quad \leftarrow \text{definition}$

# Minimizing lateness: analysis of earliest-deadline-first algorithm

---

**Theorem.** The earliest-deadline-first schedule  $S$  is optimal.

Pf. [by contradiction]

Define  $S^*$  to be an optimal schedule with the fewest inversions.

optimal schedule can  
have inversions

- Can assume  $S^*$  has no idle time. ← Observation 1
- Case 1. [  $S^*$  has no inversions ] Then  $S = S^*$ . ← Observation 3
- Case 2. [  $S^*$  has an inversion ]
  - let  $i-j$  be an adjacent inversion ← Observation 4
  - exchanging jobs  $i$  and  $j$  decreases the number of inversions by 1 without increasing the max lateness ← key claim
  - contradicts “fewest inversions” part of the definition of  $S^*$  ✎

# Greedy analysis strategies

---

**Greedy algorithm stays ahead.** Show that after each step of the greedy algorithm, its solution is at least as good as any other algorithm's.

**Structural.** Discover a simple “structural” bound asserting that every possible solution must have a certain value. Then show that your algorithm always achieves this bound.

**Exchange argument.** Gradually transform any solution to the one found by the greedy algorithm without hurting its quality.

**Other greedy algorithms.** Gale–Shapley, Kruskal, Prim, Dijkstra, Huffman, ...



# GOOGLE'S FOO.BAR CHALLENGE



Quantum antimatter fuel comes in small pellets, which is convenient since the many moving parts of the LAMBCHOP each need to be fed fuel one pellet at a time. However, minions dump pellets in bulk into the fuel intake. You need to figure out the most efficient way to sort and shift the pellets down to a single pellet at a time.

The fuel control mechanisms have three operations:

- Add 1 fuel pellet
- Remove 1 fuel pellet
- Divide the entire group of fuel pellets by 2 (due to the destructive energy released when a quantum antimatter pellet is cut in half, the safety controls will only allow this to happen if there is an even number of pellets)

Write a function called `answer(n)` which takes a positive integer `n` as a string and returns the minimum number of operations needed to transform the number of pellets to 1.